

智能体体系结构（Agent Architecture）

毛文吉

中国科学院自动化研究所

Agent Architecture

An intelligent agent is a computer system capable of *flexible autonomous action* ...

- Reactive

- Responsive to changes occurred in its environment

- Pro-active/Deliberative

- Take initiative and act purposefully toward achieving goals

Issues need to address to build agent-based systems ...

- Three types of agent architecture:

- Deliberative

- Reactive

- Hybrid

What is an Agent Architecture?

- Agent “Shell”, separates *structure* from variable *content*
- The principles/structure to be reused from task to task
 - *Why rediscover these from scratch each time?*
- Some well-known agent architectures:
 - *IRMA (Intelligent Resource-bounded Machine Arch.)*
 - *PRS, dMARS, JACK (started from PRS)*
 - *TouringMachines, InteRPaP*
 - *Soar (Cognitive Arch.)*

A Brief History

- Before ~1985, almost all agents designed within AI were symbolic reasoning agents
 - Develop a *deliberative* agent or agent architecture that contains explicit *symbolic representation* of the world, and makes decisions (e.g. about what actions to perform) via *symbolic reasoning*
- Around 1985~1990, problems with symbolic reasoning led to a reaction against this – the so-called reactive agents
- Since late 90s, a number of alternatives proposed: hybrid architectures, which attempt to combine the *best* of reasoning and reactive architectures

Outline

- Reactive agent architectures
- Deliberative agent architectures
- Hybrid agent architectures
- Soar agent architecture

Outline

- Reactive agent architectures
- Deliberative agent architectures
- Hybrid agent architectures
- Soar agent architecture

Reactive (反应式) Architectures

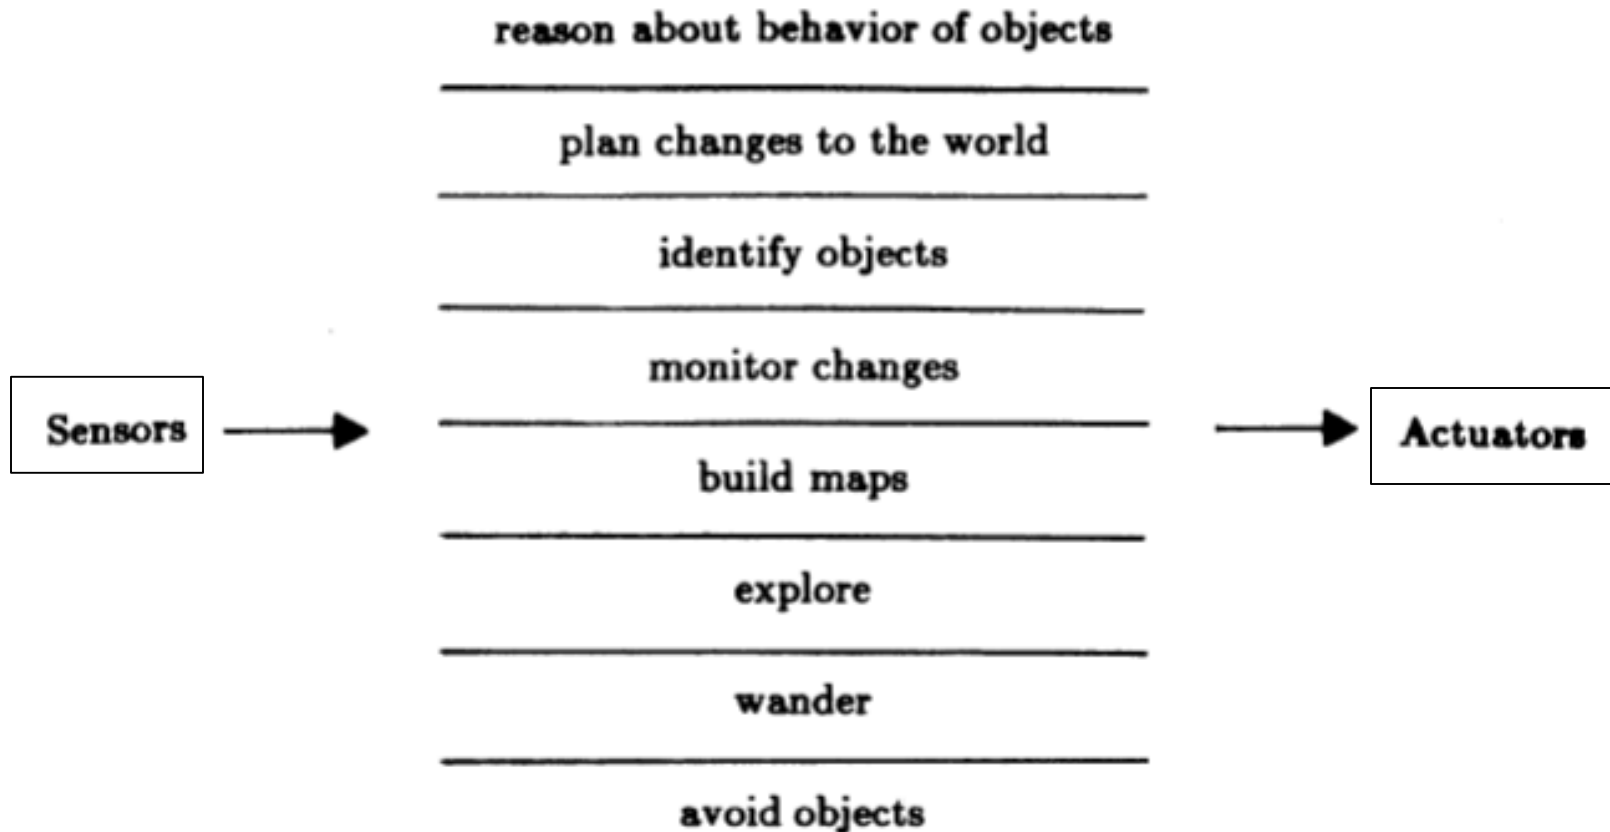
- The unsolved problems associated with symbolic AI have led to the development of reactive architectures
- Rodney Brooks (1986) is the central representative of this view, who has put forward three theses:
 - Intelligent behavior can be generated without explicit representations of the kind that symbolic AI proposes
 - Intelligent behavior can be generated without explicit abstract reasoning of the kind that symbolic AI proposes
 - Intelligence is an emergent property of certain complex systems

Brooks: Subsumption (包容式) Architecture

- A subsumption architecture is a hierarchy of task-accomplishing *behaviors*
- Each behavior is a *rather simple* rule-like structure
- Each behavior '*competes*' with others to exercise control over the agent
- Lower layers represent more primitive kinds of behavior (such as avoiding obstacles), and have *precedence* (优先权) over layers further up the hierarchy
- The resulting systems are *extremely* simple in terms of the amount of computation cost

Layered Control in Subsumption Architecture

- Task decomposition of mobile robot control system:



Reactive Architectures: Features

Advantages:

- Simplicity (fast response, real-time efficiency)
- Economy
- Computational tractability
- Robustness against failure

Limitations:

- “*Short-term*” view (短视) for decision making
- Difficult to make reactive agents learn from experience
- Hard to engineer agent with large numbers of behaviors
 - Dynamics of the interactions between different behaviors become too complex to be understood

Outline

- Reactive agent architectures
- Deliberative agent architectures
- Hybrid agent architectures
- Soar agent architecture

Deliberative（慎思式）Architectures

■ BDI-based agent architectures

- *IRMA (Intelligent Resource-bounded Machine Arch.)* [Bratman 88]
- *PRS-type architecture: PRS [87], JACK [01], dMARS [04]*
- *RAP (Reactive Action Packages)*
- *TCA (Task Control Architecture)*

■ BDI-based teamwork models

- *GRATE** [Jennings 95]
- *COLLAGEN* [Rich & Sindner 97]
- *STEAM* [Tambe 97], *TEAMCORE* [Tambe et al 00]

... ..

Outline

- Reactive agent architectures
- Deliberative agent architectures
- Hybrid agent architectures
- Soar agent architecture

Hybrid (混合式) Architectures

- Many researchers have argued that *neither* a completely deliberative *nor* completely reactive architecture is suitable for building agents
- A hybrid architecture attempts to build an agent system out of *two* (or *more*) subsystems by marrying them:
 - Deliberative sub-structure
Using a symbolic model to develop plans and make decisions
 - Reactive sub-structure
Capable of reacting to events without complex reasoning

Hybrid Architectures

- The reactive component is often given the precedence over the deliberative one
- Naturally lead to the idea of the layered architecture, of which *TouringMachines* (1992) and *InteRRaP* (1997) are representative examples
- An agent's control subsystems are arranged into a hierarchy, with higher layers dealing with information at increasing levels of abstraction

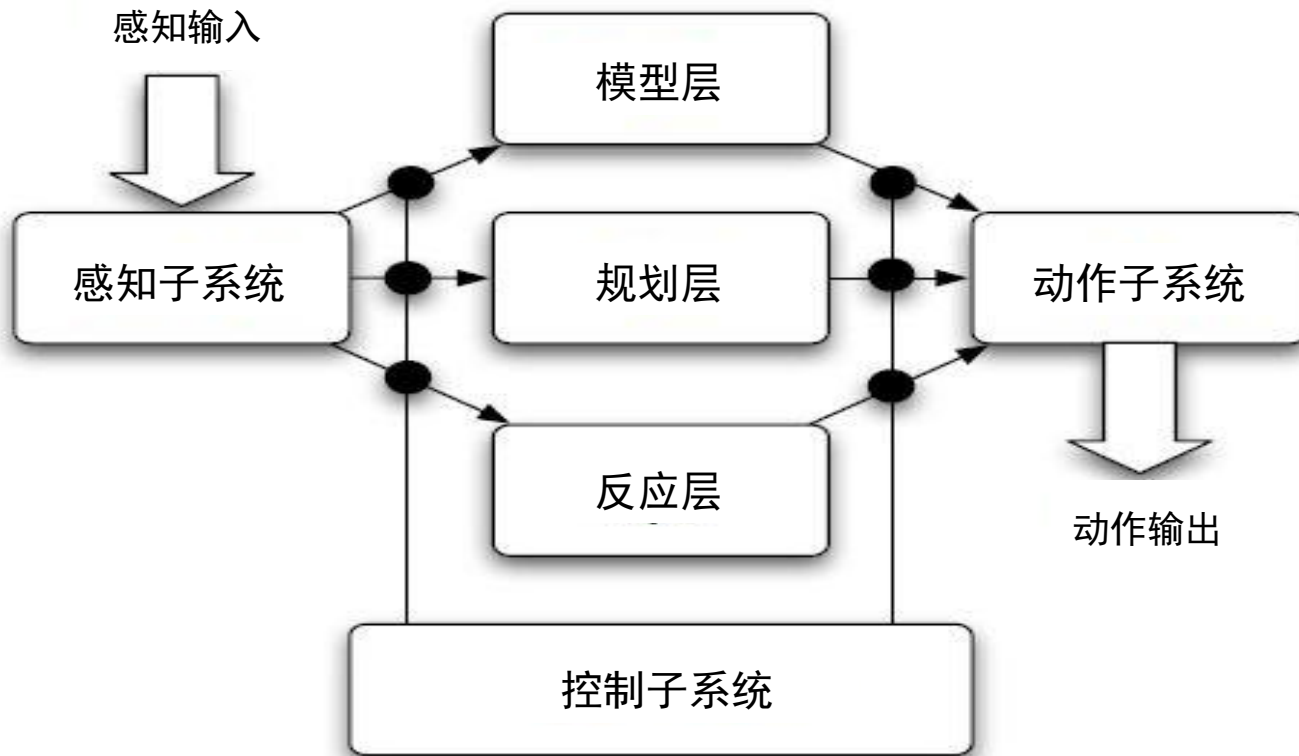
Hybrid Architectures

- Key problem: what kind of *control* framework to embed the agent's subsystems in, so as to manage *interactions* between layers
- Horizontal layering
 - Layers are each directly connected to the sensory input and action output
 - In effect, each layer itself acts like an agent, producing suggestions as to what action to perform
- Vertical layering
 - The sensory input and action output are each dealt with by at most one layer

Ferguson: TouringMachines

- The *TouringMachines* architecture consists of perception and action subsystems, which interface directly with the agent's environment, and three control layers:
 - Reactive layer
A set of situation-action rules using *subsumption architecture*
 - Planning layer
Construct *plans* and execute actions in order to achieve *goals*
 - Modeling layer
Symbolic representations of the '*cognitive state*' of other entities
- Communicate* with each other and are embedded in a control framework, which mediates between the layers using control rules


TouringMachines Architecture



Hybrid Architectures

- Suppose totally n layers with m possible actions suggested by each layer

Interactions: m^n




水平层次结构

(a) Horizontal layering


Introduce bottleneck in central control system

动作输出 Interactions: $m^2(n-1)$



感知输入

垂直层次结构
(b) Vertical layering
(One pass control)



感知输入 动作输出

(c) Vertical layering
(Two pass control)

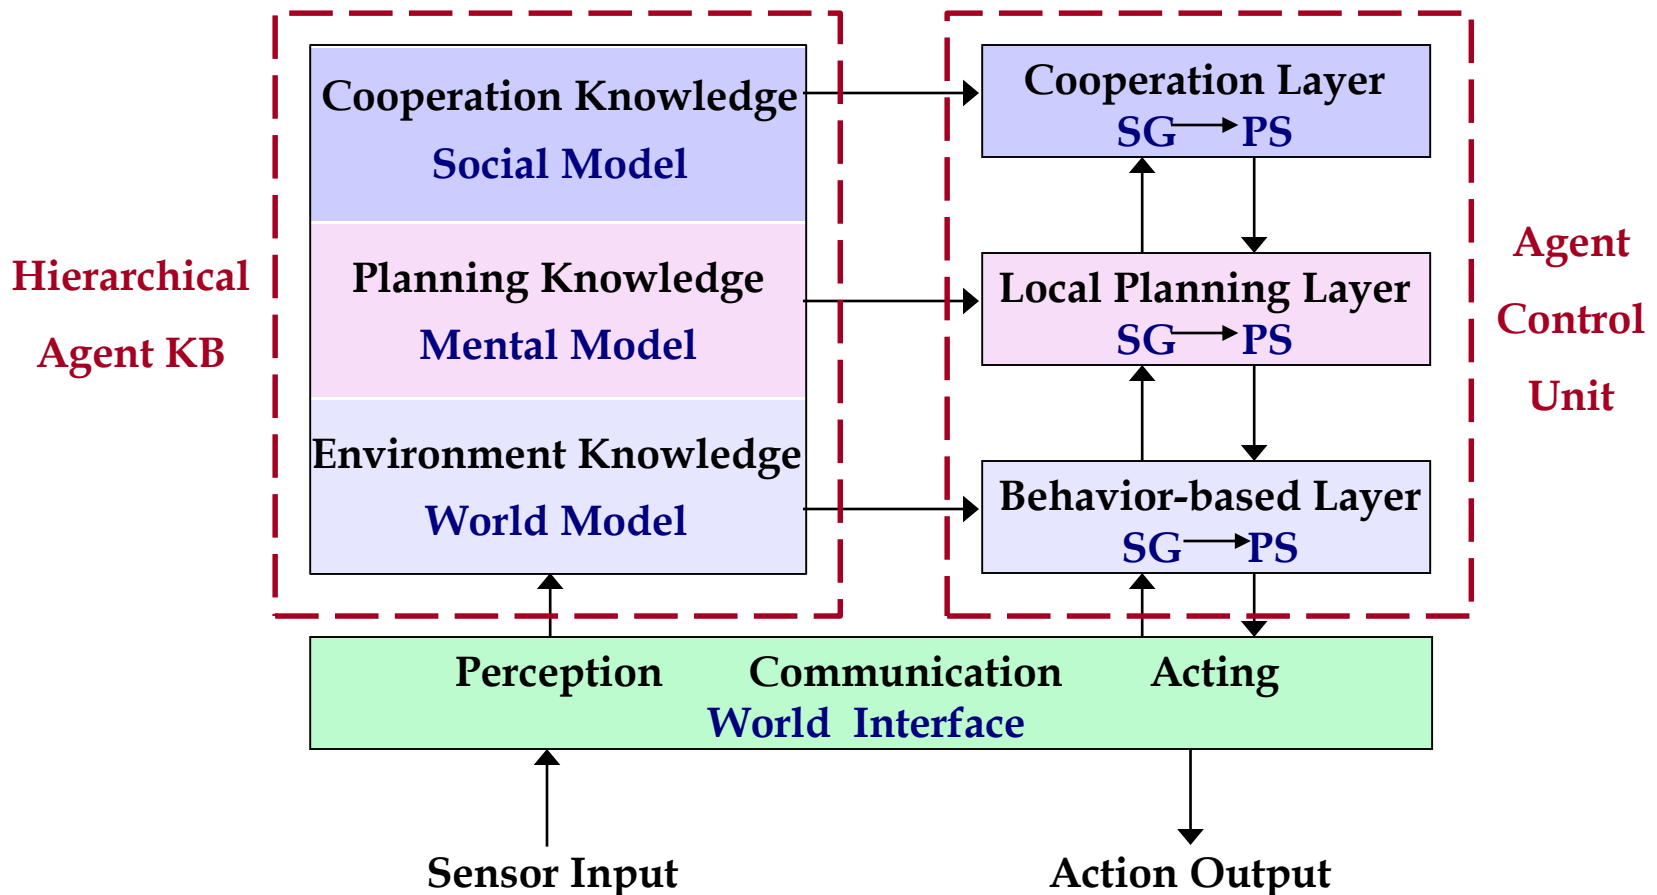
Not fault tolerant to layer failure

Müller: InteRRaP

- *InteRRaP* is a vertically layered two-pass architecture
- Agent control components
 - Each control layer consists of two subprocesses:
 - SG: situation recognition / goal activation
 - PS: planning, scheduling and execution
 - Lowest layer represents *reactive* component and higher layers *deliberative* components
- Hierarchical knowledge bases
 - World model: Information about the environment
 - Mental model: Plan knowledge and actions of agent itself
 - Social model: Maintaining models of other agents (i.e. plans and actions of other agents)

InteRRaP Architecture

- *InteRRaP* is a vertically layered two-pass architecture



InteRRaP Control Function and KB

Control Component	Corresponding Knowledge base	Function
Cooperation layer	Cooperation knowledge Plans of other agents	Generate joint plans that satisfy the goals of a number of agents, in response to request from local planning component
Local planning layer	Local planning knowledge Plan library of agent	Generate single-agent plans in response to requests from the behavior-based component
Behavior-based layer	Raw environment knowledge Beliefs of agent	Implement and control the basic reactive capability of the agent
World interface	World model Agent's perception	Manage the interface between the agent and its environment

Hybrid Architectures: Features

- Have been the most popular agent architecture available
- Interactions between layers are partly alleviated in InteRRaP

Main problem with layered architectures:

- *Lack* the conceptual and semantic *clarity*
- Arguably a pragmatic solution

Comparison: Different Approaches

	Reactive Architectures	Deliberative Architectures	Hybrid Architectures
Technology consensus	<i>No consensus</i>	<i>Yes</i>	<i>No clear consensus, but lots of similarity</i>
Methodology and theory	<i>No methodology Only isolated islands of theory</i>	<i>Clear methodology A lot of theory</i>	<i>No real methodology and theory, an ad hoc one</i>
Advantage/Disadvantage	<i>Work best using simple behavior in unknown environments</i>	<i>Cannot work well in highly dynamic domains</i>	<i>A pragmatic solution for real-world applications</i>

Outline

- Reactive agent architectures
- Deliberative agent architectures
- Hybrid agent architectures
- Soar agent architecture

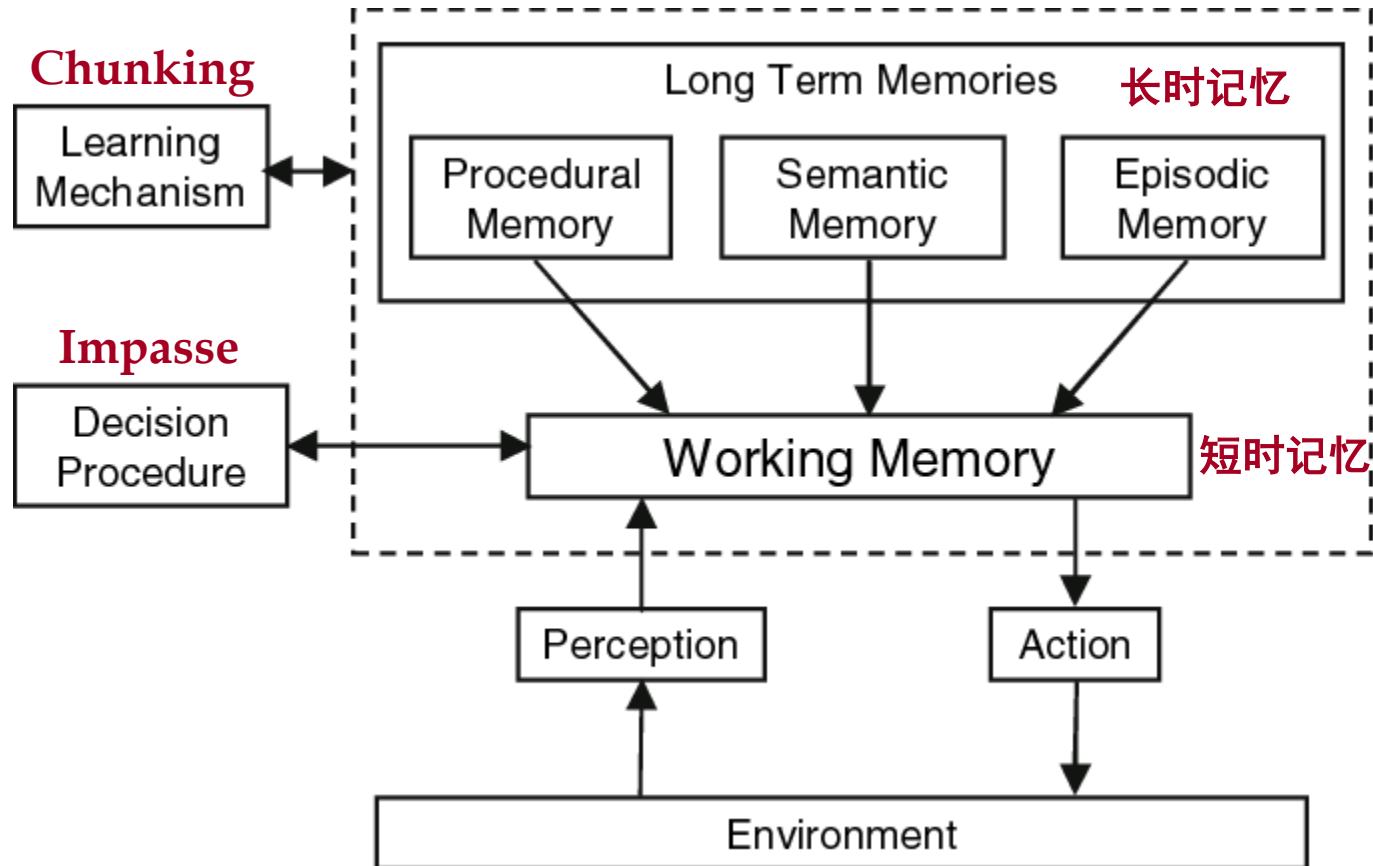
Soar Cognitive/Agent Architecture

- A production system that can be regarded as
 - A candidate Unified Theory of Cognition (*UTC*, Allen 90)
 - An architecture for intelligent behavior
- *Soar* (*State, Operator And Result*) provides an integration of
 - Knowledge, planning, reaction, search and learning
 - Within one efficient cognitive architecture
- Links
 - Learning *Soar*:
<https://soar.eecs.umich.edu>
 - Frequently ask questions:
<http://acs.ist.psu.edu/projects/soar-faq>

A Brief History of Soar

- First version operational in 1982
 - Written by Allen Newell, John Laird, Paul Rosenbloom at CMU
- Versions 1-5 written in *Lisp*
- Version 6 written in *C*
- Version 7 written in *C* with *Tcl/Tk*
- Version 7.0.4 *most commonly* used
- Version 7.2 true *multi-platform* version
- Version 8.3/4 are the latest multi-platform versions
- Version 9/9.1 has improved interfaces and more types of memories
- Version 9.6.1 is the latest current version

Original Soar Architecture



Major Design Principles

- A single framework for all tasks and subtasks (i.e. *problem spaces*)
- A single representation of permanent knowledge (i.e. *productions*)
- A single representation of temporary knowledge (i.e. *attribute/value* pairs)
- A single mechanism for generating goals (i.e. *subgoaling*)
- A single learning mechanism (i.e. *chunking*)
- All decisions are made at *run-time*, based on current sensory data and any relevant long-term knowledge

Some Definitions

Soar aims at supporting a full range of tasks and problem-solving methods, representing and using appropriate knowledge forms:

- *A goal* – is a desired situation
- *A state* – is a representation of a problem solving situation
- *A problem space* – is a set of states and operators for task
- *An operator* – transforms the state by some action

Knowledge Level

- Soar can be thought of as an engine for applying knowledge to situations to yield behavior
- Knowledge is encoded in production rules. A rule has conditions on its *LHS*, and actions on the *RHS*: $C \rightarrow A$

Example:

IF we are in the *hungry-thirsty* problem space, AND
current state is *thirsty*


THEN propose an operator to *apply* to current state
and call this operator “*drink*”

Propose drink.

```
sp {(state <s> ^problem-space <p> ^thirsty yes)
    (<p> ^name hungry-thirsty)
-> (<s> ^operator <o>)
    (<o> ^name drink)}
```

Problem Space Level

- Structure of the problem space for Eight Puzzle:



Adding Knowledge to Problem Space

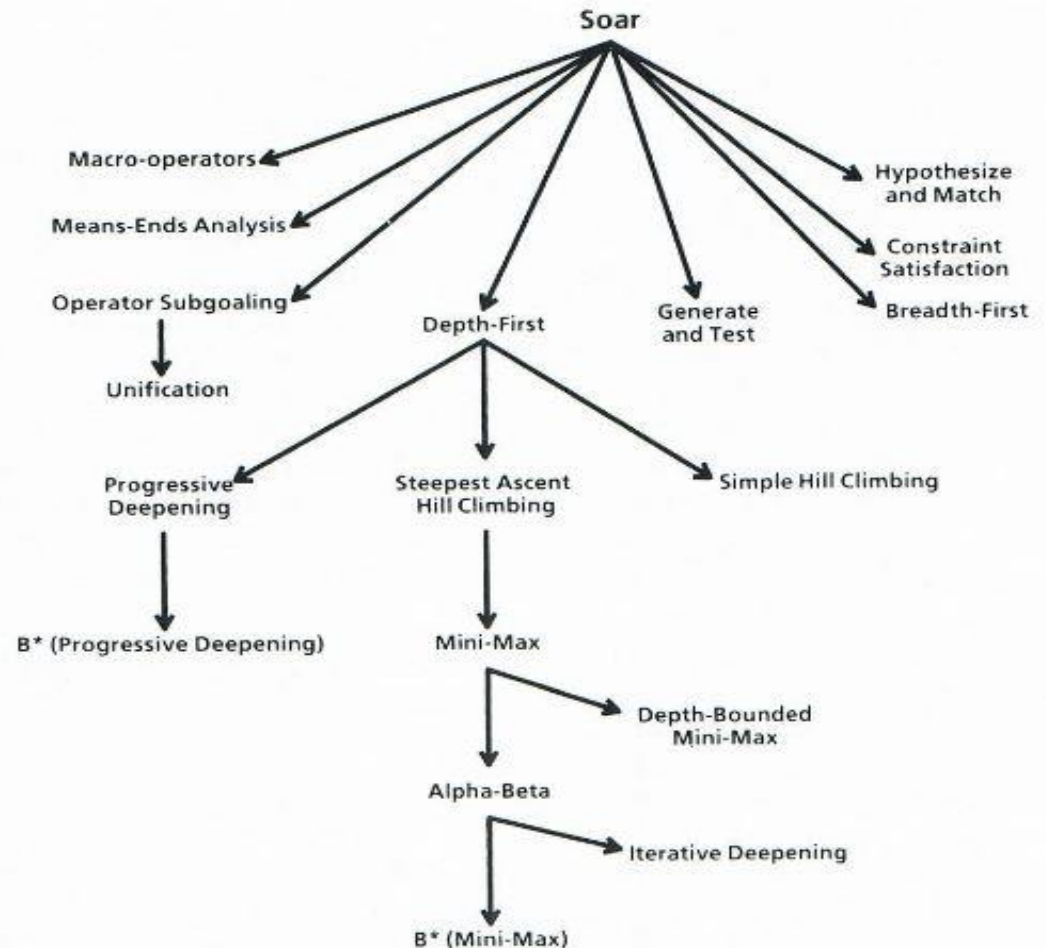
- In order to act, Soar must have knowledge of that domain (either given to it or learned)
- Domain knowledge can be divided into two categories:
 - Basic problem space knowledge:
Definitions of the state representation, the “legal move” operators, their applicability conditions, and their effects
 - Control knowledge:
Give guidance on choosing what to do, such as heuristics for solving problems in the domain

Problem Solving Methods

- Universal *weak methods* realized in Soar


- Weak methods require relatively *little knowledge* about the task

- All the standard weak methods are built on *heuristic search*



Processing Cycle and Impasse

- Once the situation is set up and running, the main activity is the repeated *selection* and then *application* of one operator after another



- If something prevents this processing cycle, e.g.
 - Soar knows of no operators to apply to that state (i.e. *no-change impasse*)
 - It knows of several, but has no knowledge of how to choose? (i.e. *tie impasse*)

In such cases, an *impasse* (僵局) occurs

Learning Mechanism: Chunking

- In Soar, *subgoals* are created *only* in response to impasse
- Resolving an *impasse* leads to learning
- Soar includes a simple, uniform learning mechanism, called *chunking* (组块)
- A *chunk* is a new production rule that summarizes the processing that was needed to resolve the *impasse*
- Whenever a result is returned from an *impasse*, a new rule is learned connecting the relevant parts of the pre-impasse situation with the result — next time a sufficiently similar situation occurs, the impasse is *avoided*

Many Applications

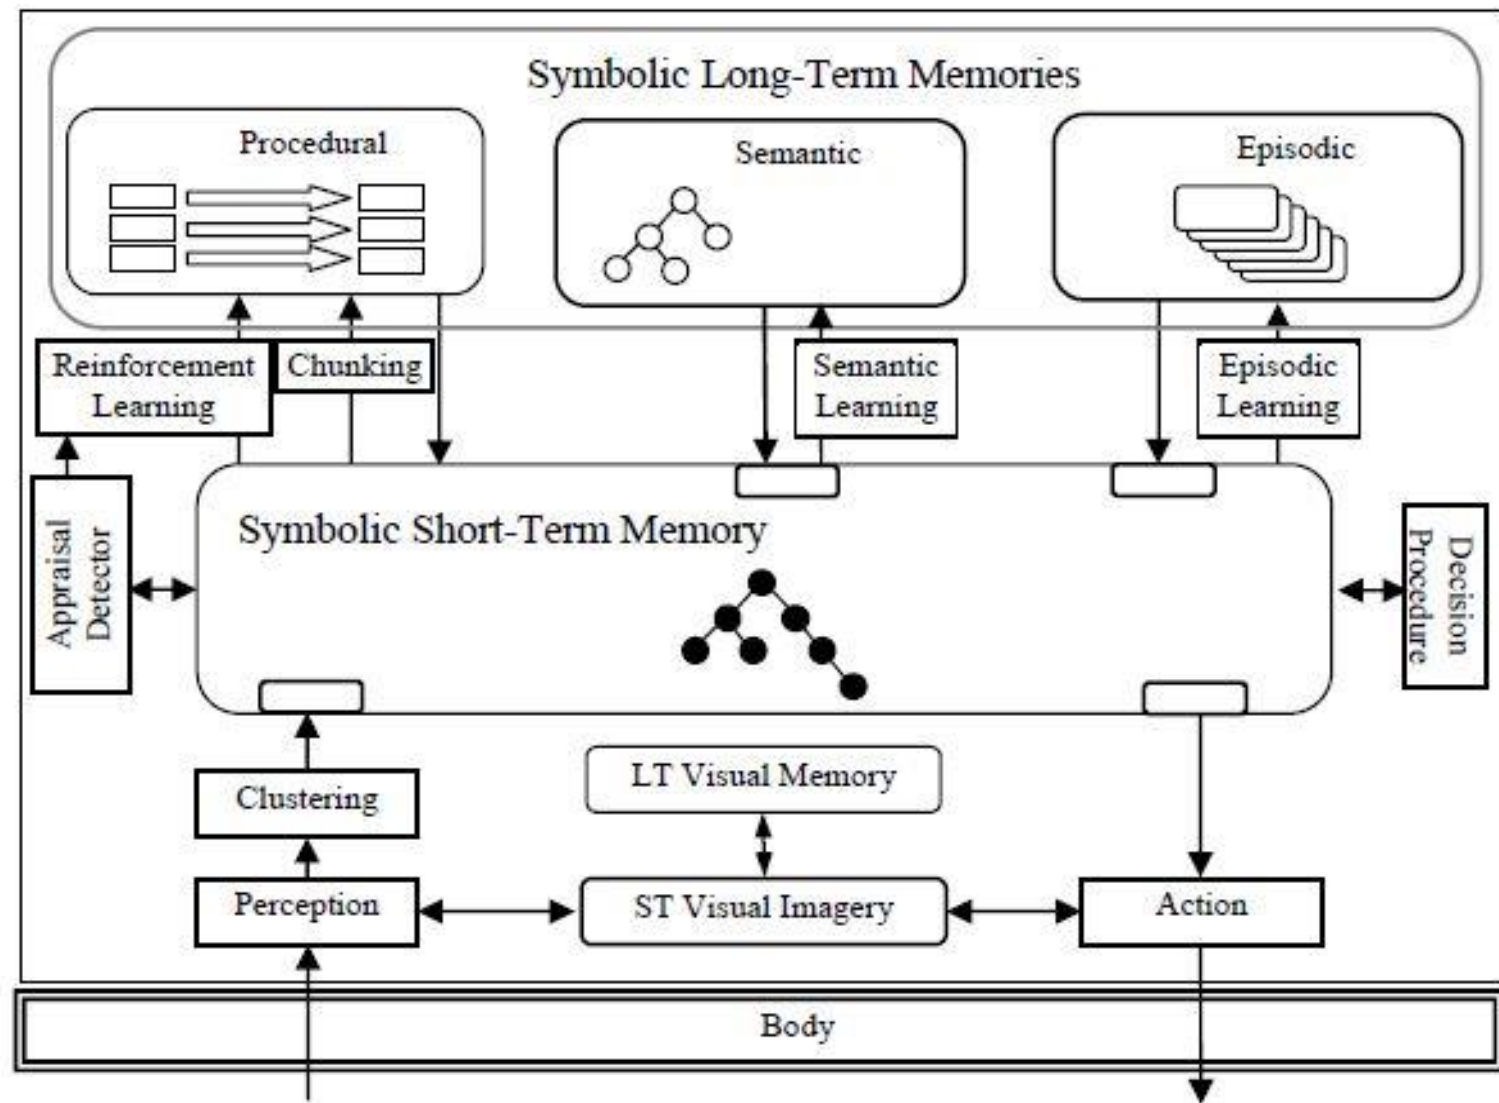
Work successfully in real-world agent-based applications:

- Computer games @ UoM
- Modeling social processes
- Learning by observation
- Virtual agents/humans – Immersive training
- Simulating emotions
- Electronic elves project – Teamwork @ USC
- Development tools
- Making Soar articulate / High-level language
- Active research community with *yearly* workshop
43nd Soar Workshop in June, 2023

Soar: Pros and Cons

- Soar remains the exemplar as a candidate unified theory of cognition
- Problem solving and chunking mechanisms are tightly *intertwined*
- Soar comes with access to Java
 - adding user-written functions easier; debugging; rapid prototyping of GUIs
- Learning has some problems
 - overgeneralization* of chunks; *expensive* chunks
- Multiple goals not naturally supported
- Non-symbolic I/O not naturally supported

Soar 9 Extension



Course Summary

Cover two important levels of agents:

- Autonomous agents (micro level)
 - BDI model
 - Planning and practical reasoning
 - Learning
 - Cognitive modeling
- Multi-agent systems (macro level)
 - Modeling other agents (plan recognition)
 - Agent communication
 - Game theory / multi-agent learning
 - Agents work together (teamwork model, cooperation, coordination and negotiation)
 - Social modeling and reasoning

其它软件和工具

- NetLogo (Northwestern University)
<https://ccl.northwestern.edu/netlogo/>
- Repast (Argonne National Lab)
<https://repast.github.io/>
- Jade (Telecom Italia Lab)
<https://jade.tilab.com/>
- GAMA
<https://gama-platform.github.io/>
- Jason
<https://jason.sourceforge.net/wp/>
- Swarm (Santa Fe Institute)
http://www.swarm.org/wiki/Swarm_main_page

References

- R. A. Brooks. *A Robust Layered Control System for a Mobile Robot*. IEEE Journal of Robotics and Automation, RA-2(1):14-23, 1986
- I. A. Ferguson. *Integrated Control and Coordinated Behaviour: A Case for Agent Models*. In: M. Wooldridge and N. R. Jennings (Eds.), *Intelligent Agents: Theories, Architectures and Languages* (LNAI 890), pp.203-218. Springer, 1995
- J. P. Müller. *A Cooperation Model for Autonomous Agents*. In: J. P. Müller, M. Wooldridge and N. R. Jennings (Eds.), *Intelligent Agents III* (LNAI 1193), pp.245-260. Springer, 1997
- J. E. Laird, P. S. Rosenbloom and A. Newell. *Soar: An Architecture for General Intelligence*. In: P. S. Rosenbloom, J. E. Laird and A. Newell (Eds.), *The Soar Papers: Research on Integrated Intelligence*. The MIT Press, 1993
- J. E. Laird. *Extending the Soar Cognitive Architecture*. Proceedings of the First Conference on Artificial General Intelligence, 2008

End.