

规划问题求解 (AI Planning)

毛文吉

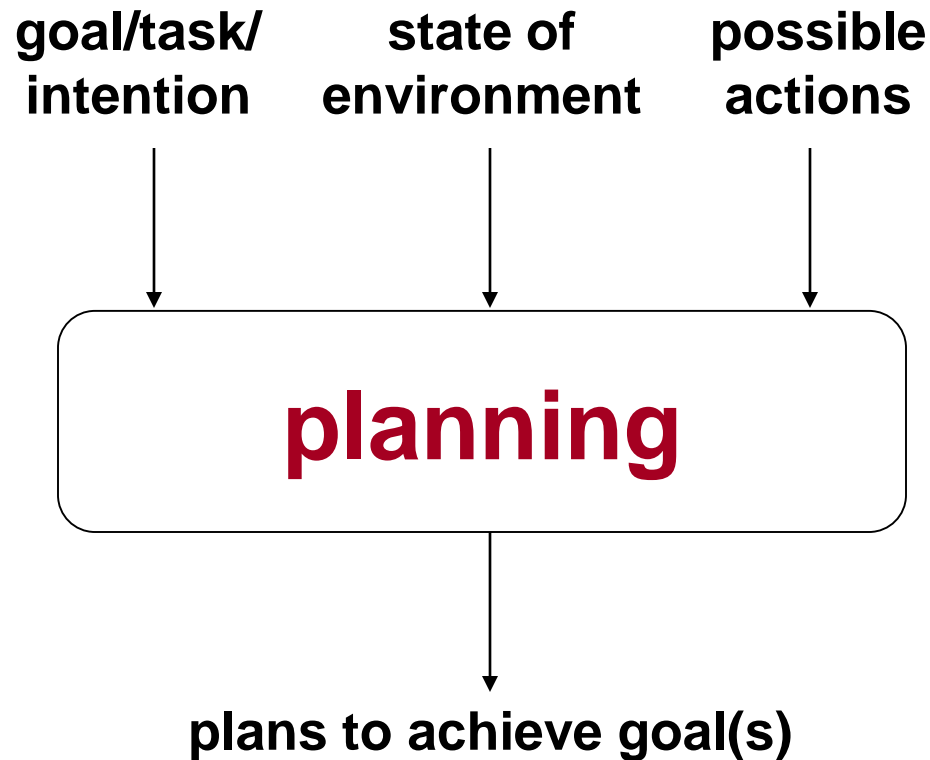
中国科学院自动化研究所

Planning agent

- Since the early 1970s [Fikes & Nilsson], the AI planning community has been closely concerned with the design of artificial agents
- AI planning system is a central component of any artificial agent
- Planning is essentially the automatic generation of a course of actions to achieve some desired goal
- Many planning algorithms have been proposed and planning has become a well-developed field in AI and agent research

What is planning

- An *automatic reasoning* process to generate plans of a sequence of actions for achieving certain goal(s)



Planning questions

Question 1: How do we *represent*. . .

- Goals to be achieved
- States of environment
- Actions available to agent
- Plans itself

Question 2: How do we use these representations to generate *plans*?

- What kind of reasoning should be involved?

Outline

- STRIPS-like plan representation
- Planning with state-space search
- Partial-order planning (POP)
- Planning graphs (GraphPlan)

Outline

- STRIPS-like plan representation
 - Planning with state-space search
 - Partial-order planning (POP)
 - Planning graphs (GraphPlan)

Strips (Fikes and Nilsson 71)

■ Highly influential representation for actions:

- Preconditions: list of propositions to be true (前提条件表)
- Delete list: list of propositions that will become false (删除表)
- Add list: list of propositions that will become true (增加表)

■ Example

Initial state: at(home), \neg have(beer), \neg have(chips)

Goal: have(beer), have(chips), at(home)

Actions:

Buy (X):

Precond: at(store)

Add: have(X)

Go (X, Y):

Precond: at(X)

Delete: at(X)

Add: at(Y)

Frame problem（框架问题）

- I go from home to the store, creating a new situation S' . In S' :
 - The store still sells chips
 - My age is still the same
 - Beijing is still the capital city of China...
- How can we efficiently represent everything that hasn't changed?

Ramification problem (分支问题)

- I go from home to the store, creating a new situation S' . In S' :
 - I am now in Zhongguancun
 - The number of people in the store went up by 1
 - The contents of my pockets are now in the store...
- Do we want to say all that in the action definition?

Strips treatment

In Strips, some facts are inferred within a world state

e.g. the number of people in the store

- Primitive facts, e.g. *at(home)* persist between states unless changed
- Inferred facts are not carried over and must be re-inferred
 - Avoids making mistakes, perhaps inefficient

Strips representation for actions:

Move-C-from-A-to-Table:


Precondition: $\text{on}(\text{C}, \text{A})$, $\text{clear}(\text{C})$

Effect:

add: $\text{on-table}(\text{C})$

$\text{clear}(\text{A})$

delete: $\text{on}(\text{C}, \text{A})$




- The explicit effects are the only changes to the state

Means-ends analysis（手段目的分析）

Strips's problem solving takes **means-ends analysis**

- Search by reducing the difference between state and goals
- What *means* (operators) are available to achieve the desired *ends* (goals)

The monkey and bananas problem（猴子香蕉问题）




Blocks world example (Sussman anomaly)


分成子任务. 把各自子任务完成.

- Strips uses *noninterleaved planner* (非交叉规划器), which cannot solve this example ...

Initial:



Goal:




Initial state: on(C, A), on-table(A), on-table(B),
clear(B), clear(C)

Exercise

- A noninterleaved planner is a planner that, when given two subgoals G1 and G2, produces either a plan for G1 concatenated with a plan for G2, or vice versa.

分成子目标 不一定完成总目标


Why a noninterleaved planner cannot solve this problem?



Outline

- STRIPS-like plan representation
- Planning with state-space search
- Partial-order planning (POP)
- Planning graphs (GraphPlan)

Search space: Blocks world




Search the space of world states

Planning as state space search


- Nodes: world states
 - Arcs: actions
 - Solution: path from the initial state to one state that satisfies the goal
-
- Progression（前向规划）: forward search
 - Regression（后向规划）: backward search

Planning algorithms 状态空间太大，效率低下

■ Progression: Forward state-space search



■ Regression: Backward state-space search



Properties of planning algorithms

■ Soundness 正确性

- A planning algorithm is **sound** if all solutions are legal plans
 - All preconditions, goals, and any additional constraints are satisfied

■ Completeness 完备

- A planning algorithm is **complete** if a solution can be found
- A planning algorithm is **strictly complete** if all solutions are included in the search space whenever one actually exists

■ Optimality 最优

- A planning algorithm is **optimal** if it maximizes a predefined measure of plan quality

Progression vs regression


- Both algorithms are
 - *sound* (they always return a valid plan)
 - *complete* (if a valid plan exists they will find one)
- Complexity $O(b^n)$ worst-case b 很大, 没有使用问题领域信息)
 - where $b = \text{branching factor}$,
n = number of “choose” operators
- Regression: often smaller b , focused by goals
- Progression: full state space to compute heuristics

Outline


- STRIPS-like plan representation
- Planning with state-space search
- Partial-order planning (POP)
- Planning graphs (GraphPlan)

Total-order vs Partial-order plans

Partial Order Plan:



Total Order Plans:



Search the space of plans

Partial-Order Planning (POP) 搜索规划空间.

Generates partial-order plans

- Nodes are partial plans
- Links are plan refinements
- Solution is a node (not a path)

Follow the least commitment principle [Weld 94]

- Don't commit to an order of actions until it is required

Partial plan representation



■ Plan = (A, O, L), where

- A: set of actions in the plan 动作
- O: *temporal orderings* between actions ($a < b$) 必要的序
- L: *causal links* linking actions via a literal 因果连接

■ Causal Link: $A_p \xrightarrow{Q} A_c$ (因果连接)

Action A_c (consumer) has precondition Q that is established in the plan by A_p (producer), e.g.

move-A-from-B-to-Table $\xrightarrow{\text{clear}(B)}$ *move-C-from-Table-to-B*



Threats to causal links and protection


Step A_t threatens link (A_p, Q, A_c) if:

- A_t has (not Q) as an effect, and
- A_t could come between A_p and A_c , i.e.

$O \cup (A_p < A_t < A_c)$ is consistent

(Ordering “ $<$ ” is not necessarily immediately before)

To protect causal link:



把 A_t 放在 A_p 之前或 A_c 之后. 保持因果连接 Q

Consistent plan in POP

A partial plan is consistent if:


- There are no cycles in the ordering constraints and 无环
- No conflicts with the causal links 无冲突

A consistent plan with no open preconditions is a **solution**:


- Every linearization (线性化) of a partial-order solution is a total-order solution whose execution from the initial state will reach a goal state

Total-order vs Partial-order plans

Partial Order Plan:



Total Order Plans: (线性化)



Initial plan

For uniformity, represent initial state and goal with two special actions:

- Start (A_0):
 - no preconditions, 无前提条件.
 - initial states as effects,
 - must be the first step in the plan
- Finish (A_∞):
 - no effects 无结果
 - goals as preconditions
 - must be the last step in the plan

Agenda: set of open conditions (议程集)

- e.g. $\{(on(A,B), A_\infty), (on(B,C), A_\infty)\}$

POP algorithm

POP((A, O, L), agenda, actions)


Initial plan: {*Start*, *Finish*} and preconditions in *Finish* as open conditions

1. If **agenda** is empty, then **return** (A, O, L) 结束条件
2. Pick (Q, A_{need}) from **agenda** (子) 目标
3. **Choose** an action A_{add} that adds effect Q 动作选择
 - If no such action exists, **fail**
 - Add the link $A_{\text{add}} \xrightarrow{Q} A_{\text{need}}$ to **L** and the ordering $A_{\text{add}} < A_{\text{need}}$ to **O**
 - If A_{add} is new, add it to **A** 规划扩充
4. Remove (Q, A_{need}) from **agenda**. If A_{add} is new, for each of its preconditions P add (P, A_{add}) to **agenda** 更新(子) 目标
5. For every action A_t in A that threatens any causal link $A_p \rightarrow A_c$ in **L**
 - **Choose** to add $A_t < A_p$ or $A_c < A_t$ to **O**
 - If neither choice is consistent, **fail**
6. POP((A, O, L), **agenda**, actions)

保护因果连接:

- 降级(Demotion): $A_t < A_p$
- 升级(Promotion): $A_c < A_t$


POP example: Sussman anomaly



on-table(C) clear(A) on(A,B) on(B,C)

Finish (\mathbf{A}_∞)

Work on open condition on(B,C)

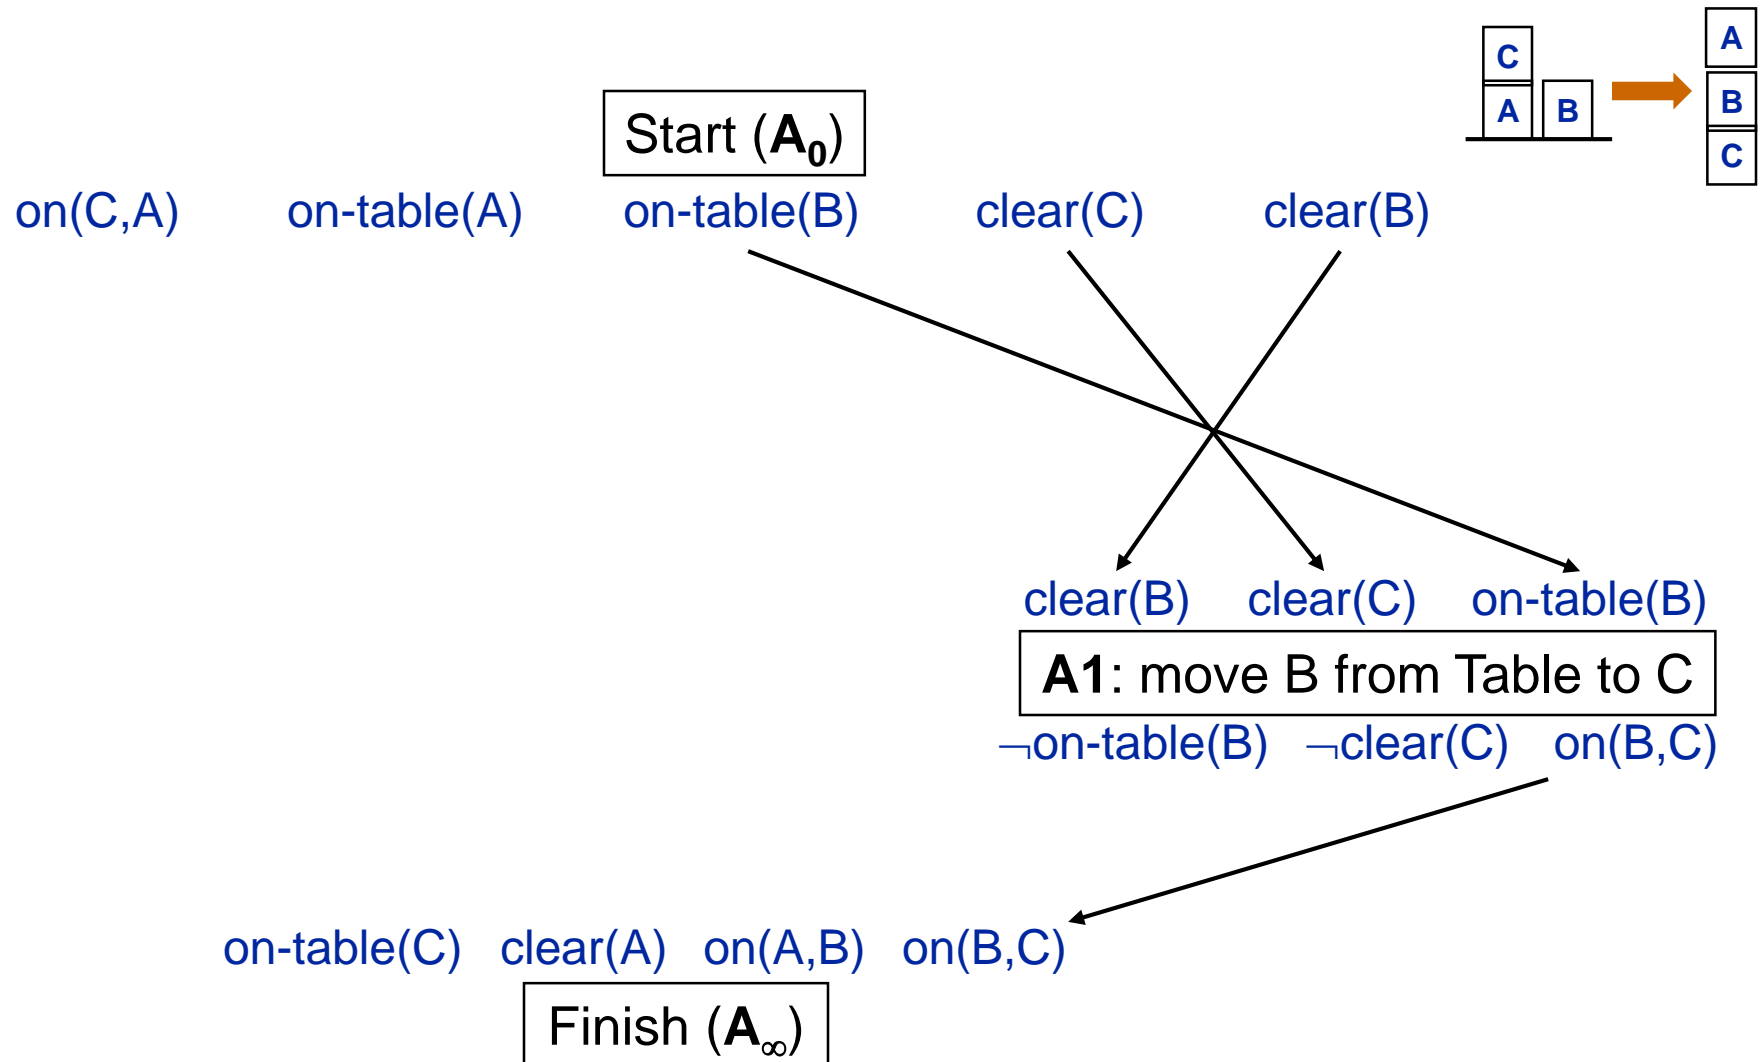


clear(B) clear(C) on-table(B)
A1: move B from Table to C
 \neg on-table(B) \neg clear(C) on(B,C)

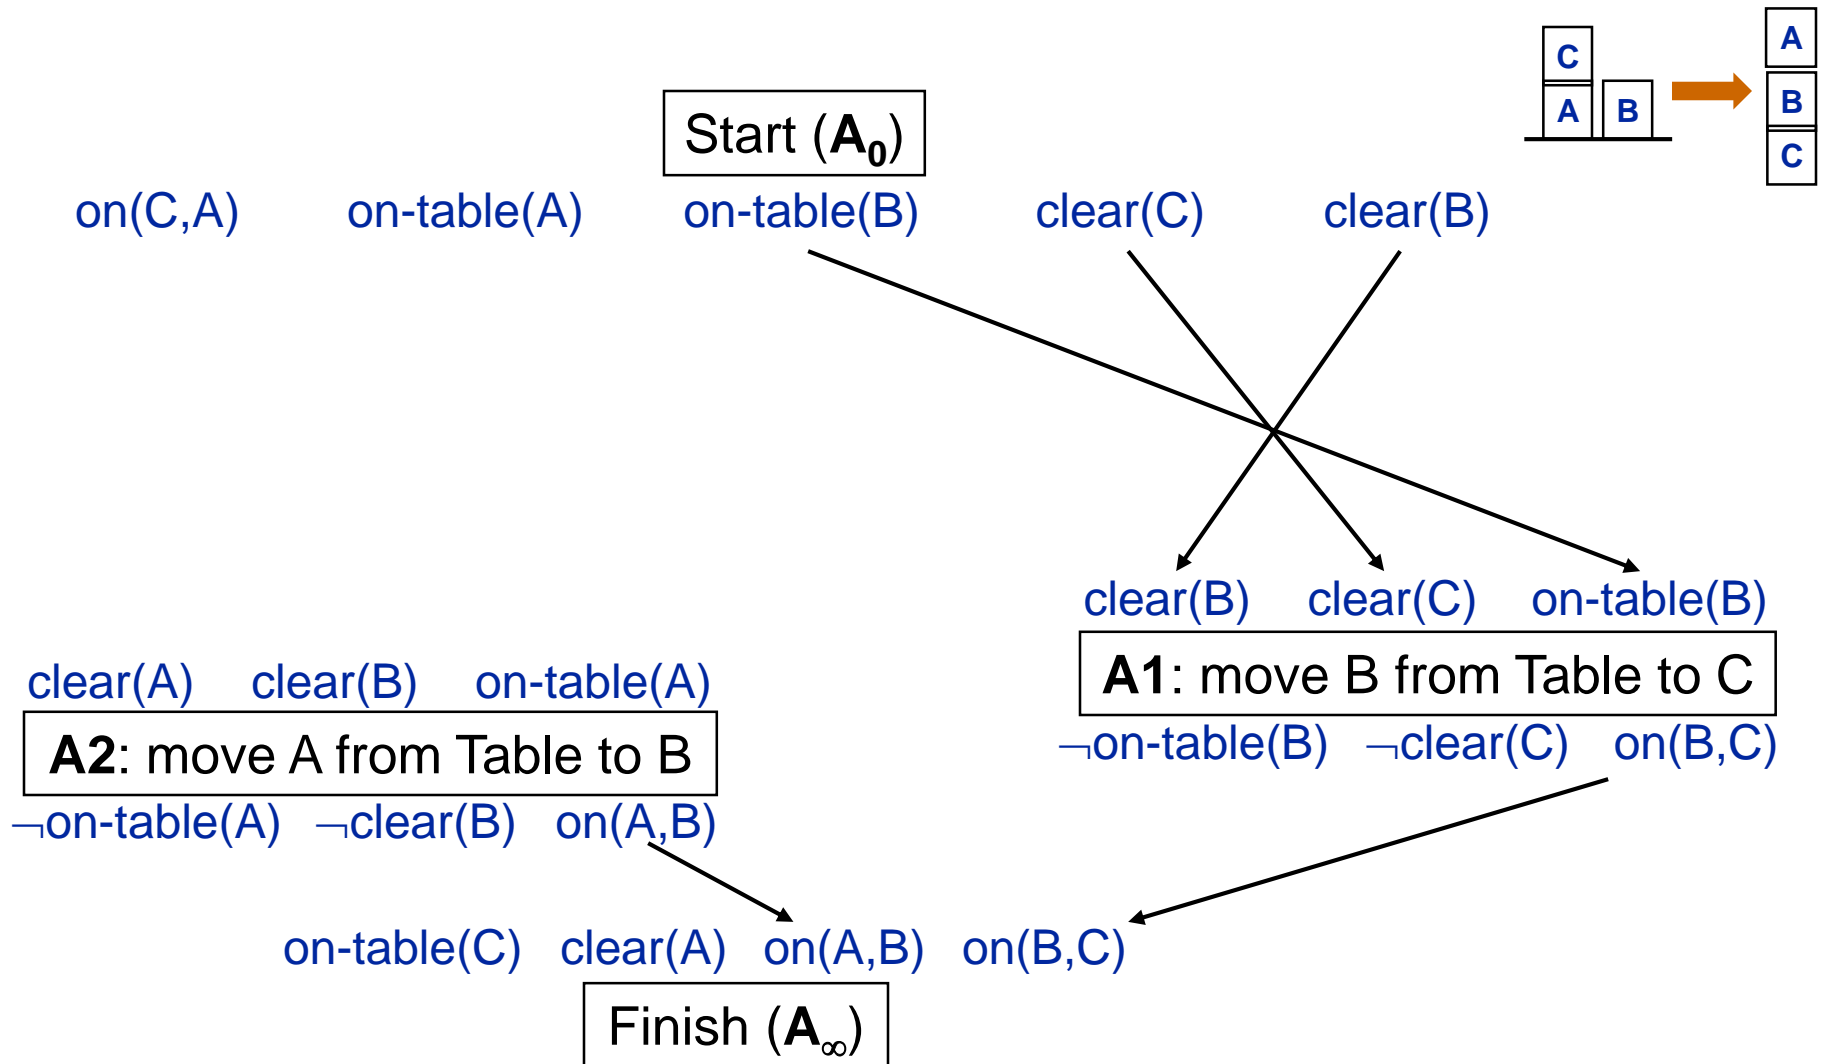
on-table(C) clear(A) on(A,B) on(B,C)

Finish (A_∞)

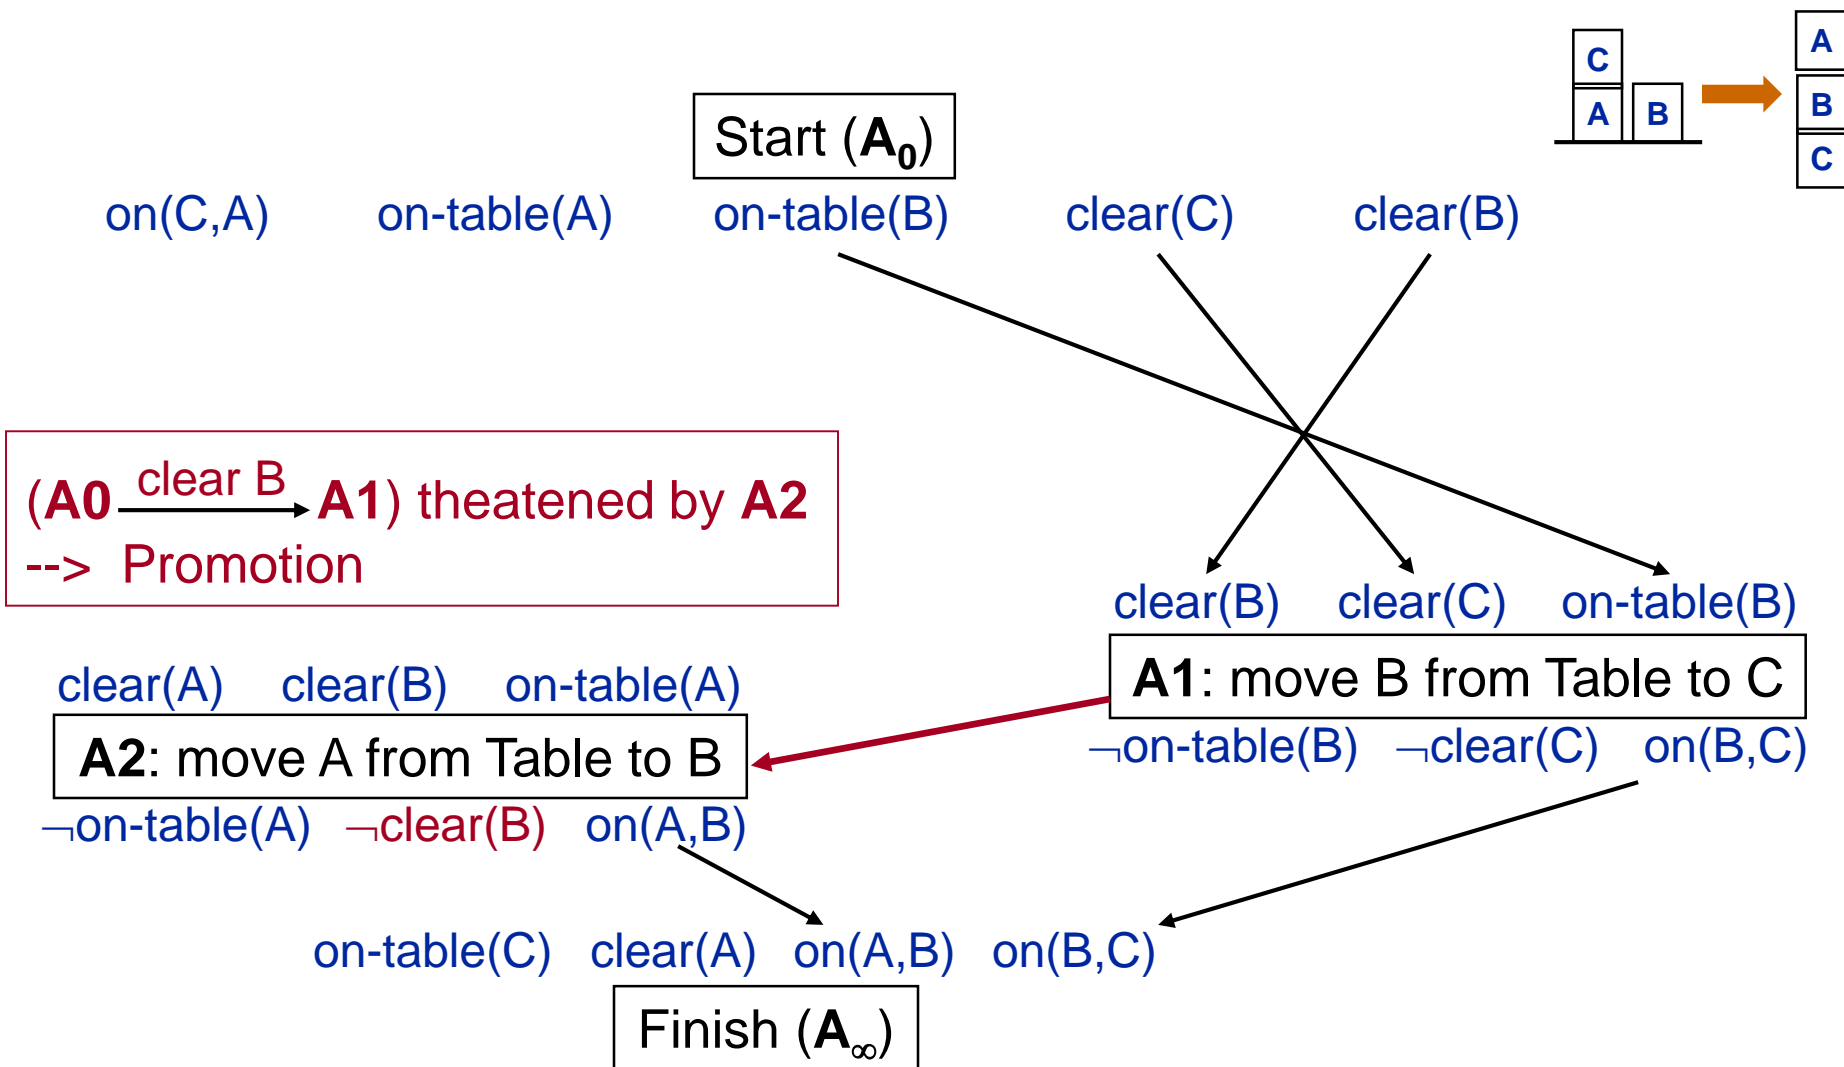
Work on open conditions clear(B), clear(C)...



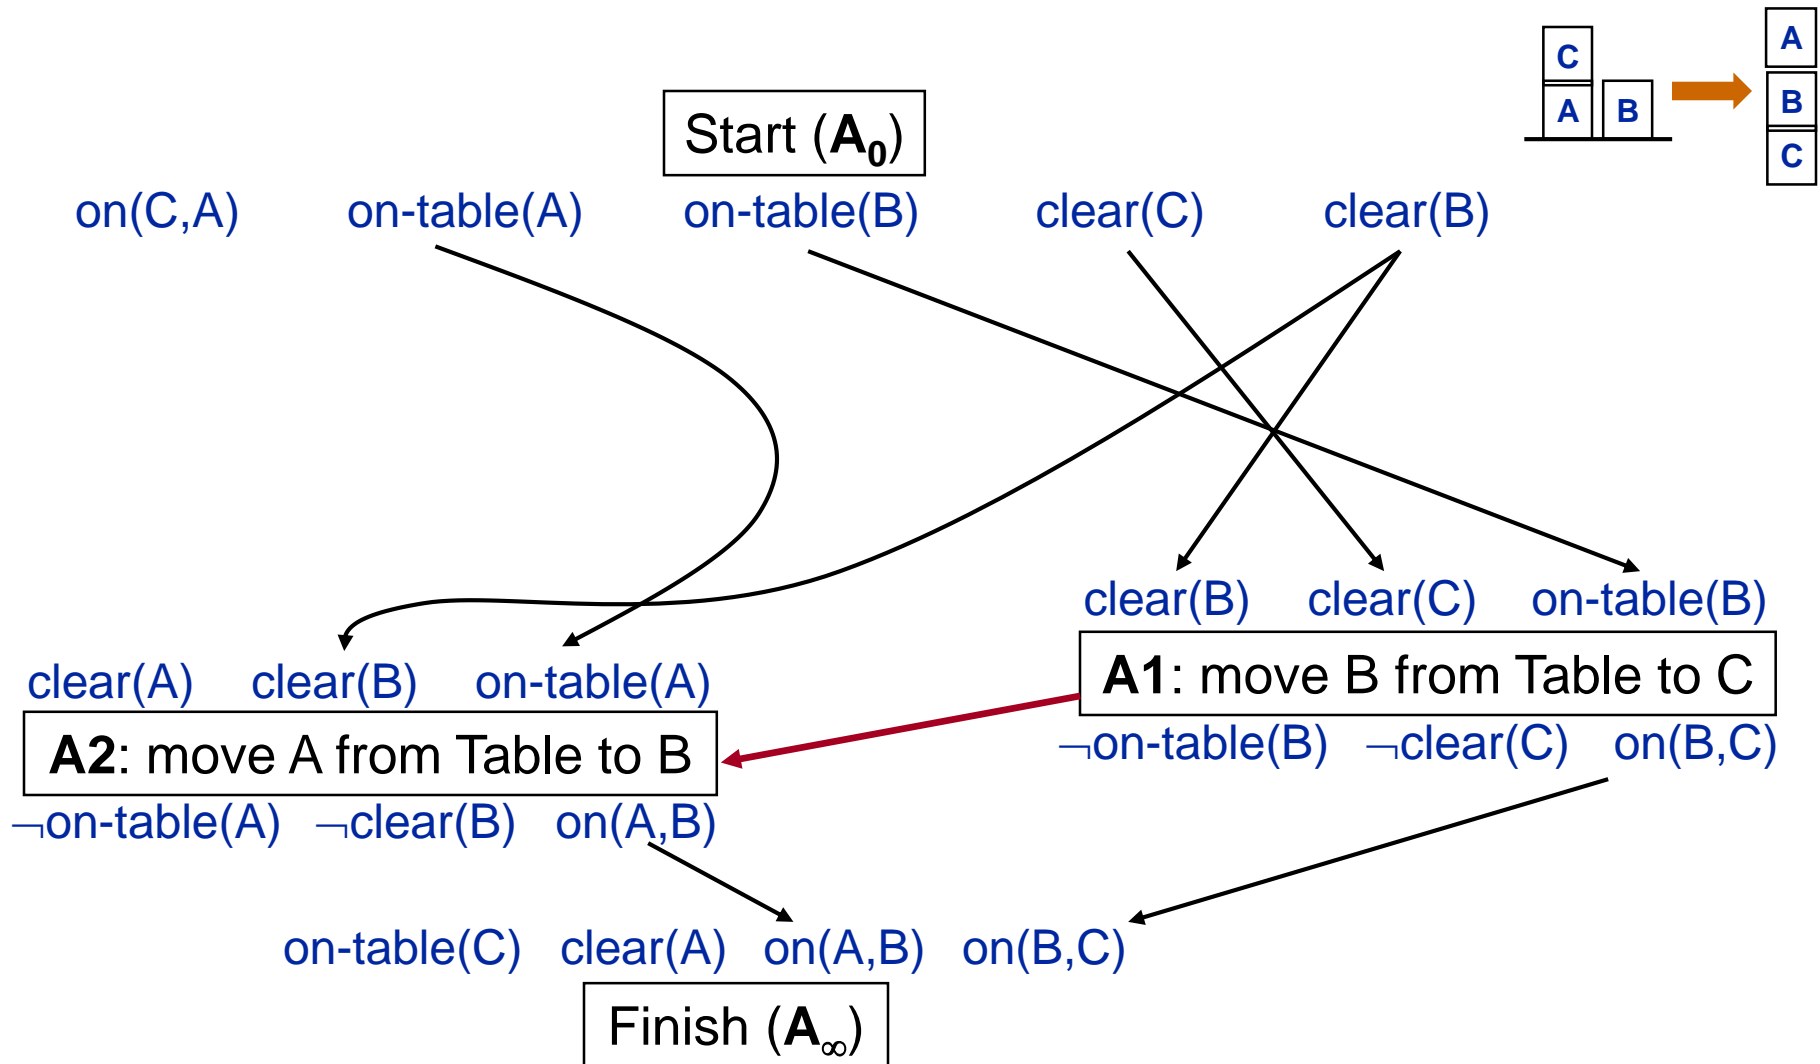
Work on open condition $\text{on}(A,B)$



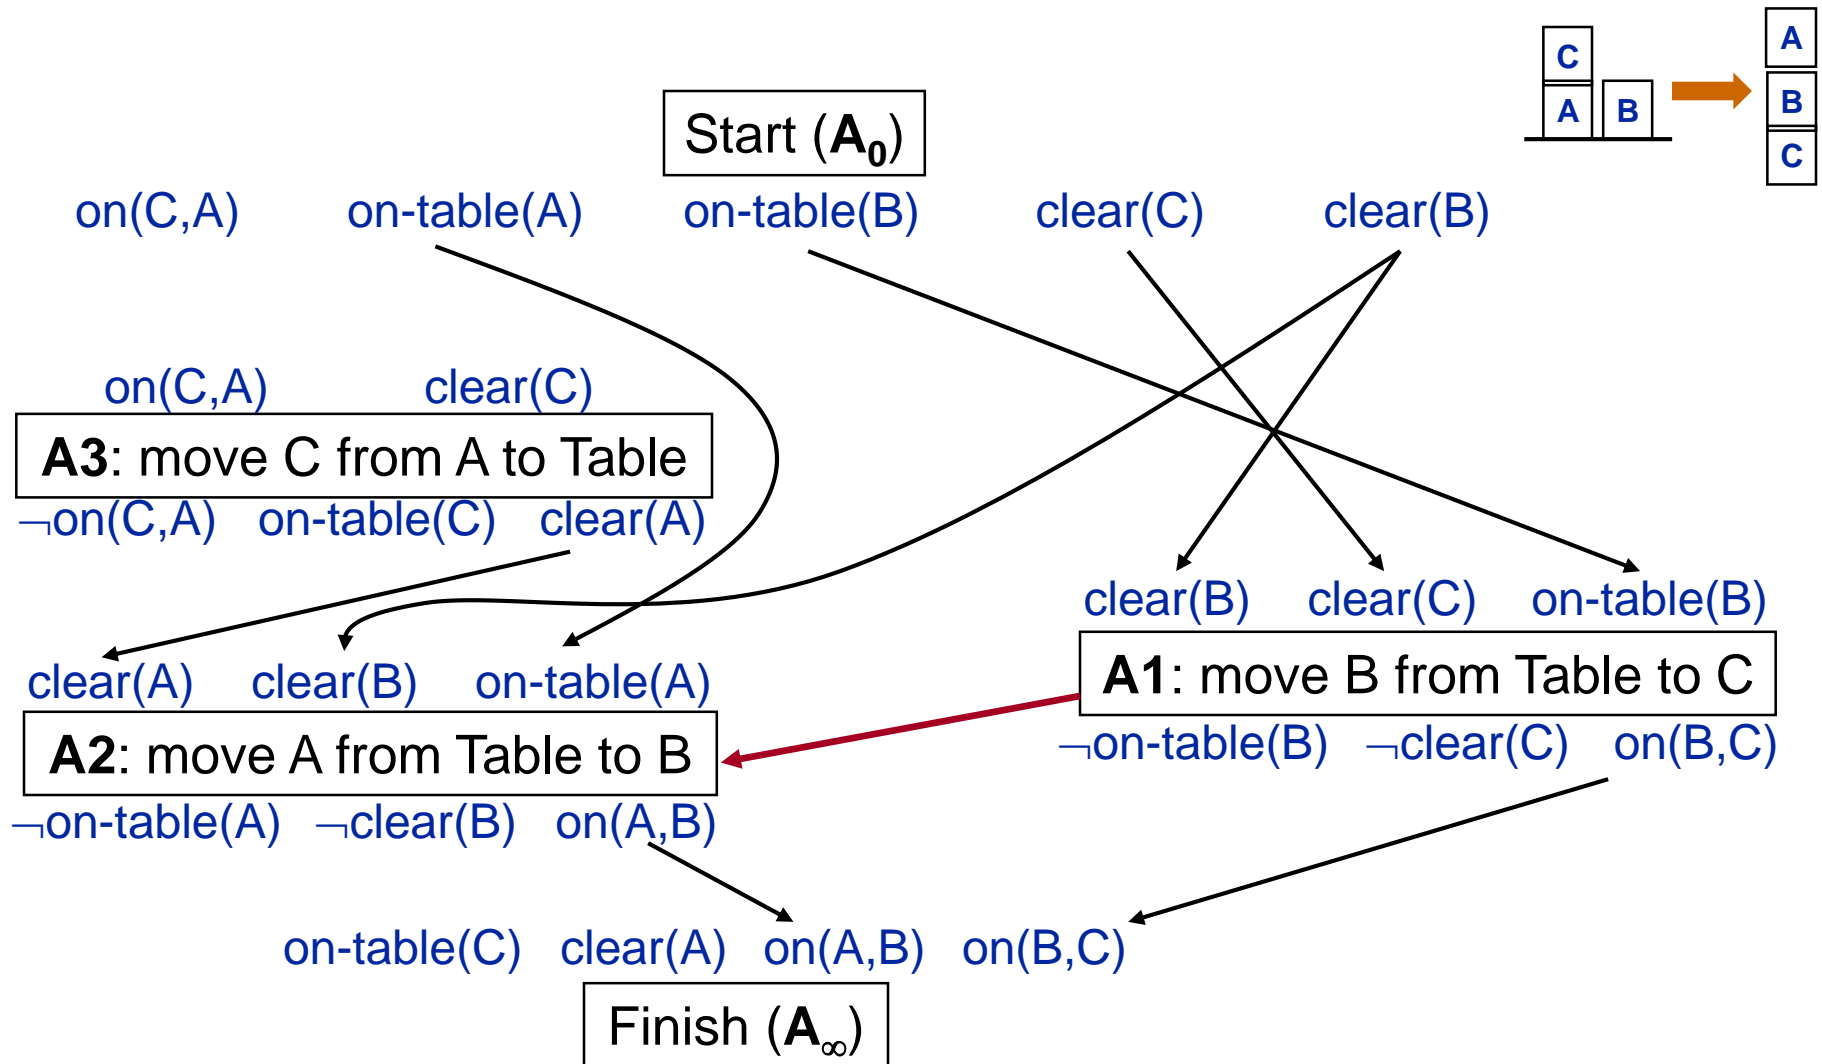
Protect causal link



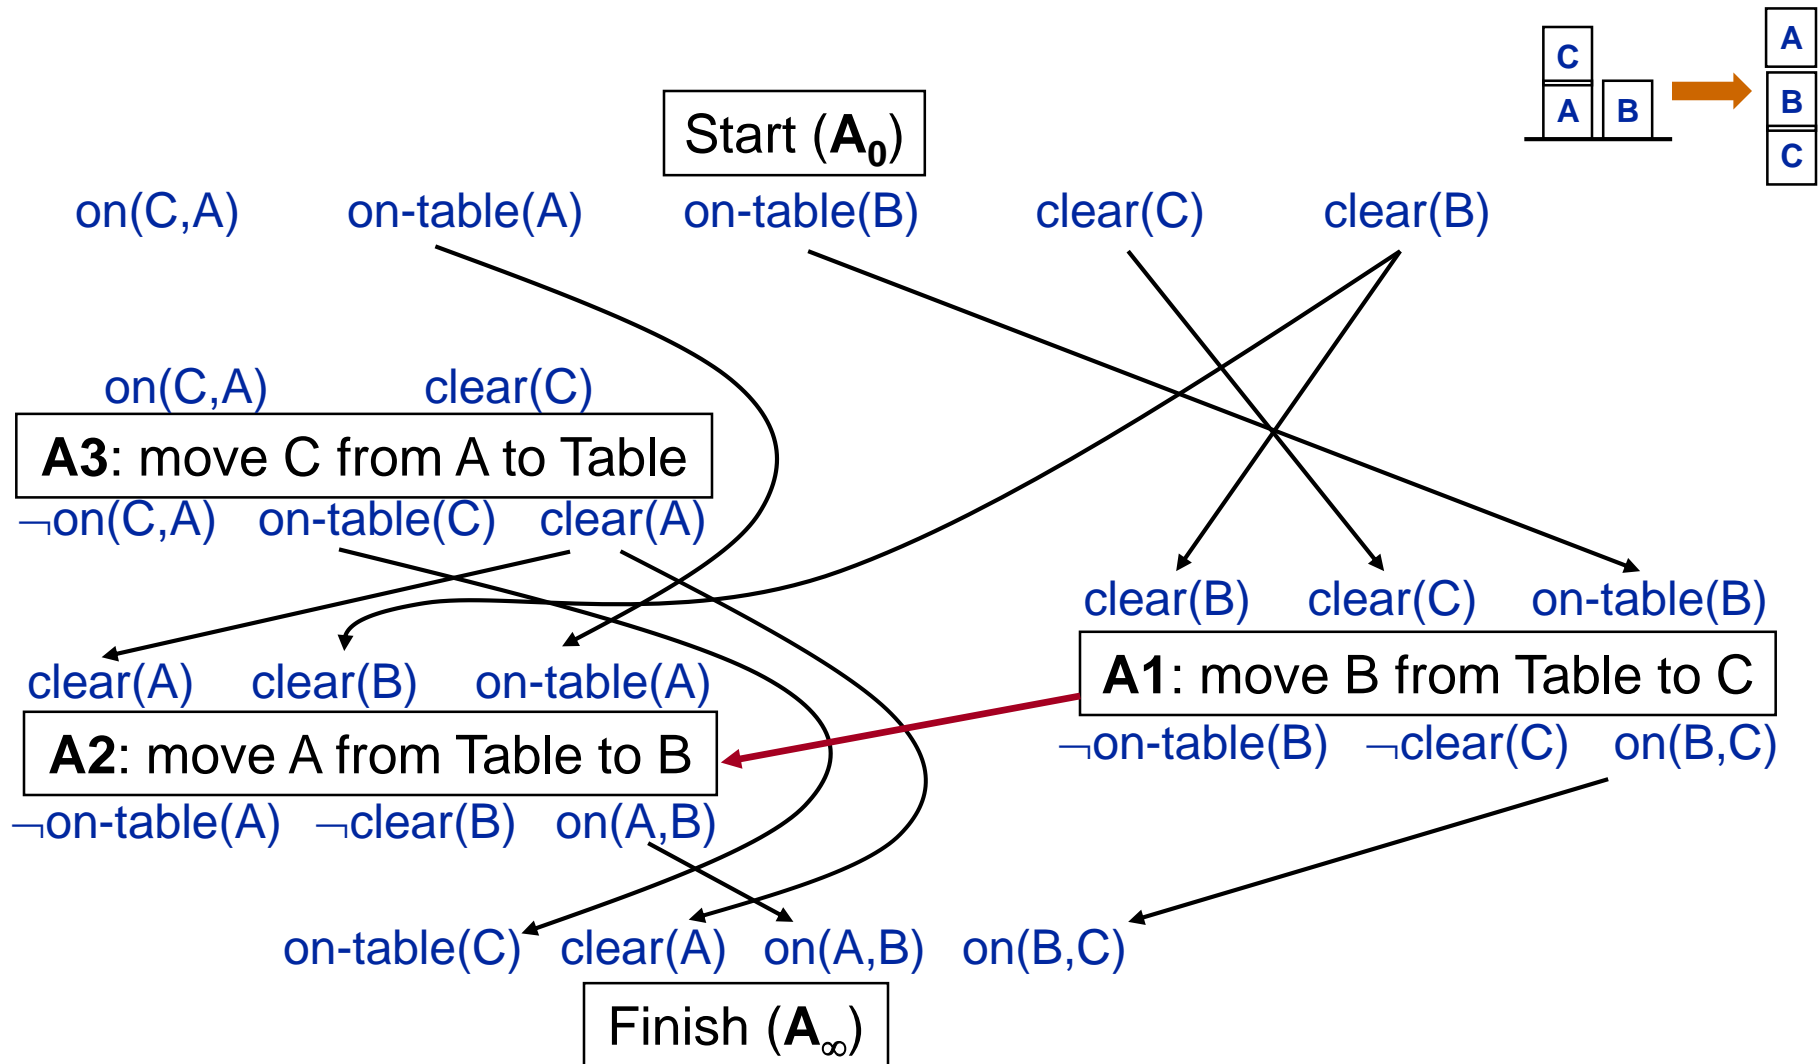
Work on conditions clear(B), on-table(A)



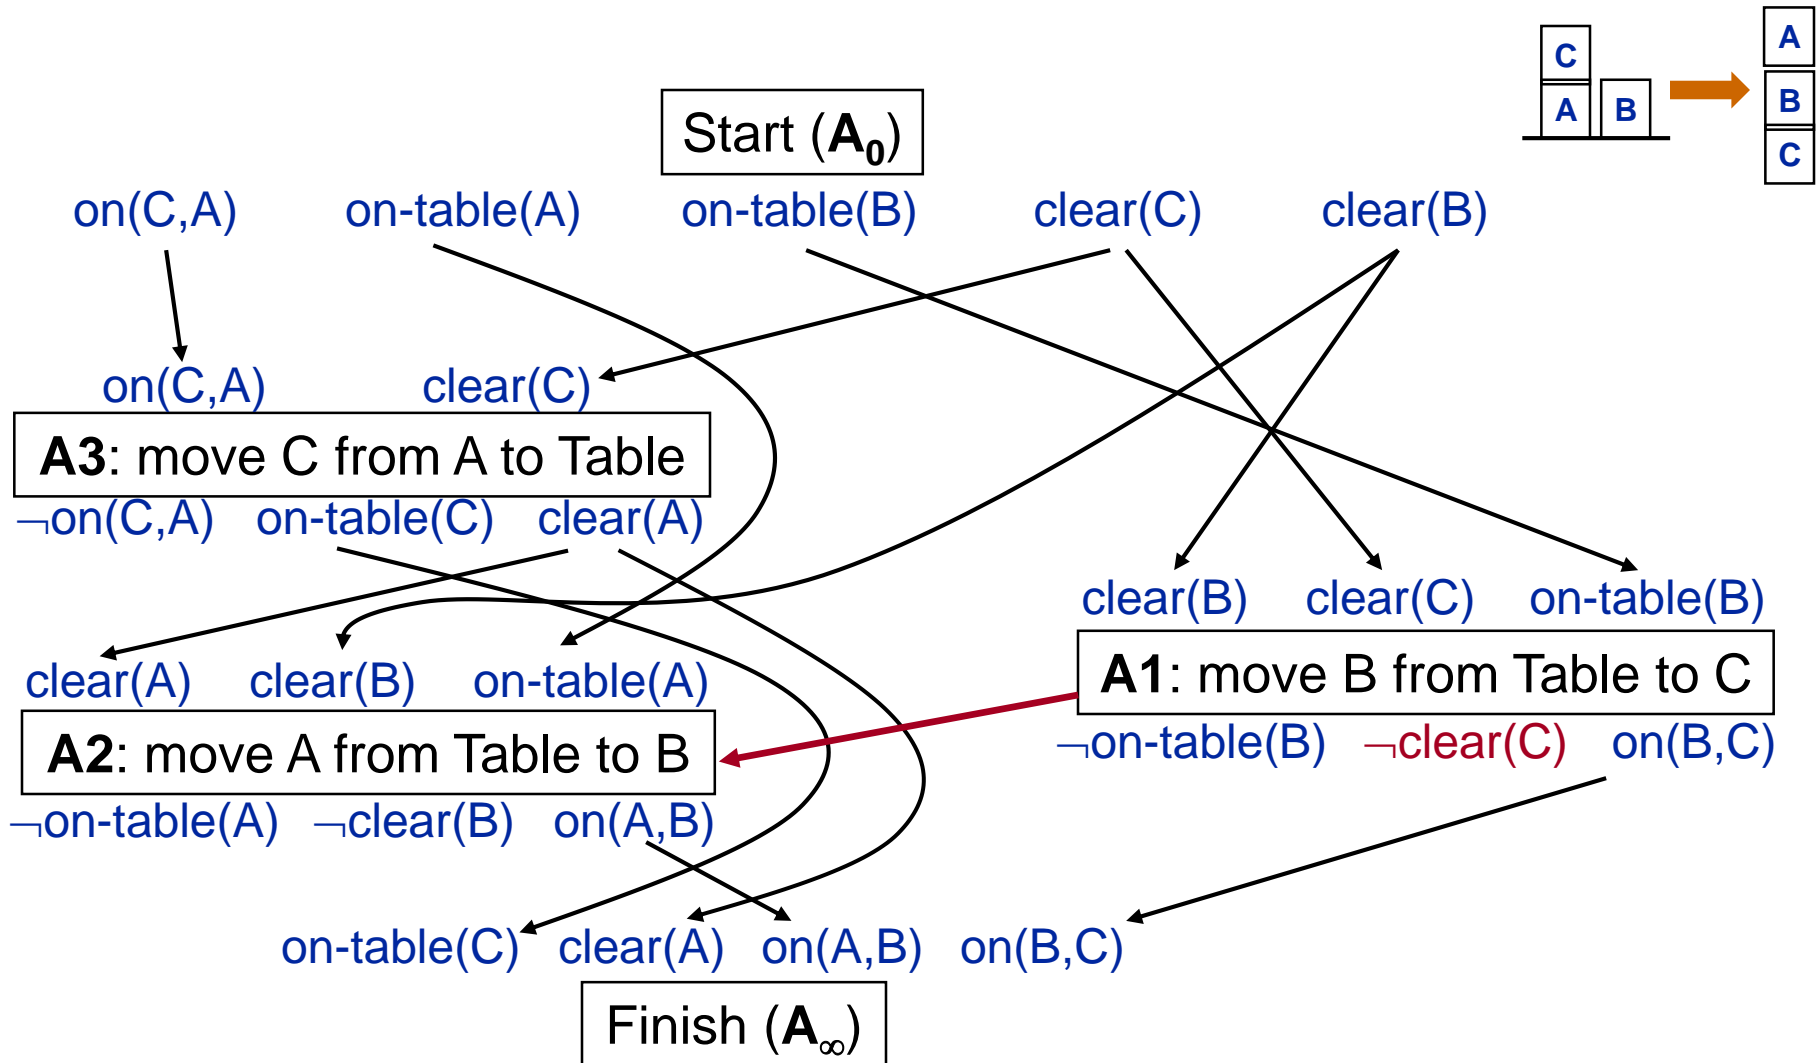
Work on condition clear(A)



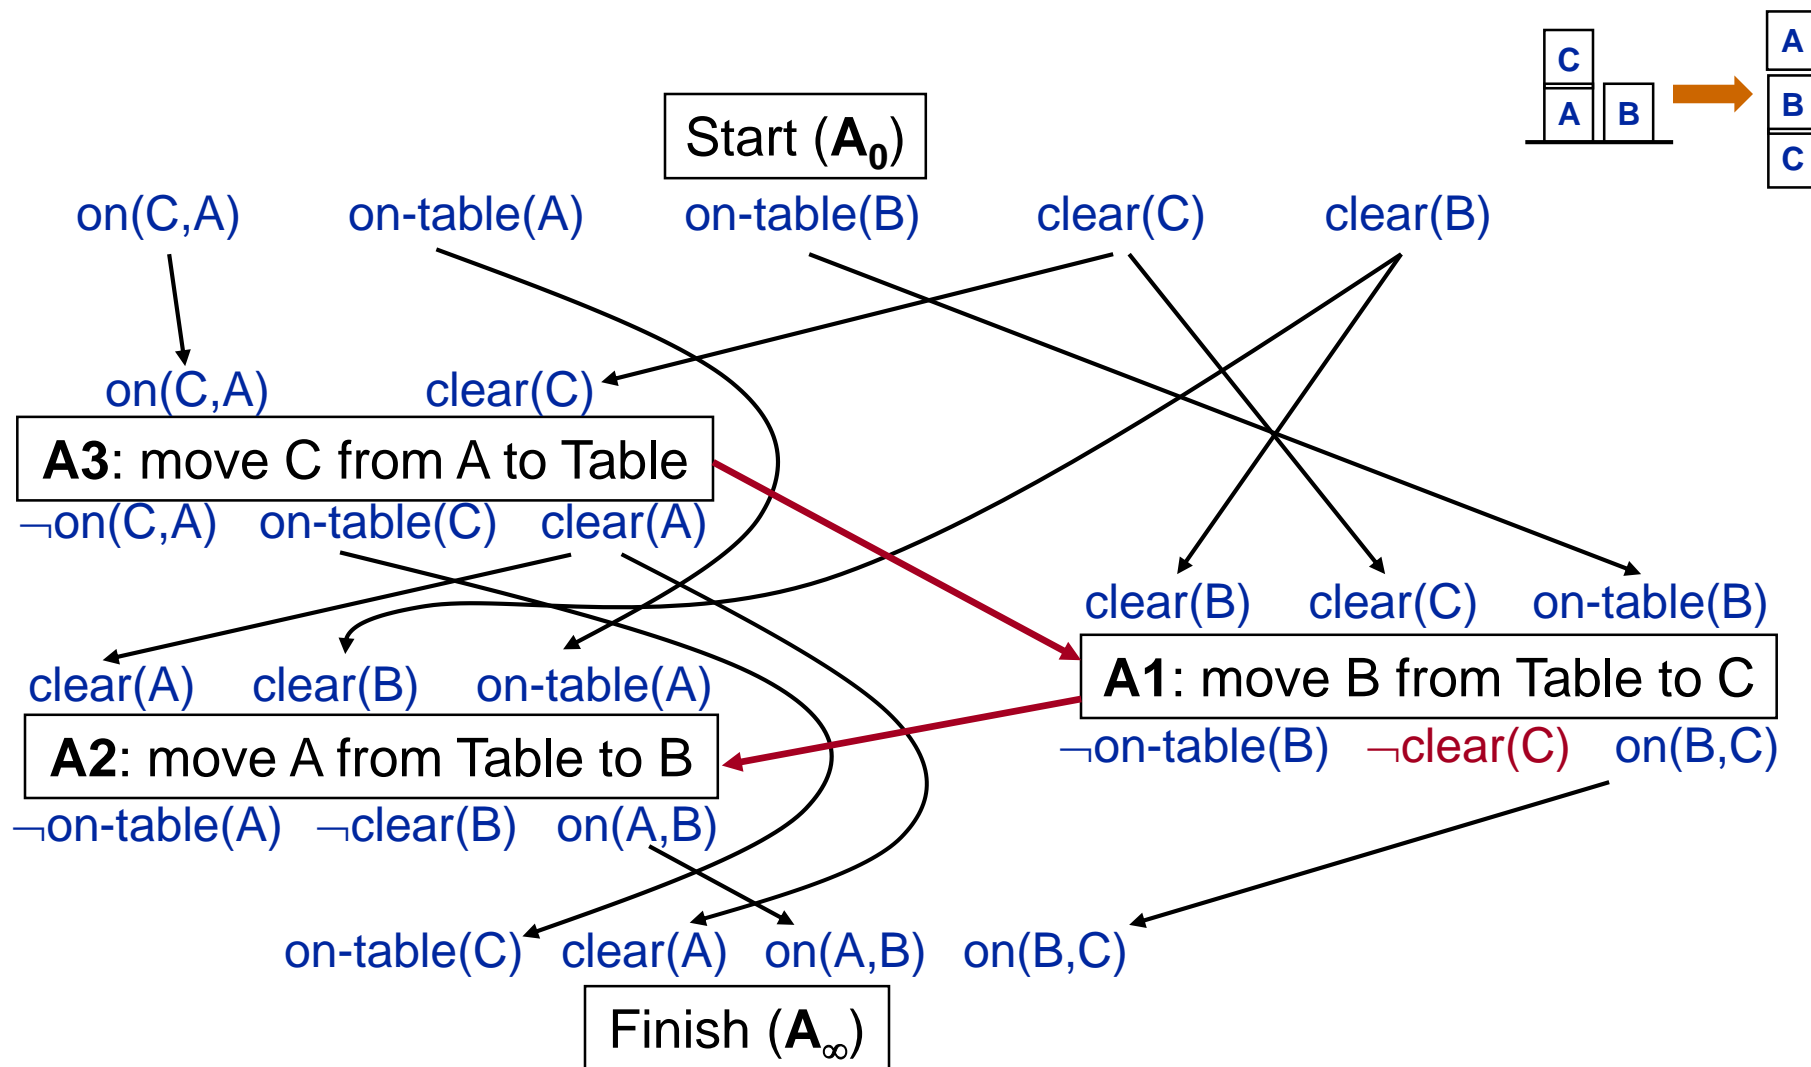
Work on conditions clear(A), on-table(C)



Work on conditions on(C,A), clear(C)



Protect causal link



Final plan step

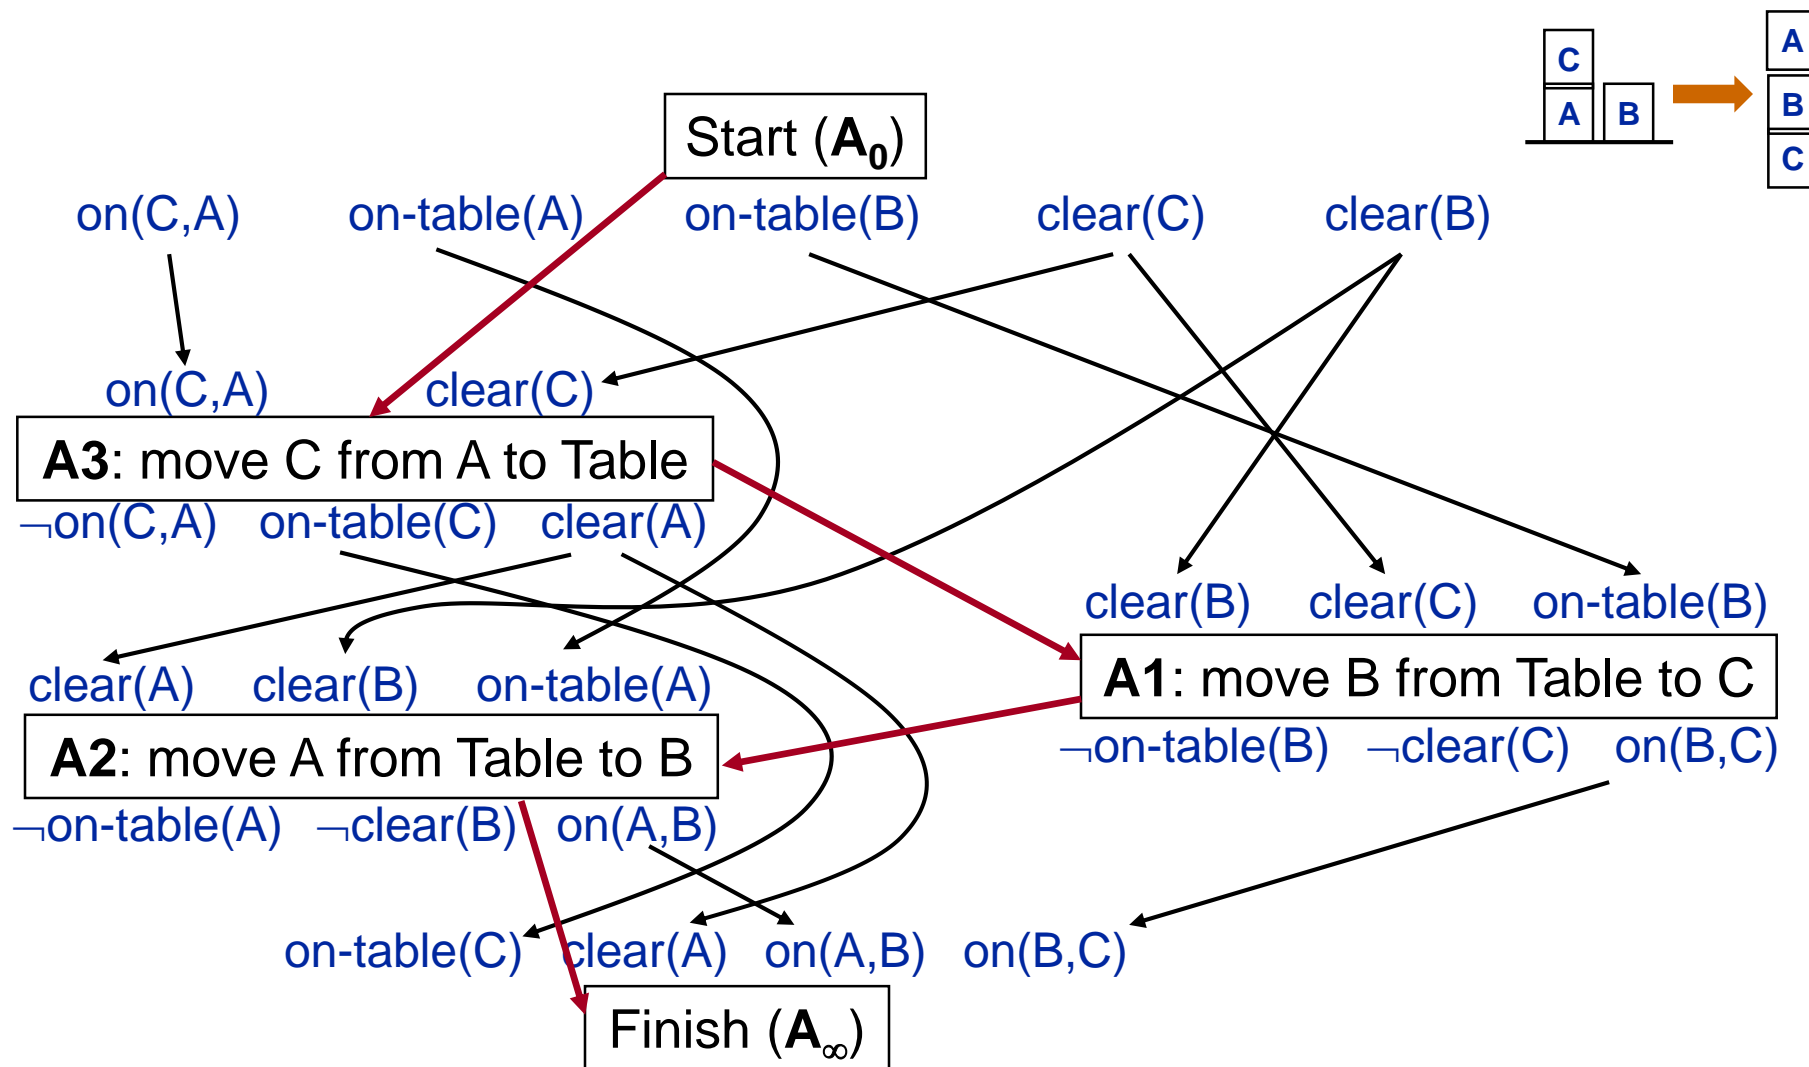
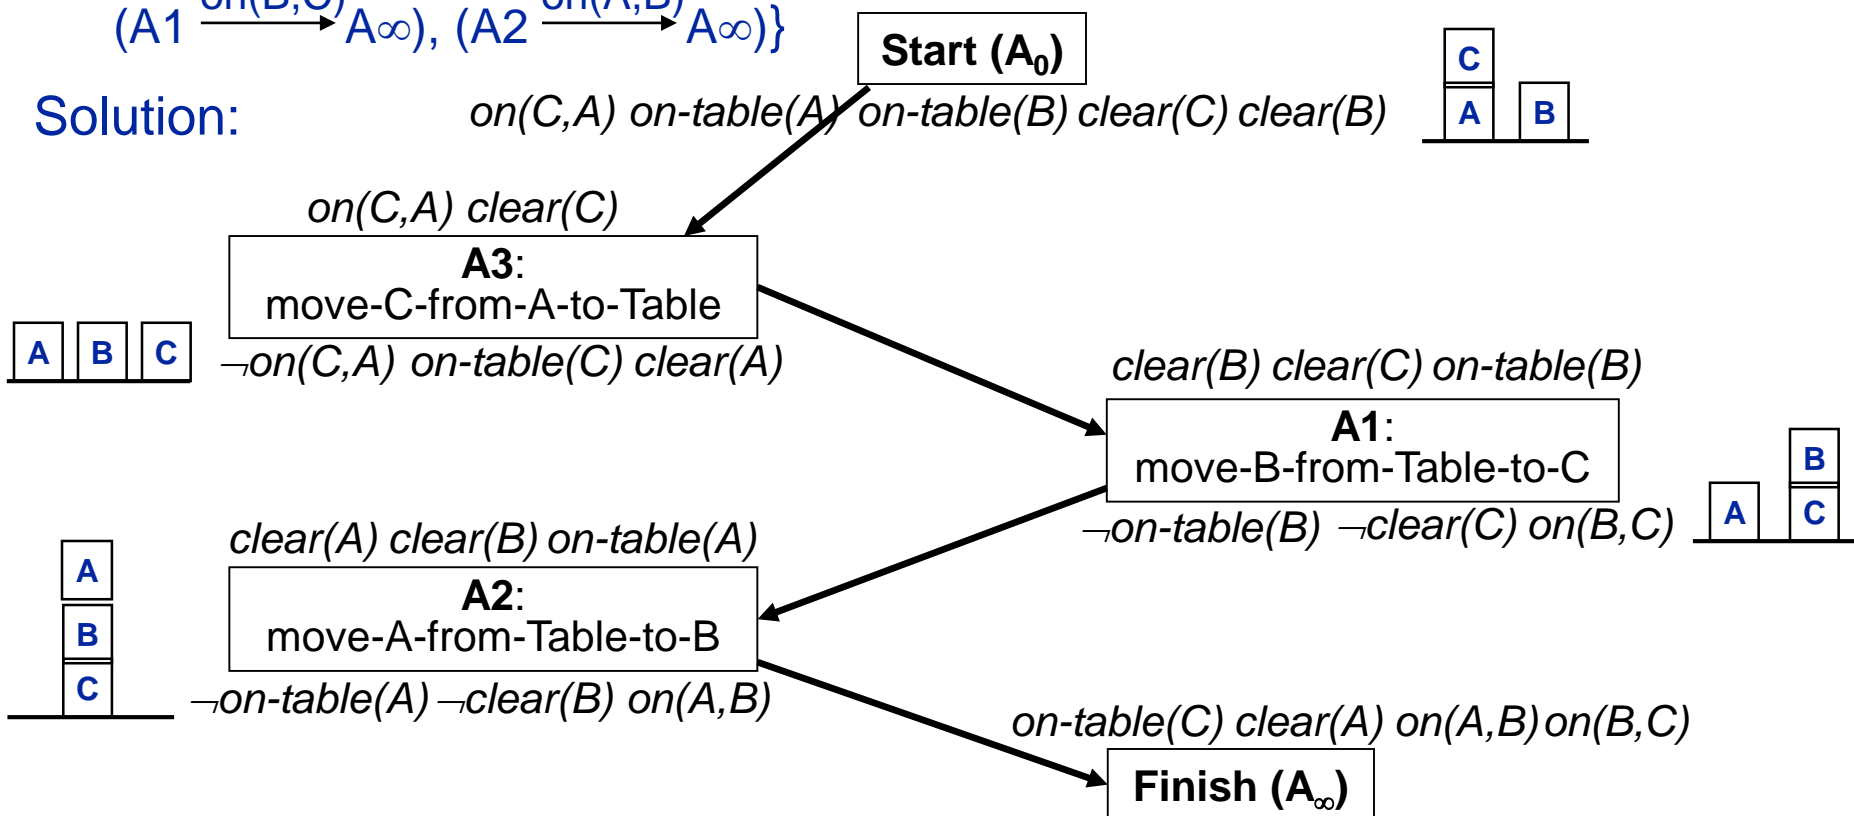


Illustration of solution plan

A: {A1, A2, A3, A0, A_∞}; **O:** {A3 < A1, A1 < A2}

L: {(A0 $\xrightarrow{\text{on}(C,A)}$ A3), (A0 $\xrightarrow{\text{clear}(C)}$ A3), (A0 $\xrightarrow{\text{clear}(B)}$ A1), (A0 $\xrightarrow{\text{clear}(C)}$ A1), (A0 $\xrightarrow{\text{on-tbl}(B)}$ A1),
 (A0 $\xrightarrow{\text{clear}(B)}$ A2), (A0 $\xrightarrow{\text{on-tbl}(A)}$ A2), (A3 $\xrightarrow{\text{clear}(A)}$ A2), (A3 $\xrightarrow{\text{on-tbl}(C)}$ A_∞), (A3 $\xrightarrow{\text{clear}(A)}$ A_∞),
 (A1 $\xrightarrow{\text{on}(B,C)}$ A_∞), (A2 $\xrightarrow{\text{on}(A,B)}$ A_∞)}

Solution:



POP exercise

- For this problem of putting one's shoes and socks, apply POP to this problem. Show the final plan, and corresponding A, O and L for this problem.

(the ordering constraints that put every other action after *Start* and before *Finish* can be omitted in O)

- The planning problem is described as follows:

Initial state: *Empty*

Goal: *RightShoeOn, LeftShoeOn*

Action *RightShoe*

Precondition: *RightSockOn*

Effect: *RightShoeOn*

Action *RightSock*

Effect: *RightSockOn*

Action *LeftShoe*

Precondition: *LeftSockOn*

Effect: *LeftShoeOn*

Action *LeftSock*

Effect: *LeftSockOn*

Properties of POP

POP is *sound* and *complete*

- Can be much faster than state-space planning, because of no need to backtrack over goal orderings (so less branching is required)
- Although it is more expensive per node and makes more choices than Regression, reduction in branching size **often** gains more
 - Larger n but smaller b

Flexibility gained by partial order

- Can be very useful to agent when the world fails to cooperate
- Make it easier to combine smaller plans into larger ones
 - Each plan can reorder its actions to avoid conflict with other plans

Partial-order (POP) vs State-space planning

Complexity: $O(b^n)$ worst-case

■ Non-deterministic choices (n):

- Progression, Regression: $n = |\text{actions}|$
- POP: $n = |\text{preconditions}| + |\text{link protection}|$ n 更大
- Generally an action has several preconditions

■ Branching factor (b)

POP has smaller b:

- No backtrack due to goal ordering
- Least commitment: no premature step ordering

Comparison

	State Space	Plan Space
Algorithm	Progression Regression	POP
Nodes	World States	Partial Plans
Edges/ Transitions	Actions E.g. in blocks world: <ul style="list-style-type: none">▪ move-A-from-B-to-C▪ move-B-from-A-to-Table▪ move-C-from-B-to-A▪ ...	Plan refinements: <ul style="list-style-type: none">▪ Step addition▪ Step reuse▪ Demotion▪ Promotion

More expressive action representation

UCPOP [Penberthy and Weld 92]

- Actions with variables

Actions: Move-C-from-A-to-Table  Move ?b from ?x to ?y
Move-A-from-B-to-C

- Conditional effects (条件结果)

Effects: (and (on ?b ?y) (not (on ?b ?x)) (clear ?x)
(when (= ?y Table) (clear ?y)))

- Disjunctive (析取) preconditions

Preconditions: (and (on ?b ?x)
(or (clear ?y) (big-and-flat ?y)))

- Universal quantification (全称量词)

Outline


- STRIPS-like plan representation
- Planning with state-space search
- Partial-order planning (POP)
- Planning graphs (GraphPlan)

History and motivation

- Before GraphPlan, mostly work on PSP-like planners:
 - POP, SNLP, UCPOP, etc
- Because GraphPlan was so much faster, many subsequent planning algorithms used ideas from it
 - IPP, STAN, GraphHTN, SGP, Blackbox, Medic, TGP, LPG
 - Many of them are much faster than the original Graphplan

Regression (backward search)

may try lots of actions
that can't be reached
from the initial state



Main ideas

- A big source of inefficiency in search algorithms is the *branching factor*
 - the number of children of each node
- GraphPlan reduces the branching factor:
 1. Create a *relaxed problem*
 - Remove some restrictions of the original problem
 - Want the relaxed problem to be easy to solve (polynomial time)
 - The solutions to the relaxed problem will include all solutions to the original problem
 2. Do a modified search for the original problem
 - Restrict its search space to include only those actions that occur in solutions to the relaxed problem

GraphPlan

GraphPlan [Blum and Furst 97]

- Preprocessing before engaging in search
- Construct a *planning graph* to record constraints on possible plans
- Forward search combined with backward search, incl. two stages:
 - Extend: extend the planning graph at each time step
 - Search: Use the planning graph to constrain search for a solution plan
- Graphplan either finds a valid plan or concludes there is no solution exists

GraphPlan algorithm

Procedure GraphPlan:

- For $k = 0, 1, 2, \dots$

<Graph expansion>: 规划图扩展

- Construct a planning graph that contains k levels
- Check whether the planning graph satisfies a necessary (but insufficient) condition for plan existence
- If it does, then


relaxed
problem

<Solution extraction>: 规划解提取

- Backward search, modified to consider only the actions in the planning graph
- If a solution is found, then return *solution* 结束条件

A planning graph (规划图)

- Consist of a sequence of **levels** that correspond to time steps in the plan, where level 0 is the initial state
- Each level contains a set of propositional literals and a set of actions, with edges connecting action preconditions and effects



Expanding a planning graph - Actions

- To expand an action-level i :
 - Add each instantiated action, for which all of its preconditions are present at proposition-level i AND *no two of its preconditions are exclusive*
 - Add all the no-op actions 空操作
- Determine the mutual exclusive (mutex, 互斥) actions

Expanding a planning graph – Propositions

- To expand a proposition-level $i+1$:
 - Add all the effects of the inserted actions at action-level i :
distinguishing *add and delete effects*
- Determine the mutual exclusive (mutex, 互斥) propositions

Determining mutex relations (互斥关系)

Two actions A and B are *mutex* at an action-level, if:

- **Interference:** A (or B) disables a *precondition* or an *effect* of B (or A)
前提和结果冲突
- **Competing needs:** A and B have *mutex preconditions*
前提条件矛盾.

Two propositions p and q are *mutex* at a proposition-level if



- **Negation:** p (or q) is the *negation* of q (or p)
互为否定
- **Inconsistent support:** *all* ways of achieving them are *mutex* (i.e. all the actions that add p are mutual exclusion of all the actions that add q)

Example: mutex actions

Two actions *A* and *B* are *mutex* at an action-level, if:

- **Interference:** *A* (or *B*) disables a *precondition* or an *effect* of *B* (or *A*)
- **Competing needs:** *A* and *B* have *mutex preconditions*

Interference (inconsistent effects): **Interference** (precondition-effect):




Example: mutex actions

Two actions *A* and *B* are *mutex* at an action-level, if:

- **Interference:** *A* (or *B*) disables a *precondition* or an *effect* of *B* (or *A*)
- **Competing needs:** *A* and *B* have *mutex preconditions*

Competing needs (mutex preconditions):




Example: mutex propositions

Two propositions p and q are *mutex* at a proposition-level if

- **Negation:** p (or q) is the *negation* of q (or p)
- **Inconsistent support:** *all* ways of achieving them are *mutex* (i.e. all the actions that add p are mutual exclusion of all the actions that add q)

Negation:




Example: mutex propositions

Two propositions p and q are *mutex* at a proposition-level if

- **Negation:** p (or q) is the *negation* of q (or p)
- **Inconsistent support:** *all* ways of achieving them are *mutex* (i.e. all the actions that add p are mutual exclusion of all the actions that add q)

Inconsistent support:



Dinner date example（为恋人准备惊喜晚餐）


- Initial state: (有)垃圾, 手清洁, 安静
- Goal: 晚餐, 惊喜, \neg 垃圾
- Actions:
 - 烹饪: precondition (手清洁)
effect (晚餐)
 - 打包: precondition (安静)
effect (惊喜)
 - 手工清理: precondition (垃圾)
effect (\neg 垃圾, \neg 手清洁)
 - 机器清理: precondition (垃圾)
effect ((\neg 垃圾, \neg 安静))

Dinner date example

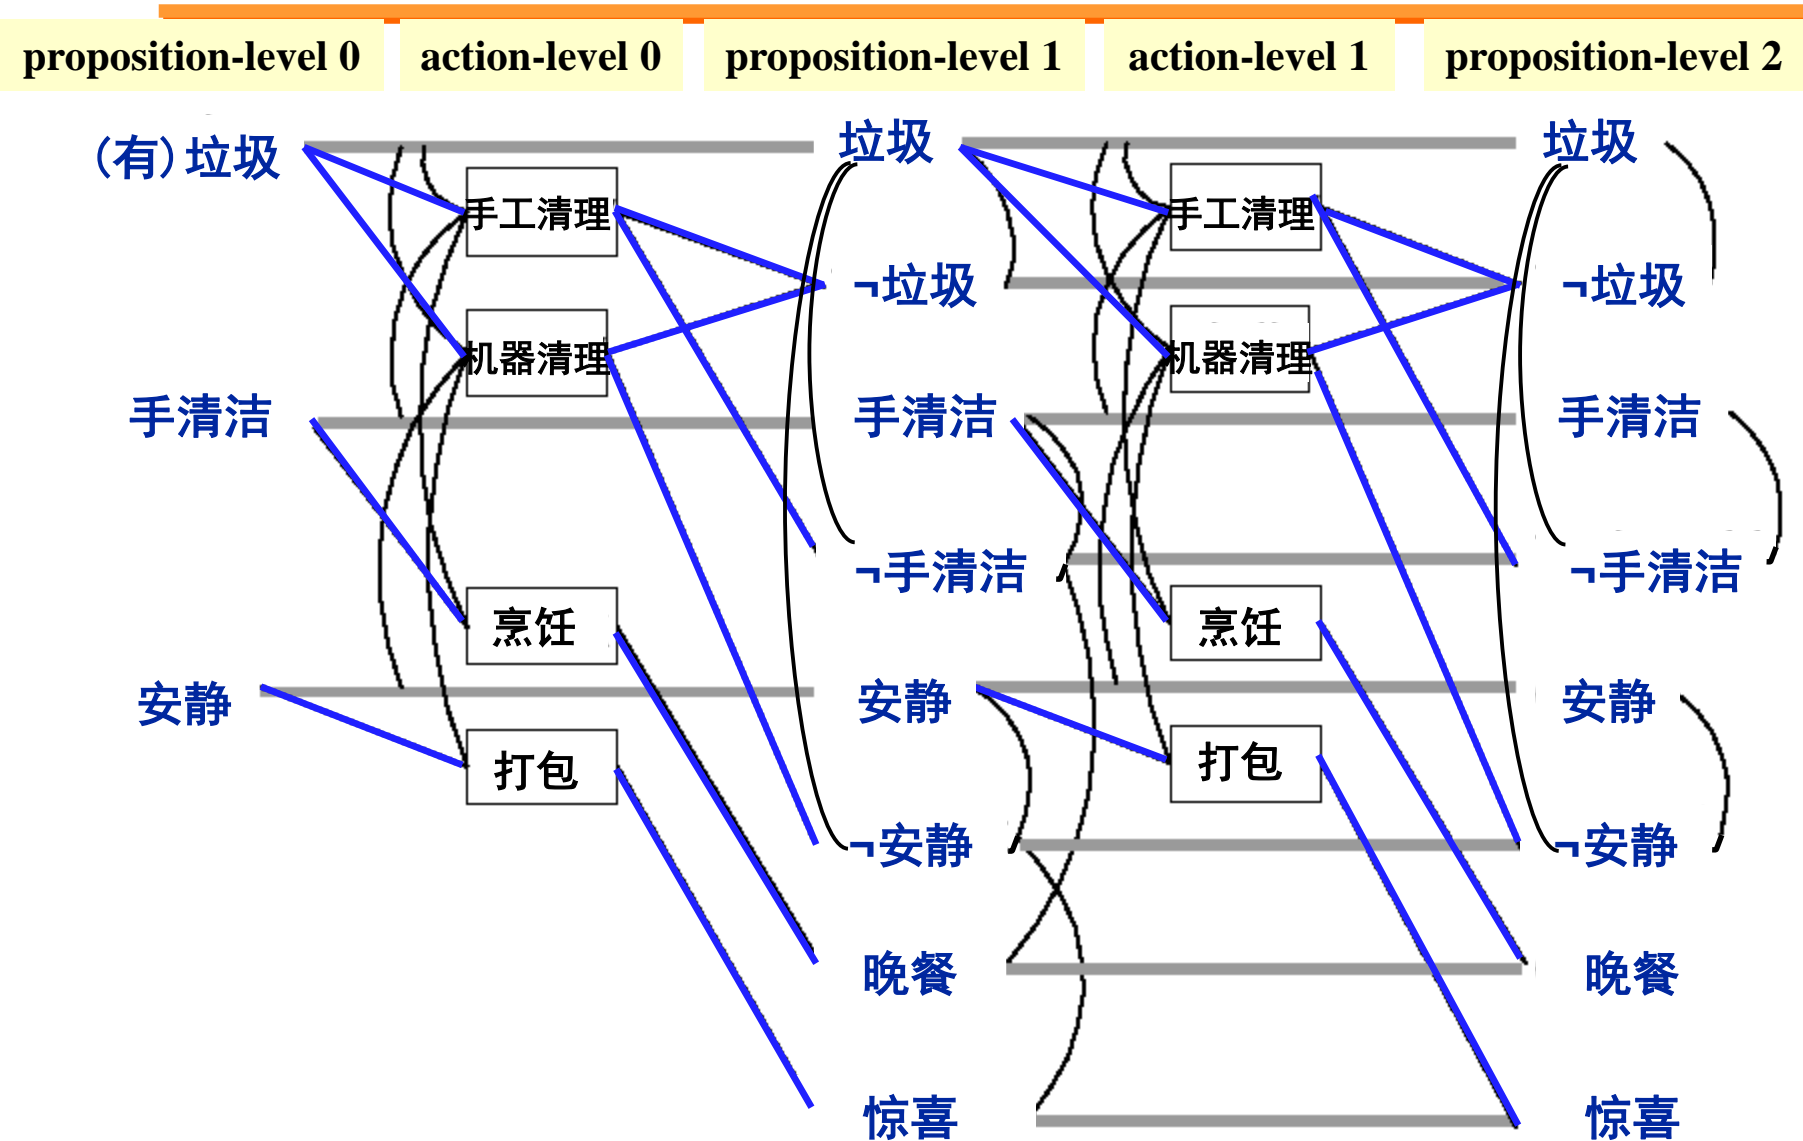
proposition-level 0

action-level 0


proposition-level 1



Dinner date example

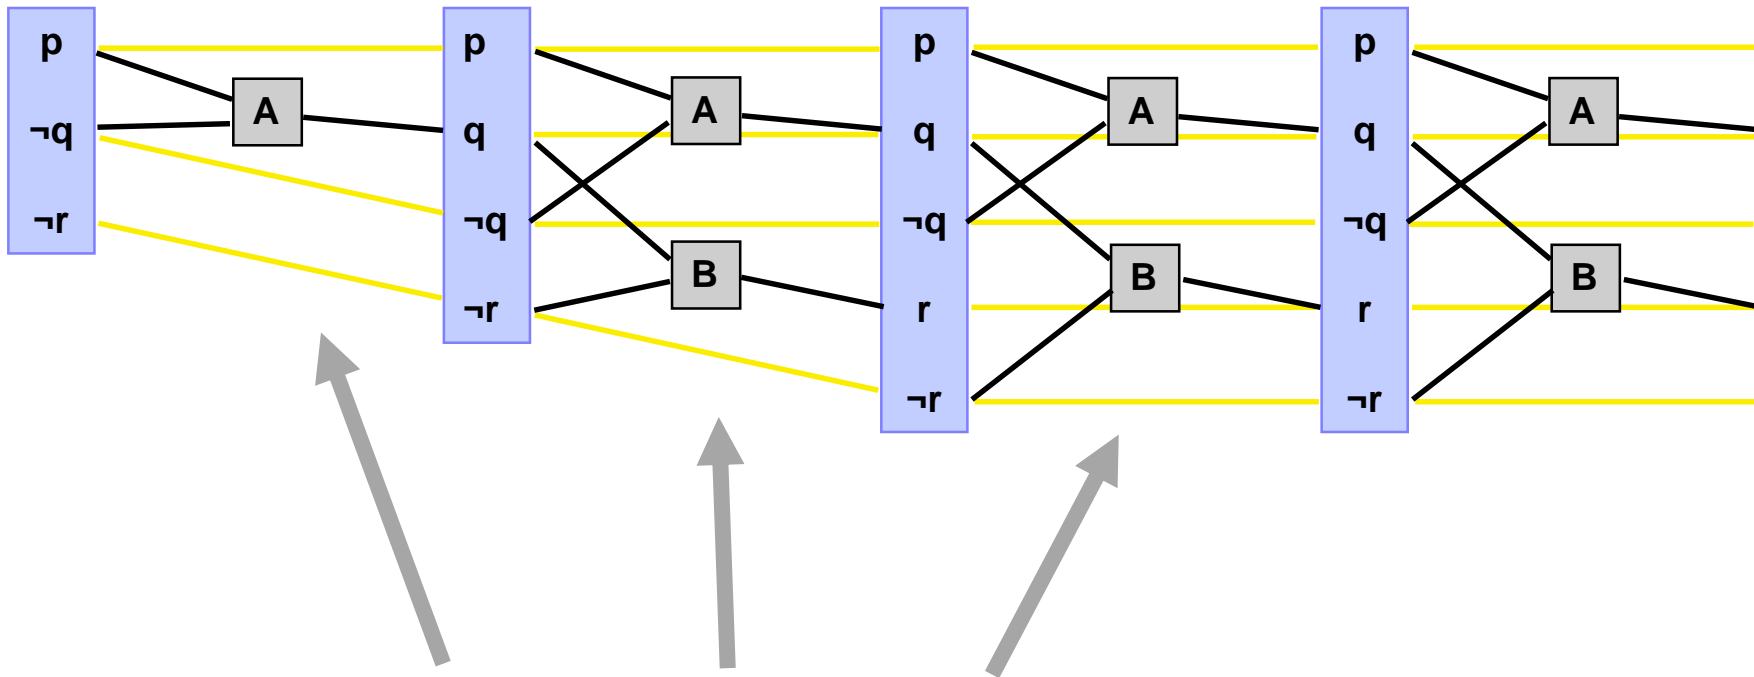


Observation 1




Propositions monotonically increase
(always carried forward by no-ops)

Observation 2



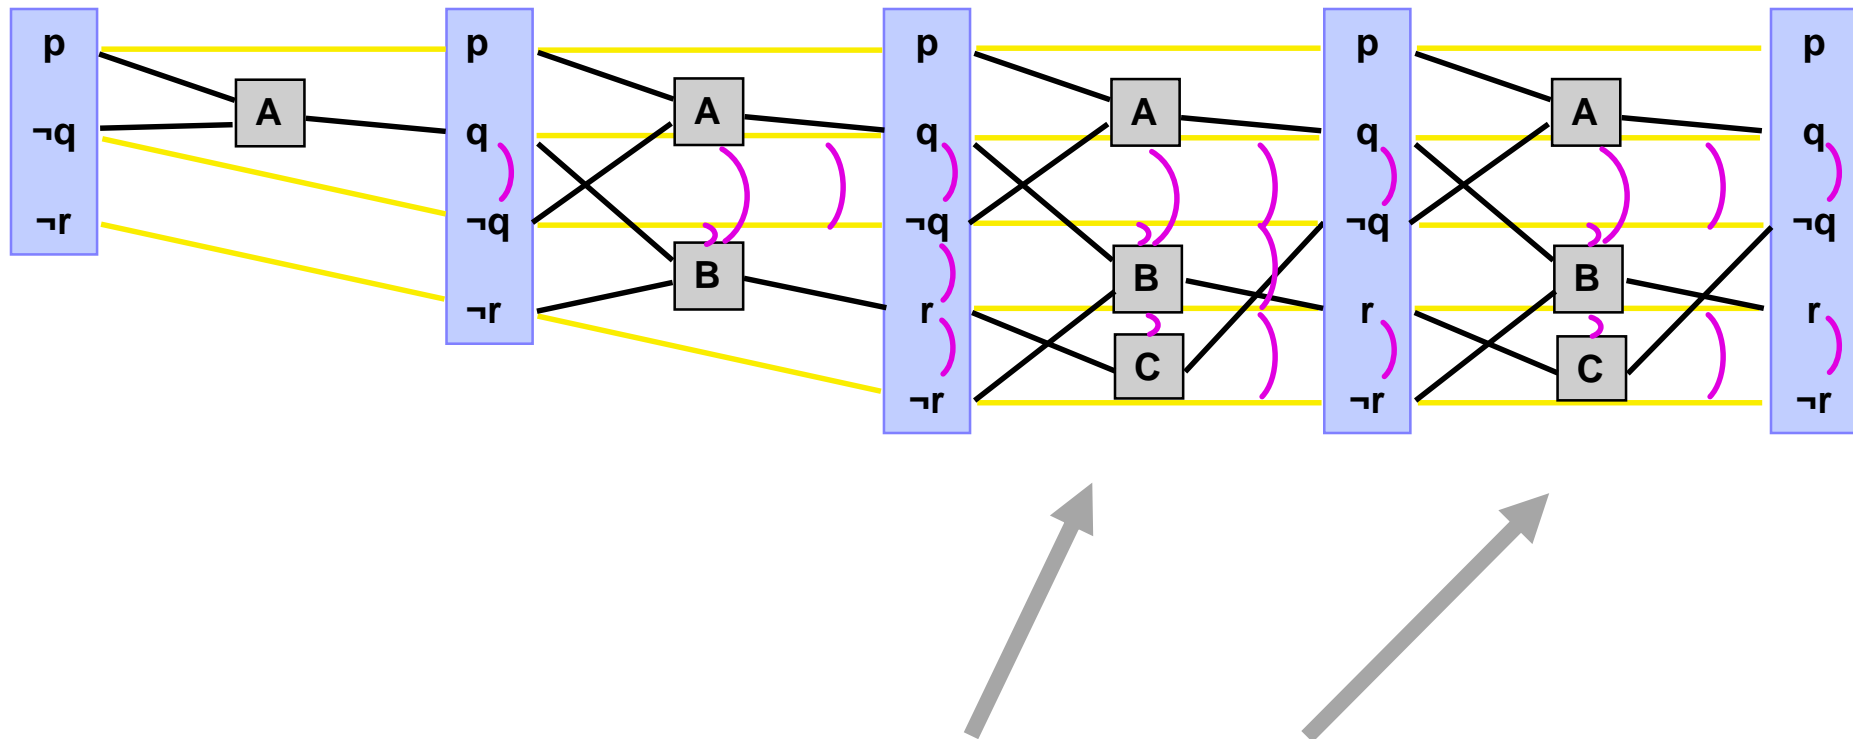
Actions monotonically increase

Observation 3



Proposition mutex relationships monotonically decrease

Observation 4



Action mutex relationships monotonically decrease

Observation 5

Planning Graph levels off (平稳)

- After some time steps all levels are identical
- Because it's a finite space, the set of propositional literals never decreases and mutexes don't reappear

Also provides a necessary and sufficient condition for termination of unsolvable problems:

- If planning graph has levelled off, yet goals are still unsatisfied

Valid plan

A *valid plan* is a subgraph of the planning graph where:

- Actions at the same level don't interfere or compete with each other (i.e. non-mutex actions)
- Each action's preconditions are made true by the plan
- Goals are satisfied

GraphPlan algorithm

Procedure GraphPlan:

- For $k = 0, 1, 2, \dots$

<Graph expansion>:

- Expand the planning graph level by level
 - If the planning graph levels off first, **fail**
until all goals are reachable and not mutex

<Solution extraction>:

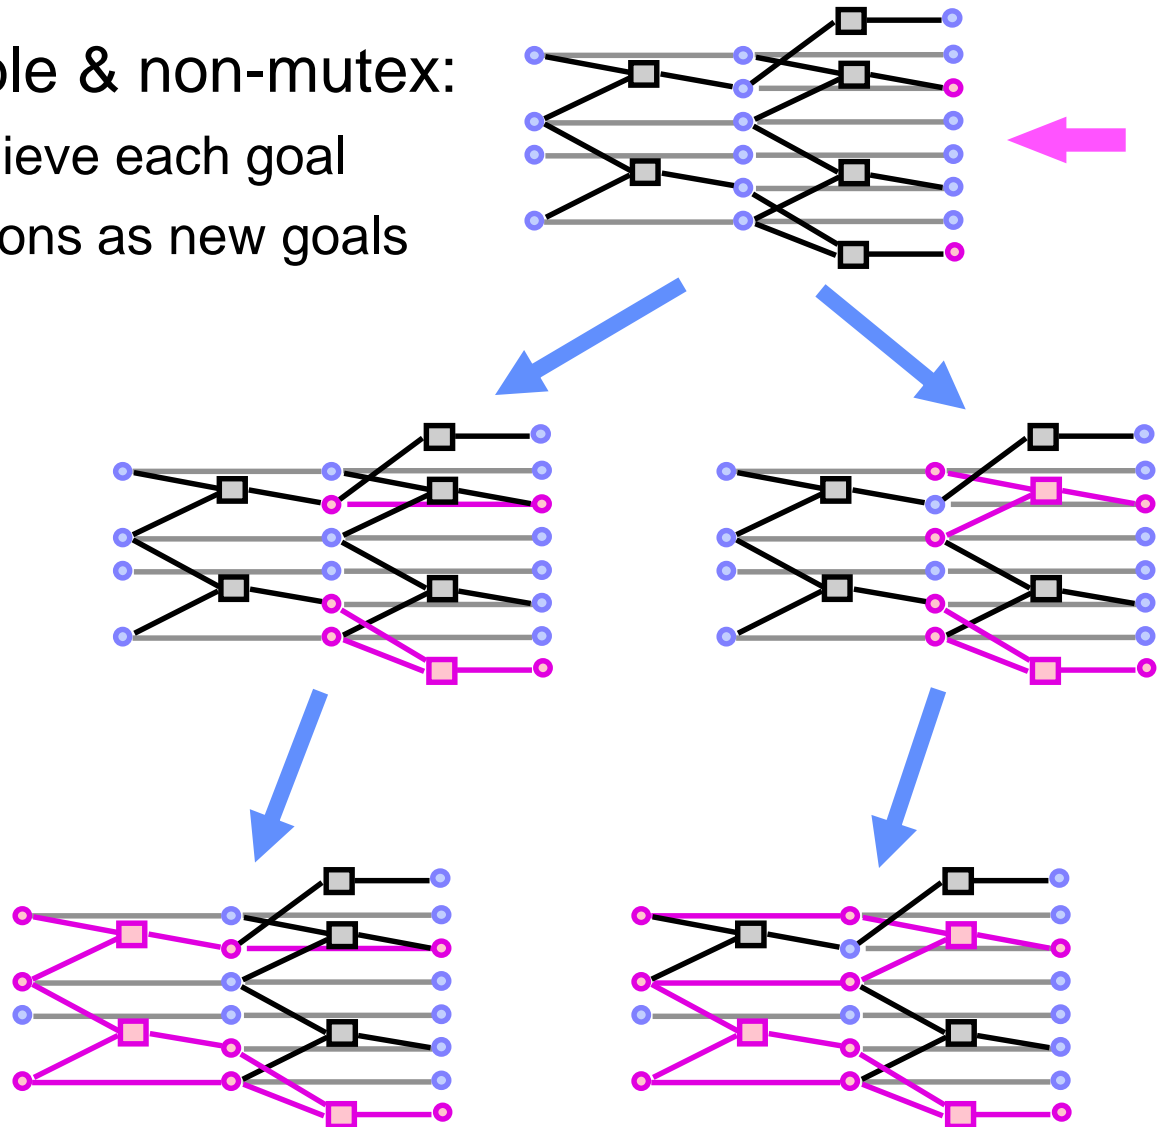
- Backward search the planning graph for a valid plan
 - If a solution is found, then **return** *solution*

Searching for a solution plan

- Level-by-level backward chaining on the planning graph using mutex constraints
- Given a set of non-mutex goals at level k , identify a subset of non-mutex actions (including no-ops) at level $k-1$ to achieve current goals. The preconditions of these actions become new goals for level $k-1$

Searching for a solution plan

- If goals are reachable & non-mutex:
 - Choose action to achieve each goal
 - Add action preconditions as new goals



Dinner date example（为恋人准备惊喜晚餐）


- Initial state: (有)垃圾, 手清洁, 安静
- Goal: 晚餐, 惊喜, \neg 垃圾
- Actions:
 - 烹饪: precondition (手清洁)
effect (晚餐)
 - 打包: precondition (安静)
effect (惊喜)
 - 手工清理: precondition (垃圾)
effect (\neg 垃圾, \neg 手清洁)
 - 机器清理: precondition (垃圾)
effect ((\neg 垃圾, \neg 安静))

Dinner date example


proposition-level 0

action-level 0

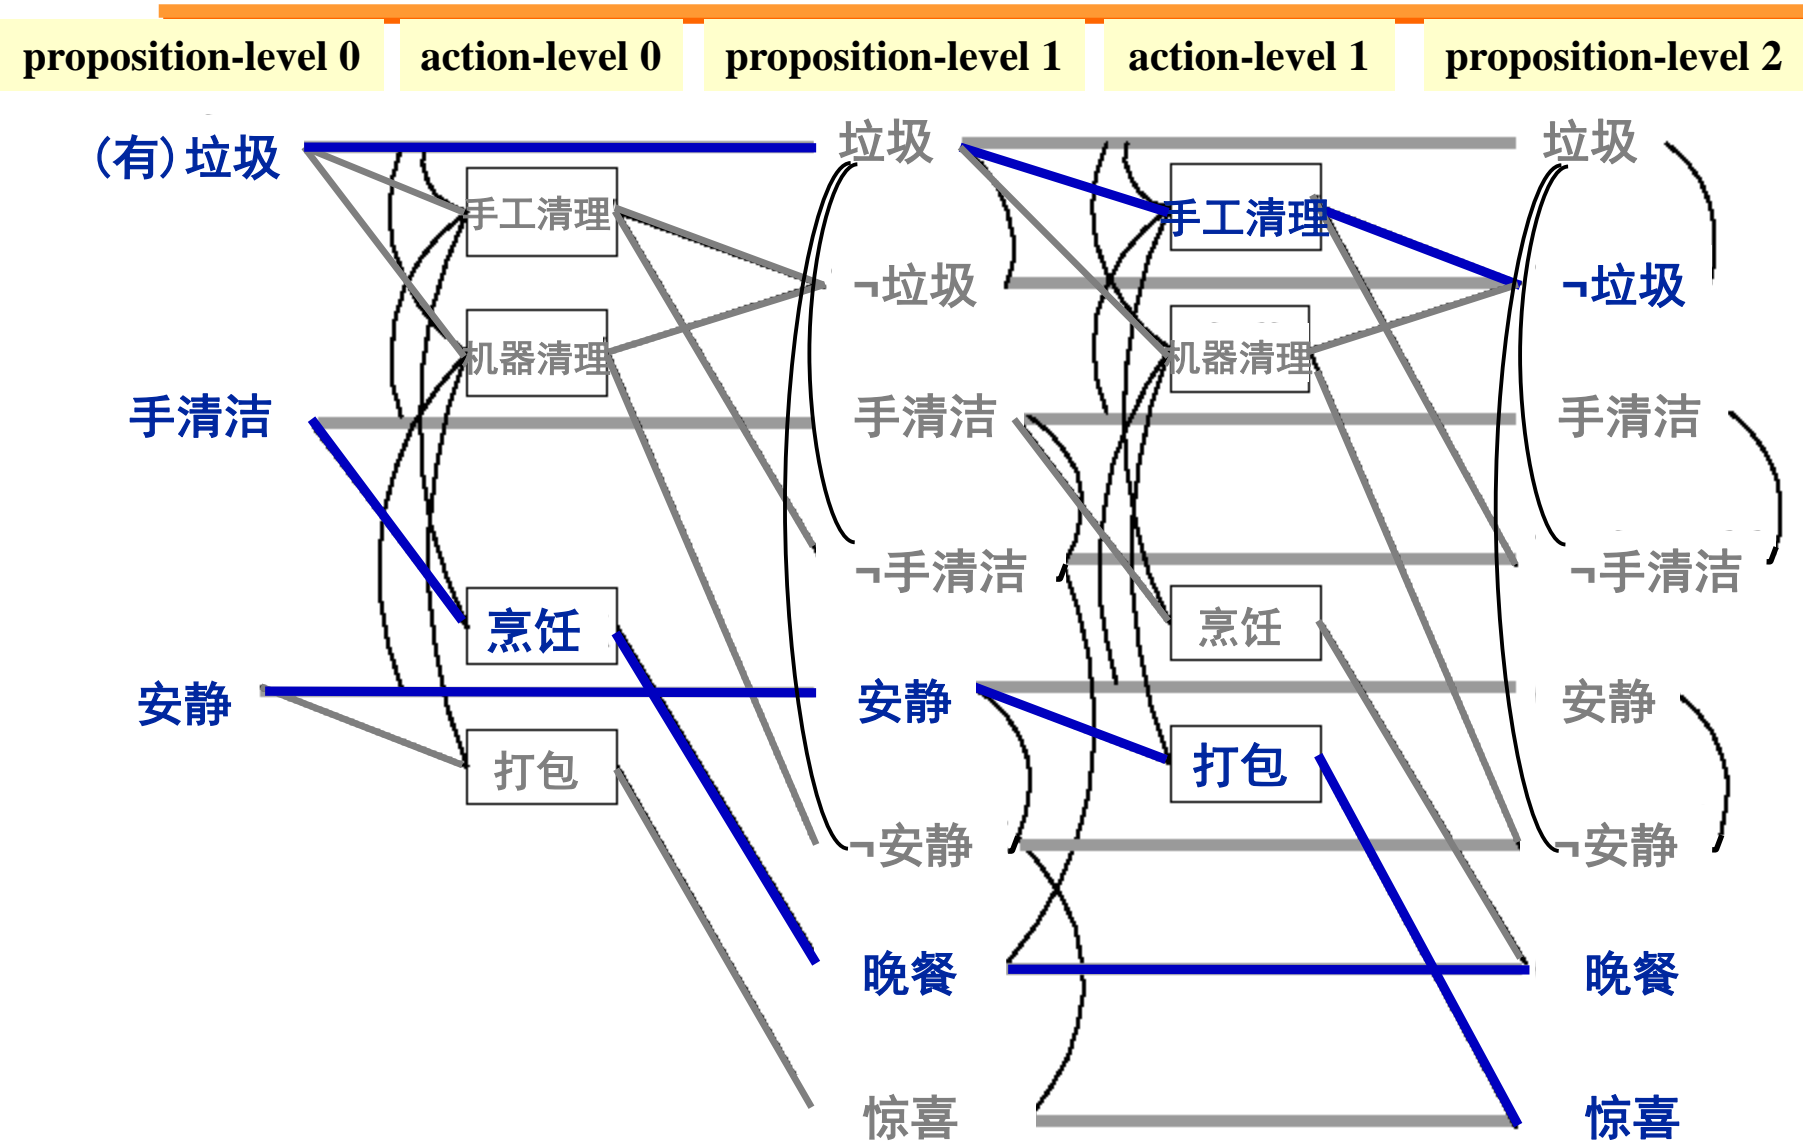
proposition-level 1



Dinner date example



Dinner date example



GraphPlan exercise

- Initial state: (有) 食材, (有) 垃圾, 手清洁, 安静
- Goal: 晚餐, 惊喜, \neg 垃圾
- Actions:
 - 烹饪: precondition (食材, 手清洁)
effect (晚餐)
 - 打包: precondition (晚餐, 安静)
effect (惊喜)
 - 手工清理: Precondition (垃圾)
effect (\neg 垃圾, \neg 手清洁)
 - 机器清理: precondition (垃圾)
effect ((\neg 垃圾, \neg 安静))

Properties of GraphPlan

Sound, complete and will terminate with failure if there is no plan

- Proved effective for solving *hard* planning problems
- Polynomial time graph construction
- Mutual exclusion (mutex) for pruning search
- Insensitivity to goal ordering
- Find “shortest parallel plan”

But work only for propositional planning problems (with no variables)

Comparison with plan-space planning

Advantage:

- The backward search of Graphplan (i.e. the hard part) only looks at actions in the planning graph
- Smaller search space than PSP, and thus faster

Disadvantage: 不能处理带变量. 量词.

- To generate the planning graph, Graphplan creates a huge number of ground atoms
- Many of them may be irrelevant

For classical planning, advantage outweighs disadvantage

FF and extensions

Fast-forward planner [Hoffmann 01]

- A* search with heuristic values from:
 - *Relaxed* planning graph – only add effects

Extension to pure Strips operators [Koehler et al 97]:

- Disjunctive preconditions
- Negated preconditions
- Conditional effects
- Universal quantification

内容回顾

Outline

- STRIPS-like plan representation
 - Planning with state-space search
 - Partial-order planning (POP)
 - Planning graphs (GraphPlan)

Strips representation for actions:

Move-C-from-A-to-Table:


Precondition: (on C A), (clear C)

Effect:

add (on-table C)

add (clear A)

delete (on C A)




- The explicit effects are the only changes to the state

Means-ends analysis（手段目的分析）

Strips's problem solving takes **means-ends analysis**

- Search by reducing the difference between state and goals
- What *means* (operators) are available to achieve the desired *ends* (goal)


The monkey and bananas problem（猴子香蕉问题）




Blocks world example (Sussman anomaly)

- Strips uses *noninterleaved planner* (非交叉规划器), which cannot solve this example ...

Initial:



Goal:




Initial state: on-table(A), on-table(B), on(C, A), clear(B), clear(C)

Goal: on(A,B), on(B,C) ...

Outline


- STRIPS-like plan representation
 - Planning with state-space search
- Partial-order planning (POP)
- Planning graphs (GraphPlan)

Search space: Blocks world




Planning algorithms

- Progression: Forward state-space search



- Regression: Backward state-space search



Progression vs regression

- Both algorithms are
 - *sound* (they always return a valid plan)
 - *complete* (if a valid plan exists they will find one)
- Complexity $O(b^n)$ worst-case
 - where $b = \text{branching factor}$,
n = number of “choose” operators
- Regression: often smaller b , focused by goals
- Progression: full state space to compute heuristics

Outline

- STRIPS-like plan representation
- Planning with state-space search
- Partial-order planning (POP)
- Planning graphs (GraphPlan)

Search the space of plans

Partial-Order Planning (POP): generates partial-order plans

- Nodes are partial plans
- Links are plan refinements
- Solution is a node (not a path)

Follow the least commitment principle

- E.g. don't commit to an order of actions until it is required


Partial plan representation

- Plan = (A, O, L), where
 - A: set of actions in the plan
 - O: *temporal orderings* between actions ($a < b$)
 - L: *causal links* linking actions via a literal

- Causal Link: $A_p \xrightarrow{Q} A_c$ (因果连接)

Action A_c (consumer) has precondition Q that is established in the plan by A_p (producer)

- To protect causal link:



POP algorithm

POP((A, O, L), agenda, actions)

Initial plan: {*Start*, *Finish*} and preconditions in *Finish* as open conditions

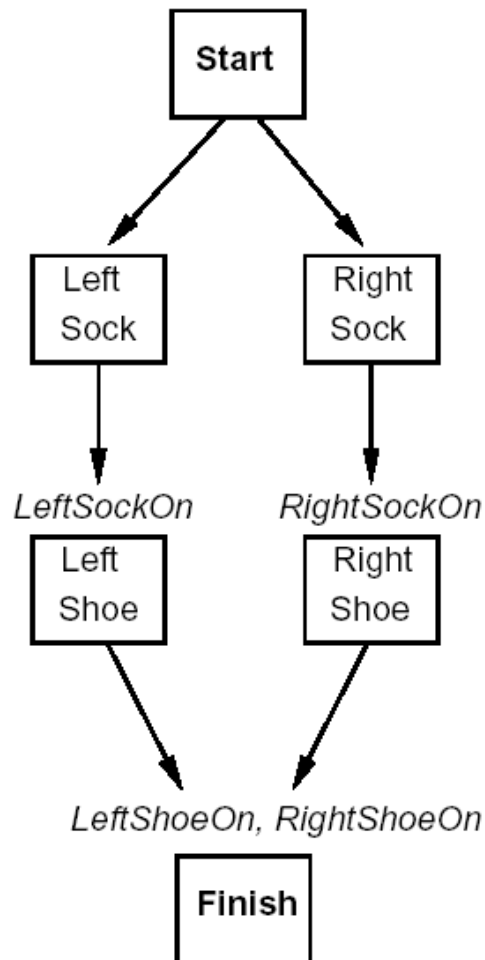
1. If **agenda** is empty, then **return** (A, O, L) 结束条件
2. Pick (Q, A_{need}) from **agenda** (子) 目标
3. **Choose** an action A_{add} that adds effect Q 动作选择
 - If no such action exists, **fail**
 - Add the link $A_{\text{add}} \xrightarrow{Q} A_{\text{need}}$ to **L** and the ordering $A_{\text{add}} < A_{\text{need}}$ to **O**
 - If A_{add} is new, add it to **A** 规划扩充
4. Remove (Q, A_{need}) from **agenda**. If A_{add} is new, for each of its preconditions P add (P, A_{add}) to **agenda** 更新(子) 目标
5. For every action A_t in A that threatens any causal link $A_p \rightarrow A_c$ in **L**
 - **Choose** to add $A_t < A_p$ or $A_c < A_t$ to **O**
 - If neither choice is consistent, **fail**
6. POP((A, O, L), **agenda**, actions)

保护因果连接:

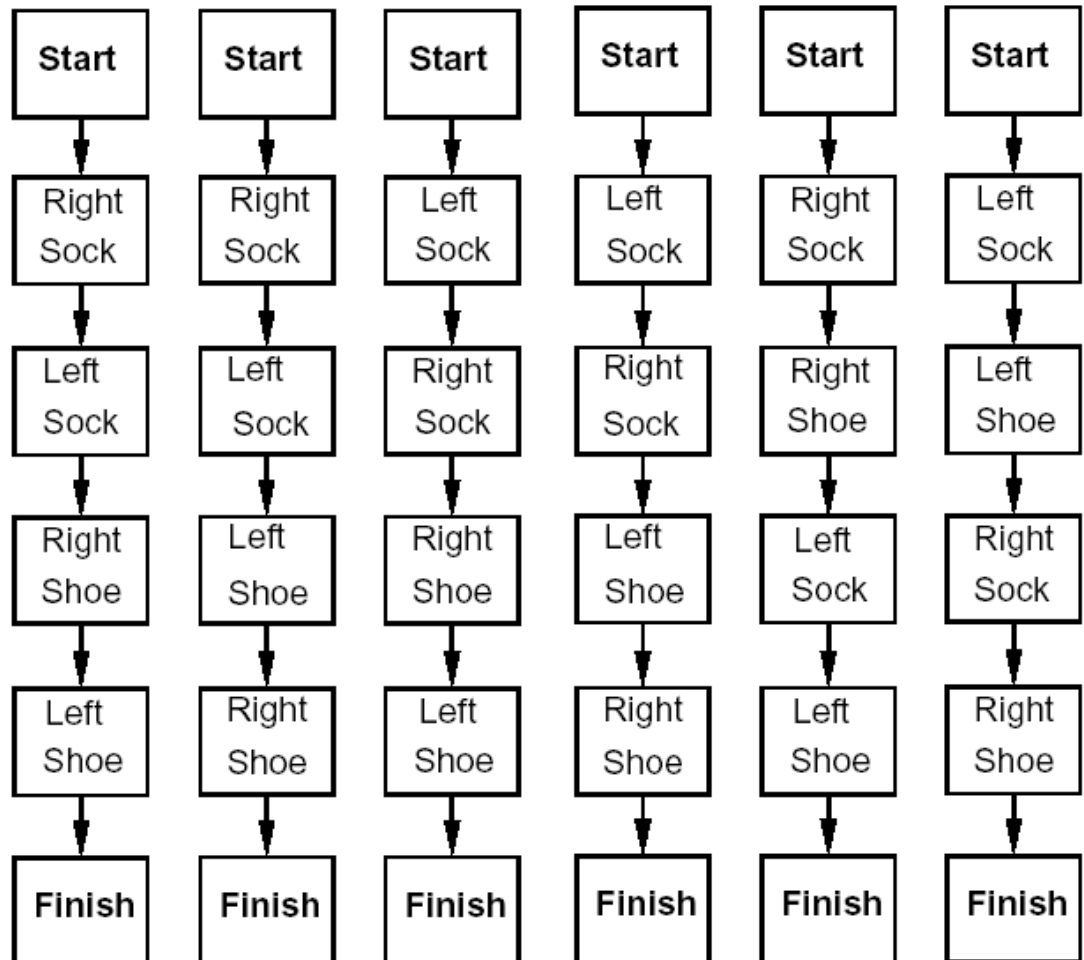
- 降级(Demotion): $A_t < A_p$
- 升级(Promotion): $A_c < A_t$

Total-order vs Partial-order plans

Partial Order Plan:



Total Order Plans: (线性化)



Properties of POP

POP is *sound* and *complete*

- Can be much faster than state-space planning, because of no need to backtrack over goal orderings (so less branching is required)
- Although it is more expensive per node and makes more choices than Regression, reduction in branching size often gains more
 - Larger n but smaller b

Flexibility gained by partial order

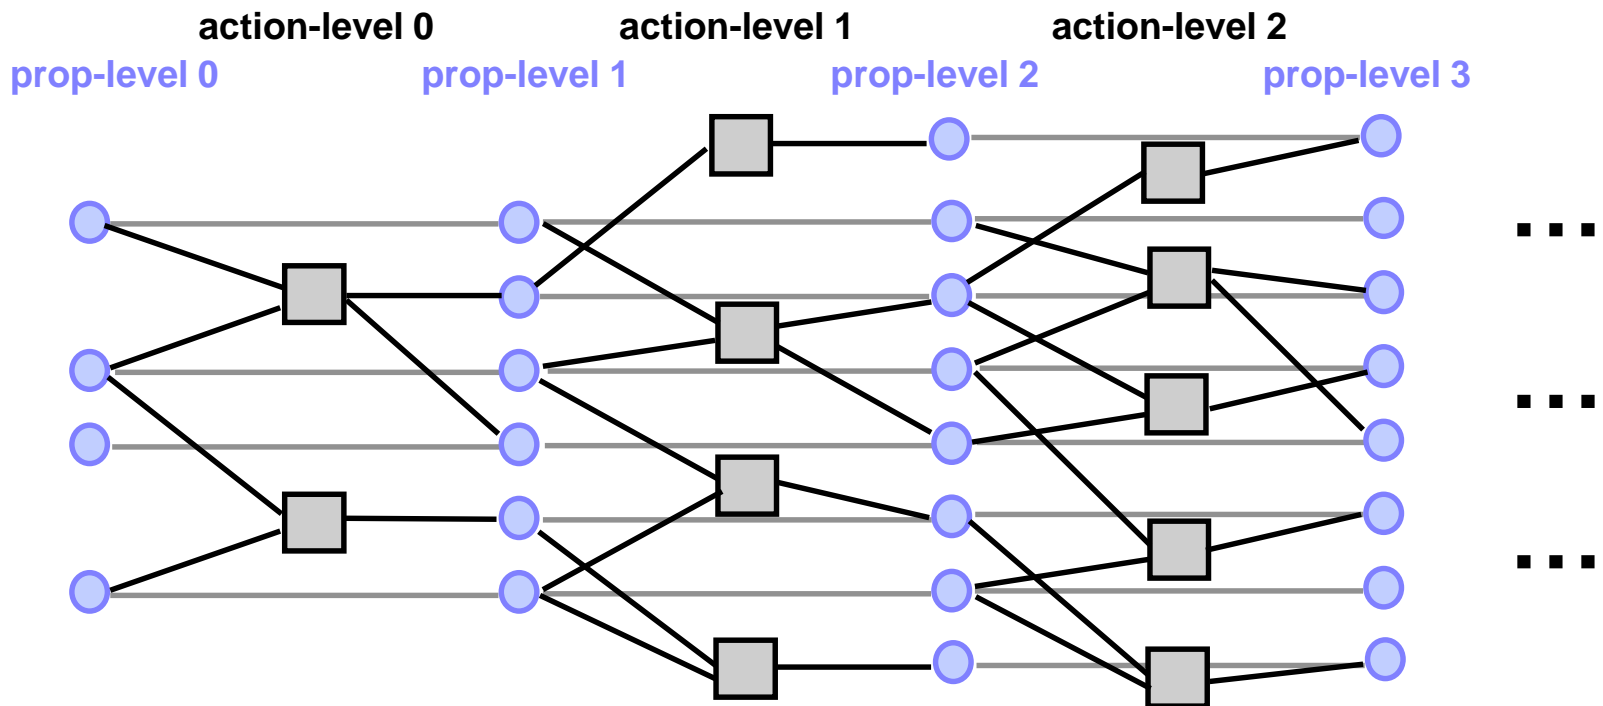
- Can be very useful to agent when the world fails to cooperate
- Make it easier to combine smaller plans into larger ones
 - Each small plan can reorder its actions to avoid conflict with the other plans

Outline

- STRIPS-like plan representation
- Planning with state-space search
- Partial-order planning (POP)
- Planning graphs (GraphPlan)

A planning graph (规划图)

- Consist of a sequence of **levels** that correspond to time steps in the plan, where level 0 is the initial state
- Each level contains a set of propositional literals and a set of actions, with edges connecting action preconditions and effects



Determining mutex relations (互斥关系)

- Two actions A and B are mutex at an action-level, if:
 - Interference: A (or B) disables a *precondition* or an *effect* of B (or A)
 - Competing needs: A and B have *mutex preconditions*
- Two propositions p and q are mutex at a proposition-level if
 - Negation: p (or q) is the *negation* of q (or p)
 - Inconsistent support: *all ways* of achieving them are *mutex* (i.e. all the actions that add p are mutual exclusion of all the actions that add q)

GraphPlan algorithm

Procedure GraphPlan:

- For $k = 0, 1, 2, \dots$

<Graph expansion>:

- Expand the planning graph level by level
 - If the planning graph levels off first, **fail**
until all goals are reachable and not mutex

<Solution extraction>:

- Backward search the planning graph for a valid plan
 - If a solution is found, then **return** *solution*

Searching for a solution plan

- Level-by-level backward chaining on the planning graph using mutex constraints
- Given a set of non-mutex goals at level k , identify a subset of non-mutex actions (including no-ops) at level $k-1$ to achieve current goals. The preconditions of these actions become new goals for level $k-1$

A *valid plan* is a subgraph of the planning graph where:

- Actions at the same level don't interfere or compete with each other (i.e. non-mutex actions)
- Each action's preconditions are made true by the plan
- Goals are satisfied

Properties of GraphPlan

Sound, complete and will terminate with failure if there is no plan

- Proved effective for solving *hard* planning problems
- Polynomial time graph construction
- Mutual exclusion (mutex) for pruning search
- Insensitivity to goal ordering
- Find “shortest parallel plan”

But work only for propositional planning problems (with no variables)

References

- R. Fikes and N. Nilsson. STRIPS: A New Approach to the Application of Theorem Proving to Problem Solving. *Artificial Intelligence*, 2(3-4), 1971
- D. Weld. An Introduction to Least-Commitment Planning. *AI Magazine*, 15(4), 1994 (POP)
- J. S. Penberthy and D. Weld. UCPOP: A Sound, Complete, Partial-Order Planner for ADL. *Proceedings of KR*, 1992
- A. Blum and M. Furst. Fast Planning Through Planning Graph Analysis. *Artificial Intelligence*, 90(1-2), 1997 (GraphPlan)
- J. Hoffmann. FF: The Fast-Forward Planning System. *AI Magazine*, 22(3), 2001

Resources

- GraphPlan Planner:
<https://en.wikipedia.org/wiki/Graphplan> (External links)
- UCPOP Planner:
<https://www.swmath.org/software/20687>
- Action Description Language (ADL):
https://handwiki.org/wiki/Action_description_language
- Planning Domain Definition Language (PDDL):
<https://planning.wiki/guide/whatis/pddl>

End.