

独立于算法的机器学习

山世光

中国科学院计算技术研究所

sgshan@ict.ac.cn



中国科学院计算技术研究所
Institute of Computing Technology, Chinese Academy of Sciences



课前思考

- 你们学过的模型中哪个最好？为什么？如何比较两个不同模型的优劣？
- Bias和variance分别描述了算法的什么性质？
- 如果有很多可选算法，怎么集成它们？
- 数据多样、规模极大，如何利用好它们？
- 特征维度特别高，如何利用好它们？



参考文献

- 第九章 R. Duda, P. Hart, D. Stork, Pattern Classification (Second edition), John Wiley & Sons, New York, USA, 2000



What's in This Chapter?

- Algorithm-Independent by definition
 - to those mathematical foundations that do not depend upon the particular classifier or learning algorithm used.
 - techniques that can be used in conjunction with different learning algorithms, or provide guidance in their use.



Problems to Answer

- Many algorithms/techniques in hand
 - ☐ Which is the “best”?
 - ☐ Are there any reasons to favor one algorithm over another?
 - ☐ Can we even find an algorithm that is overall superior to (or inferior to) random guessing?
 - ☐ Which one to choose given one problem?




Outline of This Chapter

- Some philosophy in PR/ML
 - No Free Lunch Theorem
 - Ugly Duckling Theorem
 - Minimum Description Length principle
 - Occam's razor
- Resampling for classifier design
 - Bagging
 - Boosting
 - AdaBoost
 - Active Learning
- Estimating and comparing classifiers
 - Cross validation
- Self-paced learning and curriculum learning

No Free Lunch Theorem

- [Wolpert, 1996] shows that
 - In a noise-free scenario where the loss function is the misclassification rate, if one is interested in *off-training-set error*, there are no a priori distinctions between learning algorithms.





No Free Lunch Theorem

- All algorithms are equivalent, on average, by any of the following measures of error: $E(L/D)$, $E(L/n)$, $E(L/f, D)$, or $E(L/f, n)$, where
 - D = training set;
 - n = number of elements in training set;
 - f = ‘target’ input-output relationships;
 - h = hypothesis (the algorithm's guess for f made in response to D); and
 - L = off-training-set ‘loss’ associated with f and h (‘generalization error’)



Implications of NFL

- There are no i and j such that, for all F ,

$$E_i(E/F, n) < E_j(E/F, n)$$

if all target functions $F(\mathbf{x})$ are equally likely.

- Furthermore, even if we know D , averaged over all target functions, no learning algorithm yields an **off-training set error** that is superior to any other.
- All statements of the form “**learning/recognition algorithm 1 is better than algorithm 2**” are ultimately statements about the **relevant target functions**.
- It is the **assumptions** about the learning domains that are relevant.

Ugly Duckling Theorem

- Problem to answer
 - In the absence of prior information, is there a principled reason to judge any two distinct patterns as more or less similar than two other distinct patterns?
- Ugly Duckling Theorem [Watanabe, 1969]





Ugly Duckling Theorem

■ Problem to answer

- In the absence of prior information, is there a principled reason to judge any two distinct patterns as more or less similar than two other distinct patterns?

■ Ugly Duckling Theorem [Watanabe, 1969]

- All things being equal. An ugly duckling is just as similar to a swan as two swans are to each other.
- 丑小鸭与白天鹅之间的区别和两只白天鹅之间的区别一样大（依赖于分类标准或依据）

Watanabe, Satoshi (1969). *Knowing and Guessing: A Quantitative Study of Inference and Information*. New York: Wiley. pp. 376–377.



Ugly Duckling Theorem

■ Implications

- In the absence of assumptions there is no privileged or “best” feature representation.
 - There is no problem-independent or privileged or “best” set of features or feature attributes.
- Even the apparently simple notion of similarity between patterns is fundamentally based on implicit assumptions about the problem domain



MDL Principle

- Aims at finding some **irreducible, smallest representation**
- We should minimize the sum of the *model's algorithmic complexity* and the description of the training data with respect to that model, i.e.,

$$K(h, D) = K(h) + K(D \text{ using } h).$$

with $K(\cdot)$ the Kolmogorov complexity, a measure of the **incompressibility**.



MDL Principle

- Example: decision tree classifiers
 - The **algorithmic complexity** of the model is *proportional to the number of nodes*.
 - The **complexity of the data** given the model can be expressed in terms of the *weighted sum of the entropies of the data at the leaf nodes*.
 - Thus, if the tree is **pruned** based on an entropy criterion, it is using MDL.
- Example: Neural Network
 - Deep network compression by pruning
 - Removal of some connections between neurons



MDL Principle

- Theoretically classifiers designed with an MDL principle are guaranteed to converge to the ideal or true model *in the limit of more and more data*.
- The MDL principle states that simple models (*smaller $K(h)$*) are to be preferred, and thus amounts to a bias to *simplicity*.



Occam's Razor

- Philosophy Principle Occam's Razor
 - “Entities” (or explanations) should not be multiplied beyond necessity.
如无必要，勿增实体
 - Among competing hypotheses, *the one with the fewest assumptions* should be selected.
 - For PR/ML, **NOT use machines that are more complicated than necessary**
 - “Necessary” can be determined by the quality of fitting to the training data.



Occam's Razor

- Techniques to avoid overfitting
 - Simplicity
 - Pruning
 - Regularization
 - Inclusion of penalty terms
 - Minimizing a description length...
- Seems conflict with NFL?
 - For a given training error, why do we generally prefer simple classifiers with fewer features and parameters?




Occam's Razor

- Not conflict with NFL, but imply that problems addressed so far *favor simpler classifiers*. Why?
- Evolution bias: strong selection pressure on our pattern recognition apparatuses to be computationally simple
 - Fewer neurons
 - Less time
 - Less energy cost


Bias and Variance Dilemma

- Two ways measuring the “match” or “alignment” of the model to the problem
 - **Bias**: accuracy/quality of the match
 - **Variance**: precision/specificity of the match

Overfitting-Classification



Overfitting-Regression





Bias and Variance Dilemma

- **Bias**: model fits training data well,
 - Low bias: favor complex models
- **Variance**: model has capacity to accommodate different testing data
 - Low variance: favor simpler models
- **Discussion**
 - How about deep learning?
 - Why can network be compressed but with accuracy preserved?
 - Why not train the simpler network directly?



Outline of This Chapter

- Some philosophy in PR/ML
 - ☐ No Free Lunch Theorem
 - ☐ Ugly Duckling Theorem
 - ☐ Minimum Description Length principle
 - ☐ Occam's razor
- Resampling for classifier design
 - ☐ Bagging
 - ☐ Boosting
 - ☐ AdaBoost
 - ☐ Active Learning
- Estimating and comparing classifiers
 - ☐ Cross validation

■ What?

- Sample a (sub)set from original training set
 - Jackknife (leave one out)
 - **Bootstrap**: randomly selecting n points from the training set D , with replacement
 - Reweighting each points

■ Why?

- Yield a more informative estimate of a general statistic.
- Good for improve classifiers.



Arcing methods

- Arcing: adaptive reweighting and combining
 - Techniques by reusing or selecting data in order to improve classification
 - Bagging: bootstrap aggregating
 - Independently bootstrap data sets
 - Boosting
 - Dependently bootstrap data sets
 - AdaBoost

- Bagging: bootstrap aggregating
 - Proposed by [Breiman, 1996]
 - Derived from bootstrap [Efron, 1993]
- Basic idea
 - Create classifiers using training sets bootstrapped independently (drawn with replacement)
 - Average results of each component classifiers

■ Algorithm

Given a training set

$$D = \{(x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)\}$$

- 1. Sample m sets D_1, D_2, \dots, D_m of n elements from D (with replacement)
- 2. Train a component classifier/regression f_i from each D_i
- 3. The final classifiers is
$$f(x) = \text{sum/vote}(f_1(x), f_2(x), \dots, f_m(x))$$

Bagging Example (Opitz, 1999)

- Bootstrap data sets
 - With replacement
 - Independently resampled

Original training set	1	2	3	4	5	6	7	8
Training set 1	2	7	8	3	7	6	3	1
Training set 2	7	8	5	6	4	2	7	1
Training set 3	3	6	2	7	5	6	2	2
Training set 4	4	5	1	4	6	4	3	8



Bagging

- Discussion
 - Why?

- Component classifiers selection
 - Generally of the same general form
 - SVM, NN, ANN, tree...
- Effects
 - Improves recognition for *unstable* classifiers since it effectively averages over such **discontinuities**
 - Unstable (related to *high variance*)

$$f(D) \approx (D + \Delta D)$$
 - “small” changes in the training data lead to significantly different classifiers and relatively “large” changes in accuracy.
 - 鲁棒性差：易被攻击， $f(x) \approx f(x + \Delta x)$



Arcing methods

- Arcing: adaptive reweighting and combining
 - Techniques by reusing or selecting data in order to improve classification
 - Bagging: bootstrap aggregating
 - Independently bootstrap data sets
 - Boosting
 - Dependently bootstrap data sets
 - AdaBoost



Boosting

- Powerful technique for combining multiple weak “base” learners to form a **committee** whose performance can be significantly better than any of the base classifiers
 - Originated from [Schapire, 1989]
- Basic idea
 - **Sequential production of classifiers**: each classifier dependent on the previous one, and focuses on the previous one’s failures
 - **Examples incorrectly predicted** in previous classifiers say louder in the next round



A Formal Description of Boosting

- Given training set $X = \{(x_1, y_1), \dots, (x_n, y_n)\}$
 $y_i \in \{+1, -1\}$ is the label of instance x_i
- for $t = 1, \dots, T$:

- Construct a **new** distribution D_t from X
- Find weak classifier

$$h_t: X \rightarrow \{+1, -1\}$$

with small error ε_t on D_t :

$$\varepsilon_t = \Pr_{i \sim D_t}[h_t(x_i) \neq y_i]$$

- Output final classifier $H_{\text{final}} = \text{weighted sum}(h_t)$

Example Boosting Setting

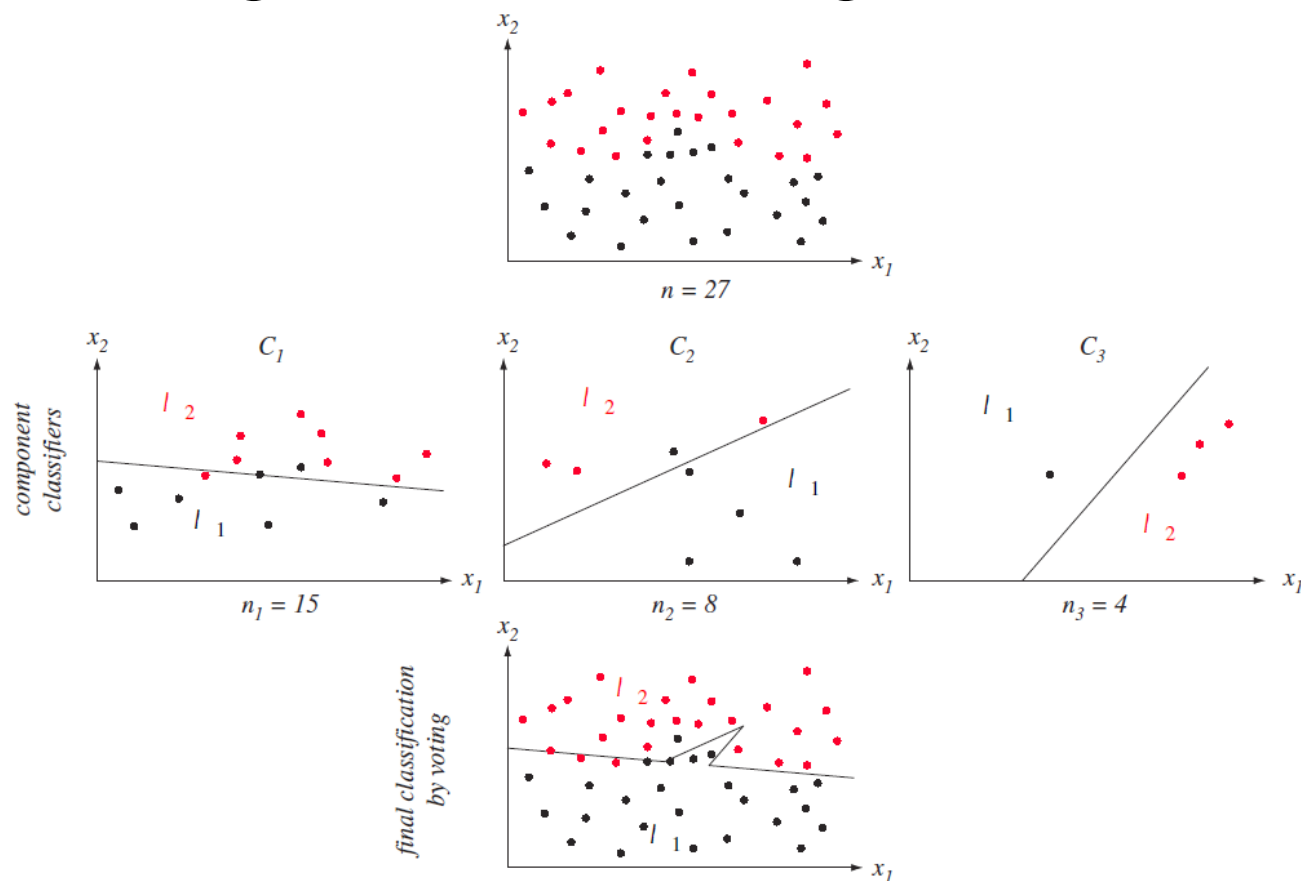
- $D_1 = \text{randomly select a subset of } X$
- $D_2 = \text{select from } X/D_1, \{ \text{half } \textcolor{blue}{\text{correctly}} \text{ classified by } h_1 \} + \{ \text{half } \textcolor{red}{\text{incorrectly}} \text{ classified by } h_1 \}$
- $D_3 = \{x_i \in (X/D_1 \cup D_2) \text{ and } h_1(x_i) \neq h_2(x_i)\}$

- The final classifier:

$$h_{\text{final}}(x) = \begin{cases} h_1(x); & \text{if } h_1(x) == h_2(x) \\ h_3(x); & \text{otherwise} \end{cases}$$

Example Boosting Setting

- Component classifiers: LMS
- Sampling: basic boosting procedure





Many Variations

- AdaBoost.M1, AdaBoost.MR, FilterBoost, GentleBoost, GradientBoost, MadaBoost, LogitBoost, LPBoost, MultiBoost, RealBoost, RobustBoost, ...



From Bagging to Boosting

- Base classifiers are trained in sequence
- Each base classifier trained using a weighted form of the dataset
 - Weighting coefficient depends on the performance of the previous classifiers
 - Points **misclassified** by previous classifiers are **given more weights** in training next classifier
- Decisions are combined using a **weighted majority voting** scheme



Arcing methods

- Arcing: adaptive reweighting and combining
 - Techniques by reusing or selecting data in order to improve classification
 - Bagging: bootstrap aggregating
 - Independently bootstrap data sets
 - Boosting
 - Dependently bootstrap data sets
 - AdaBoost



AdaBoost

- Proposed by [Freund & Schapire'95]:
 - Strong practical advantages over previous boosting algorithms
 - With amazing generalization ability
- Answer the open problem
 - An open problem [Kearns & Valiant, STOC'89]:
“weakly learnable” ?= “strongly learnable”
 - In intuitive words, whether a “weak” learning algorithm that works just **slightly better than random guess** can be “boosted” into an arbitrarily accurate “strong” learning algorithm!



The Born of AdaBoost

- Amazingly, in 1990 Schapire proves that the answer is “yes”. More importantly, the proof is a construction! This is the first Boosting algorithm
- In 1993, Freund presents a scheme of combining weak learners by majority voting in PhD thesis at UC Santa Cruz

However, these algorithms are not practical!

- Later, at AT&T Bell Labs, Freund & Schapire published the 1997 journal paper (the work was reported in EuroCOLT'95), which proposed the AdaBoost algorithm, a practical algorithm.

New Resampling Mechanism

- Given training set $X = \{(x_1, y_1), \dots, (x_m, y_m)\}$

$y_i \in \{+1, -1\}$ is the label of instance x_i

- $D_1(i) = \frac{1}{m}$; given D_t and h_t , then

$$D_{t+1}(i) = \frac{D_t(i)}{Z_t} \times \begin{cases} e^{-\alpha_t}, & \text{if } y_i = h_t(x_i) \\ e^{\alpha_t}, & \text{if } y_i \neq h_t(x_i) \end{cases} = \frac{D_t(i)}{Z_t} \exp(-\alpha_t y_i h_t(x_i))$$

where Z_t is a normalization factor, and

$$\alpha_t = \frac{1}{2} \ln \left(\frac{1 - \varepsilon_t}{\varepsilon_t} \right) > 0, \text{ with } \varepsilon_t = P_{i \sim D_t} [h_t(x_i) \neq y_i] < 0.5$$

- Final classifier

$$H_{final}(x) = \text{sign} \left(\sum_t \alpha_t h_t(x) \right)$$



AdaBoost Algorithm

■ Weights of misclassified samples are increase in (t+1)th iteration.

- given training set $(x_1, y_1), \dots, (x_m, y_m)$
where $x_i \in X$, $y_i \in \{-1, +1\}$
- initialize $D_1(i) = 1/m$ ($\forall i$)
- for $t = 1, \dots, T$:
 - train weak classifier $h_t : X \rightarrow \{-1, +1\}$ with error $\epsilon_t = \Pr_{i \sim D_t} [h_t(x_i) \neq y_i]$
 - $\alpha_t = \frac{1}{2} \ln \left(\frac{1 - \epsilon_t}{\epsilon_t} \right)$
 - update $\forall i$:

$$D_{t+1}(i) = \frac{D_t(i)}{Z_t} \exp(-\alpha_t y_i h_t(x_i))$$


where Z_t = normalization factor

- $H_{\text{final}}(x) = \text{sign} \left(\sum_{t=1}^T \alpha_t h_t(x) \right)$




Toy Example of AdaBoost

■ Initialize




Toy Example of AdaBoost

■ Round 1




Toy Example of AdaBoost

■ Round 1




Toy Example of AdaBoost

■ Round 2




$$\epsilon_1 = 0.30$$

$$\alpha_1 = 0.42$$




$$\epsilon_2 = 0.21$$

$$\alpha_2 = 0.65$$




Toy Example of AdaBoost

Round 3




$$\epsilon_1 = 0.30$$

$$\alpha_1 = 0.42$$



$$\epsilon_2 = 0.21$$

$$\alpha_2 = 0.65$$

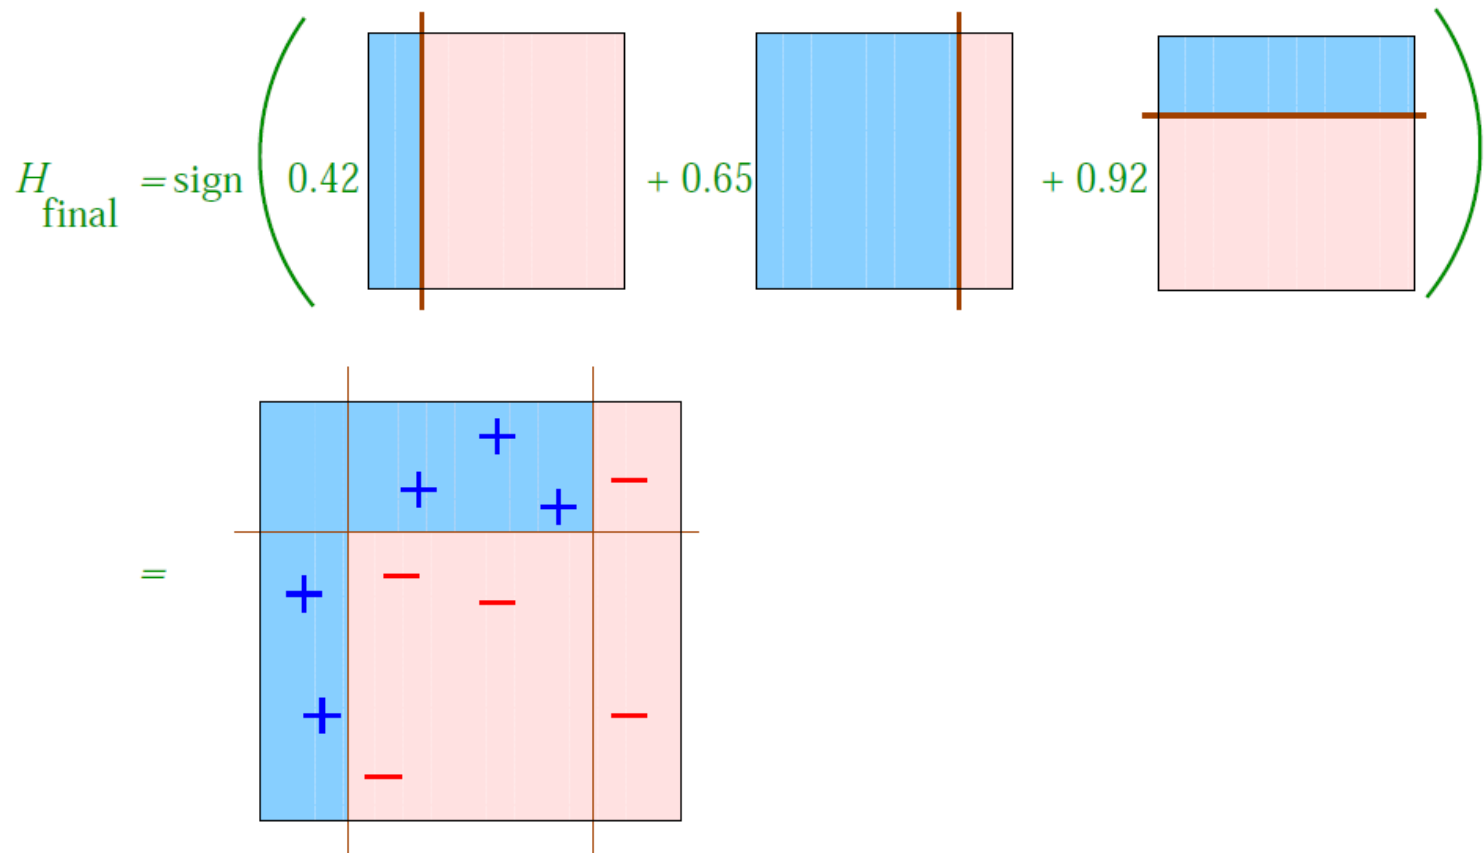


$$\epsilon_3 = 0.14$$

$$\alpha_3 = 0.92$$

Toy Example of AdaBoost

■ Final Strong Classifier



AdaBoost Training Error

- Theorem:

- write ϵ_t as $\frac{1}{2} - \gamma_t$ [$\gamma_t = \text{"edge"}$]
- then

$$\begin{aligned} \text{training error}(H_{\text{final}}) &\leq \prod_t \left[2\sqrt{\epsilon_t(1-\epsilon_t)} \right] \\ &= \prod_t \sqrt{1-4\gamma_t^2} \\ &\leq \exp \left(-2 \sum_t \gamma_t^2 \right) \end{aligned}$$

- so: if $\forall t : \gamma_t \geq \gamma > 0$

then $\text{training error}(H_{\text{final}}) \leq e^{-2\gamma^2 T}$


Decrease with
increase of T

- AdaBoost is adaptive:

- does **not** need to know γ or T a priori

How Will Test Error Behave??

- Expect (a first guess)
 - Training error to continue to drop (or reach 0)
 - Test error increases when H_{final} becomes “too complex”
 - Occam’s razor
 - Overfitting: hard to know when to stop training






Theoretically...

- With high probability

$$\text{generalization error} \leq \text{training error} + \tilde{O}\left(\sqrt{\frac{dT}{m}}\right)$$

- bound depends on
 - m = # training examples
 - d = “complexity” of weak classifiers, VC-dimension
 - T = # rounds
- Generalization error = $E[\text{test error}]$
- **Should overfit with T increase...**



- Test error does **not** increase, even after 1000 rounds
 - Total size > 2,000,000 nodes
- Test error continues to drop even after training error is zero!

	# rounds		
	5	100	1000
train error	0.0	0.0	0.0
test error	8.4	3.3	3.1

- **Occam's razor wrongly predicts "simpler" rule is better?**

Margin Theory Explanation


Based on the concept of margin, Schapire et al. [1998] proved that, given any threshold $\theta > 0$ of margin over the training data D , with probability at least $1 - \delta$, the generalization error of the ensemble $\epsilon_D = P_{\mathbf{x} \sim \mathcal{D}}(f(\mathbf{x}) \neq H(\mathbf{x}))$ is bounded by

$$\begin{aligned} \epsilon_D &\leq P_{\mathbf{x} \sim D}(f(\mathbf{x})H(\mathbf{x}) \leq \theta) + \tilde{O} \left(\sqrt{\frac{d}{m\theta^2} + \ln \frac{1}{\delta}} \right) \\ &\leq 2^T \prod_{t=1}^T \sqrt{\epsilon_t^{1-\theta}(1 - \epsilon_t)^{1+\theta}} + \tilde{O} \left(\sqrt{\frac{d}{m\theta^2} + \ln \frac{1}{\delta}} \right) \end{aligned}$$

- This bound implies that, when other variables are fixed, **the larger the margin over the training data, the smaller the generalization error**

Margin Theory Explanation

- Why AdaBoost tends to be resistant to overfitting? the margin theory answers:
 - It can increase the **ensemble margin** even after the training error reaches zero!






Margin Theory Explanation

中科院计算所

■ Key ideas


- Training error only measures whether classifications are right or wrong
- Should also consider **confidence** of classifications
- Recall: H_{final} is weighted majority vote of weak classifiers

- ## ■ Measure confidence by **margin** = strength of the vote = (weighted fraction voting correctly) - (weighted fraction voting incorrectly)



Empirical Evidence: Margin Distribution

- Margin distribution
 - Cumulative distribution of margins of training examples



	# rounds		
	5	100	1000
train error	0.0	0.0	0.0
test error	8.4	3.3	3.1
% margins ≤ 0.5	7.7	0.0	0.0
minimum margin	0.14	0.52	0.55



Theoretical Evidence

- Theorems
 - Larger margins \Rightarrow better bound on generalization error (independent of number of rounds)
 - Boosting tends to increase margins of training examples (**given weak learning assumption**)
- Tighter bound with margin distribution
 - Minimum margin, media margin, average margin, margin variance...



More...

- Predicts good generalization with no overfitting if:
 - weak classifiers not **too complex** relative to size of training set
 - weak classifiers have large edges (implying large margins)
- For example
 - Boosting decision trees resistant to overfitting since trees often have large edges and limited complexity
- Overfitting may occur if:
 - Overly complex weak classifiers
 - Small edges (**underfitting**)



Practical Advantages of AdaBoost

- Very fast, simple and easy to program
- Few parameters to tune (except T)
- Flexible
 - can combine with (m)any learning algorithm
- Little prior knowledge needed about weak learner
 - Provably effective, provided can consistently find rough rules of thumb
 - Shift in mindset — goal now is merely to **find classifiers barely better than random guessing**
- Versatile
 - can use with data that is textual, numeric, discrete, etc.
 - has been extended to learning problems well beyond binary classification



Disadvantages of AdaBoost

- Performance of AdaBoost depends on data and weak learner.
- Consistent with theory, AdaBoost can fail if
 - weak classifiers are too complex
 - Due to possible overfitting
 - weak classifiers too weak ($\gamma_t \rightarrow 0$ too quickly)
 - Due to underfitting
 - low margins \rightarrow overfitting
- empirically, AdaBoost seems especially susceptible to uniform noise

AdaBoost for Face Detection

(separate slides)



中国科学院计算技术研究所

Institute of Computing Technology, Chinese Academy of Sciences



Outline of This Chapter

- Some philosophy in PR/ML
 - No Free Lunch Theorem
 - Ugly Duckling Theorem
 - Minimum Description Length principle
 - Occam's razor
- Resampling for classifier design
 - Bagging
 - Boosting
 - AdaBoost
 - Active Learning
- Estimating and comparing classifiers
 - Cross validation



Active Learning

- Problem to address
 - Semi-supervised learning
- a.k.a
 - Learning with query; Interactive learning
- Method
 - Human in the loop to label “self-proposed” unlabeled informative samples
 - How to self-proposed samples?
 - pattern that the current classifier is least **certain**
 - pattern that yields the greatest **disagreement** among the committee



Outline of This Chapter

- Some philosophy in PR/ML
 - No Free Lunch Theorem 没有最好的算法/学习器
 - Ugly Duckling Theorem 没有最优的特征
 - Minimum Description Length principle 描述越短越好
 - Occam's razor 越简单越好
- Resampling for classifier design
 - Bagging
 - Boosting
 - AdaBoost
 - Active Learning
- Estimating and comparing classifiers
 - Cross validation



Cross Validation

- Hard if not impossible to **theoretically** assess and compare classifiers
- Heuristic method
 - Cross validation
 - $\rightarrow m$ -fold cross-validation
 - Jackknife (leave-one-out)




Cross Validation

- Wrong!!
 - Turing parameters on the testing set
 - Validating on the training set
- Cross-validation
 - Randomly split the set of labeled training samples D into two parts
 - one as the traditional training set for adjusting model parameters in the classifier
 - The other set — the *validation set* — is used to estimate the generalization validation error

Cross Validation

- Typically, the error on the validation set decreases, but then increases
 - Indication that the classifier may be **overfitting** the training data
- Training or parameter adjustment is **stopped** at the first minimum of the validation error.



- How to split the original training set?
 - heuristics for choosing the portion γ of D to be used as a validation set ($0 < \gamma < 1$)
 - Generally a small portion for validation $\gamma < 0.5$
 - If a classifier has a large number of free parameters or degrees of freedom *e.g.* $\gamma=0.1$
- m -fold Cross Validation
 - training set is randomly divided into m disjoint sets of equal size n/m
 - One set for validation, the other $m-1$ sets for training. And repeat m times...
 - If $m=n$, jackknife (leave-one-out)



延伸阅读

- Boosting & Deep Learning
 - MoE
 - Furong Huang 1 Jordan T. Ash 2 John Langford 3 Robert E. Schapire. Learning Deep ResNet Blocks Sequentially using Boosting Theory. ICML2018
 - M Moghimi, et al., Boosted Convolutional Neural Networks, BMVC2016



谢谢！