

# Supervised Learning

Hong Chang

Institute of Computing Technology,  
Chinese Academy of Sciences

Pattern Recognition and Machine Learning (Fall 2022)

# Outline I

- 1 Linear Models for Regression
  - Linear Regression
  - Probabilistic Interpretation
  - Generalized Linear Regression
  
- 2 Discriminative Classification
  - Logistic Regression
  
- 3 Generative Classification
  - Gaussian Discriminative Analysis
  - Naive Bayes

# A Regression Example

- Consider the house price problem as follows:

living area ( $m^2$ )	location	price (RMB10,000)
65	Zhongguancun	350
100	Shangdi	320
105	Shangdi	380
110	Changping	168
$\vdots$	$\vdots$	$\vdots$

- Input samples:  $\mathcal{X} = \{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(N)}\}$ 
  - Features include  $x_1^{(i)}$ : living area;  $x_2^{(i)}$ : location
- Output target:  $\mathcal{Y} = \{y^{(1)}, \dots, y^{(N)}\}$
- Linear regression model:**

$$f_{\mathbf{w}}(\mathbf{x}) = w_0 + w_1 x_1 + \dots + w_D x_D = \sum_{j=0}^D w_j x_j = \mathbf{w}^T \mathbf{x}$$

- $\mathbf{w} = (w_0, w_1, \dots, w_D)^T$ ,  $\mathbf{x} = (1, x_1, \dots, x_D)^T$

# Least Mean Square Algorithm

- Cost function:

$$J(\mathbf{w}) = \frac{1}{2} \sum_{i=1}^N (f_{\mathbf{w}}(\mathbf{x}^{(i)}) - y^{(i)})^2$$

- Gradient descent:

$$w_j := w_j - \alpha \frac{\partial}{\partial w_j} J(\mathbf{w})$$

- For a single training example:

$$w_j := w_j + \alpha (y^{(i)} - f_{\mathbf{w}}(\mathbf{x}^{(i)})) x_j^{(i)}$$

# Least Mean Square Algorithm (2)

- For a training set of more than one example:

- **Batch gradient descent:**

Repeat until convergence {

$$w_j := w_j + \alpha \sum_{i=1}^N (y^{(i)} - f_w(\mathbf{x}^{(i)})) x_j^{(i)}$$

}

- **Stochastic gradient descent:**

Loop {

For  $i = 1$  to  $N$  {

$$w_j := w_j + \alpha (y^{(i)} - f_w(\mathbf{x}^{(i)})) x_j^{(i)}$$

}

}

# Probabilistic Assumptions

- Relate the inputs and target via:

$$y^{(i)} = \mathbf{w}^T \mathbf{x}^{(i)} + \epsilon^{(i)}$$

with  $\epsilon^{(i)}$  IID Gaussian distributed, i.e.,

$$p(\epsilon^{(i)}) = \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(\epsilon^{(i)})^2}{2\sigma^2}\right)$$

- Thus,

$$p(y^{(i)}|\mathbf{x}^{(i)}; \mathbf{w}) = \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(y^{(i)} - \mathbf{w}^T \mathbf{x}^{(i)})^2}{2\sigma^2}\right)$$

- Likelihood function:**

$$L(\mathbf{w}) = L(\mathbf{w}; \mathbf{X}, \mathbf{y}) = p(\mathbf{y}|\mathbf{X}; \mathbf{w}) = \prod_{i=1}^N p(y^{(i)}|\mathbf{x}^{(i)}; \mathbf{w})$$

# Maximum Likelihood

- Log likelihood

$$\begin{aligned}\mathcal{L}(\mathbf{w}) = \log L(\mathbf{w}) &= \log \prod_{i=1}^N \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(y^{(i)} - \mathbf{w}^T \mathbf{x}^{(i)})^2}{2\sigma^2}\right) \\ &= N \log \frac{1}{\sqrt{2\pi}\sigma} - \frac{1}{\sigma^2} \cdot \frac{1}{2} \sum_{i=1}^N (y^{(i)} - \mathbf{w}^T \mathbf{x}^{(i)})^2\end{aligned}$$

- Maximizing  $\mathcal{L}(\mathbf{w})$  gives the same answer as minimizing

$$\frac{1}{2} \sum_{i=1}^N (y^{(i)} - \mathbf{w}^T \mathbf{x}^{(i)})^2$$


- Under the previous probabilistic assumptions on the data, **least-square regression** corresponds to finding the **maximum likelihood estimate** of  $\mathbf{w}$ .

# Locally Weighted Linear Regression (LWR)

- The LWR algorithm:
  - 1 Fit  $\mathbf{w}$  to minimize  $\sum_i \theta^{(i)} (y^{(i)} - \mathbf{w}^T \mathbf{x}^{(i)})^2$
  - 2 Output  $\mathbf{w}^T \mathbf{x}$
- A standard choice for the weights is:

$$\theta^{(i)} = \exp\left(-\frac{\|\mathbf{x}^{(i)} - \mathbf{x}\|^2}{2\tau^2}\right)$$




- More general form is possible
- The weights depend on the particular point  $\mathbf{x}$  at which we are trying to evaluate.
- LWR is an example of **non-parametric** algorithm.





# Locally Weighted Linear Regression (LWR)

- The effect of  $\tau$  in the weight function:




# Linear Regression with Nonlinear Basis

- Consider linear combinations of **nonlinear basis functions** of the input variables:

$$f_{\mathbf{w}}(\mathbf{x}) = w_0 + \sum_{j=1}^{M-1} w_j \phi_j(\mathbf{x}) = \mathbf{w}^T \boldsymbol{\phi}(\mathbf{x})$$

- $\boldsymbol{\phi} = (\phi_0, \phi_1, \dots, \phi_{M-1})$ ,  $\phi_0(\mathbf{x}) = 1$
- Examples of basis functions:
  - $\phi(x) = (1, x, x^2, \dots, x^{M-1})$
  - $\phi_j(x) = \exp(-\frac{(x-\mu_j)^2}{2s^2})$
- Can still use least squares method to estimate
- Linear method can **model nonlinear functions** through using nonlinear basis functions.

# Geometry of Least Squares



- $\varphi_{j+1} = [\phi_j(\mathbf{x}^{(1)}), \dots, \phi_j(\mathbf{x}^{(N)})] \in \mathbb{R}^N, j = 0, \dots, M - 1$
- If  $M < N$ , vectors  $\varphi_j$  span a linear subspace  $\mathcal{S}$  of dimensionality  $M$
- $f_{\mathbf{w}}(\mathbf{X}) \in \mathbb{R}^N$  is linear combination of vectors  $\varphi_j$
- The least-squares regression function is the **orthogonal projection of the target vector  $\mathbf{y}$  onto the subspace  $\mathcal{S}$**

# Two-Class Example

- Let  $P(y = 1|\mathbf{x}) = t, P(y = 0|\mathbf{x}) = 1 - t$
- Classification rule:

$$\text{Choose} = \begin{cases} y = 1 & \text{if } t > 0.5 \\ y = 0 & \text{otherwise} \end{cases}$$

- Equivalent test for classification rule: choose  $y = 1$  if

$$\frac{t}{1-t} > 1 \text{ or } \log \frac{t}{1-t} > 0$$


where  $\frac{t}{1-t}$  is called the **odds** of  $t$  and  $\log \frac{t}{1-t}$  is called the **log odds** or **logit function** of  $t$ .

# Logistic Regression

- The posterior probability (or the hypothesis) is written as:

$$P(y = 1|\mathbf{x}) = f_{\mathbf{w}}(\mathbf{x}) = g(\mathbf{w}^T \mathbf{x}) = \frac{1}{1 + \exp(-\mathbf{w}^T \mathbf{x})}$$

where  $g(\cdot)$  is called **logistic** or **sigmoid function**.



- The **derivative** of the sigmoid function:

$$g'(z) = g(z)(1 - g(z))$$

- Note that this is a model for classification rather than regression.

# Parameter Estimation

- Since

$$\begin{aligned}P(y = 1|\mathbf{x}; \mathbf{w}) &= f_{\mathbf{w}}(\mathbf{x}) \\P(y = 0|\mathbf{x}; \mathbf{w}) &= 1 - f_{\mathbf{w}}(\mathbf{x}),\end{aligned}$$

thus

$$P(y|\mathbf{x}; \mathbf{w}) = (f_{\mathbf{w}}(\mathbf{x}))^y (1 - f_{\mathbf{w}}(\mathbf{x}))^{1-y}.$$

- The **likelihood** of the parameters (given  $N$  training examples):

$$L(\mathbf{w}) = \prod_{i=1}^N P(y^{(i)}|\mathbf{x}^{(i)}; \mathbf{w}) = \prod_{i=1}^N (f_{\mathbf{w}}(\mathbf{x}^{(i)}))^{y^{(i)}} (1 - f_{\mathbf{w}}(\mathbf{x}^{(i)}))^{1-y^{(i)}}$$

- The **log likelihood** function:

$$\mathcal{L}(\mathbf{w}) = \log L(\mathbf{w}) = \sum_{i=1}^N \{y^{(i)} \log f_{\mathbf{w}}(\mathbf{x}^{(i)}) + (1 - y^{(i)}) \log(1 - f_{\mathbf{w}}(\mathbf{x}^{(i)}))\}$$

## Parameter Estimation (2)

- As for one training example  $(\mathbf{x}, y)$ , maximizing the log likelihood gives

$$\frac{\partial}{\partial w_j} \mathcal{L}(\mathbf{w}) = (y - f_{\mathbf{w}}(\mathbf{x})) x_j$$

- Use the fact that  $f_{\mathbf{w}}(\mathbf{x}) = g(\mathbf{w}^T \mathbf{x})$  and  $g'(z) = g(z)(1 - g(z))$
- This therefore gives the **stochastic gradient ascent** rule:

$$w_j := w_j + \alpha (y^{(i)} - f_{\mathbf{w}}(\mathbf{x}^{(i)})) x_j^{(i)}$$

- The solution is **identical to least mean squares update rule for linear regression!**

# Logistic Discrimination

- Logistic discrimination does not model the class-conditional densities  $p(\mathbf{x}|C_k)$  but their **ratio**.
- **Parametric classification** approach (studied later) learns the classifier by estimating the parameters of  $p(\mathbf{x}|C_k)$ ; **logistic discrimination estimates the parameters of the discriminant function directly**.



# Discriminative vs. Generative classification (Revisited)

- **Discriminative** classification: model the posterior probabilities or learn discriminant function directly.
  - Compute  $P(C_k|\mathbf{x})$ : such as logistic regression.
  - Compute discriminant function: such as perceptron, SVM.
  - Makes an assumption on the form of the discriminants, but not the densities.
- **Generative** classification: explicitly or implicitly model the distribution of inputs as well as outputs.
  - Assume a model for the class-conditional probability densities  $p(\mathbf{x}|C_k)$ .
  - Estimate  $p(\mathbf{x}|C_k)$  and  $P(C_k)$  from data, and apply Bayes' rule to compute the posterior probabilities  $P(C_k|\mathbf{x})$ .
  - Perform optimal classification based on  $P(C_k|\mathbf{x})$ .

# Bernoulli and Multinomial

- **Bernoulli**: the distribution for a single binary variable  $x \in \{0, 1\}$ , governed by a single continuous parameter  $\beta \in [0, 1]$  which represents the probability of  $x = 1$ .

$$\text{Bern}(x|\beta) = \beta^x(1 - \beta)^{1-x}$$


- **Binomial**: gives the probability of observing  $m$  occurrences of  $x = 1$  in a set of  $N$  samples from a Bernoulli distribution.

$$\text{Bin}(m|N, \beta) = \binom{N}{m} \beta^m(1 - \beta)^{N-m}$$

- **Multinomial**: multivariate generalization of binomial, gives the distribution over counts  $m_k$  for  $K$ -state discrete variable to be in state  $k$ :

$$\text{Mult}(m_1, \dots, m_K|N, \beta) = \binom{N}{m_1 m_2 \dots m_K} \prod_{k=1}^K \beta_k^{m_k}$$

# Multivariate Normal Distribution



$$\mathbf{x} \sim \mathcal{N}_D(\boldsymbol{\mu}, |\boldsymbol{\Sigma}|)$$
$$p(\mathbf{x}) = \frac{1}{(2\pi)^{D/2} |\boldsymbol{\Sigma}|^{1/2}} \exp\left(-\frac{1}{2}(\mathbf{x} - \boldsymbol{\mu})^T \boldsymbol{\Sigma}^{-1}(\mathbf{x} - \boldsymbol{\mu})\right)$$

## Multivariate Normal Distribution (2)

- **Mahalanobis distance** measures the distance from  $\mathbf{x}$  to  $\mu$  in terms of  $\Sigma$ :

$$(\mathbf{x} - \mu)^T \Sigma^{-1} (\mathbf{x} - \mu)$$

- $(\mathbf{x} - \mu)^T \Sigma^{-1} (\mathbf{x} - \mu) = c^2$  is the **D-dimensional hyperellipsoid** centered at  $\mu$ . Its shape and orientation are defined by  $\Sigma$ .
- **Euclidean distance** is a special case of Mahalanobis distance when  $\Sigma = s^2 \mathbf{I}$ ; the hyperellipsoid degenerates into a **hypersphere**.

# Gaussian Discriminant Analysis (GDA)

- GDA models  $p(\mathbf{x}|y)$  using a **multivariate normal distribution**:

$$\begin{aligned}y &\sim \text{Bernoulli}(\beta) \\ \mathbf{x}|y=0 &\sim \mathcal{N}(\boldsymbol{\mu}_0, \boldsymbol{\Sigma}) \\ \mathbf{x}|y=1 &\sim \mathcal{N}(\boldsymbol{\mu}_1, \boldsymbol{\Sigma})\end{aligned}$$

- The parameters of the model are  $\{\beta, \boldsymbol{\mu}_0, \boldsymbol{\mu}_1, \boldsymbol{\Sigma}\}$ . Note that common covariance matrix is usually used.

# Gaussian Discriminant Analysis (2)


- The **log likelihood** of the data is:

$$\begin{aligned}\mathcal{L}(\beta, \mu_0, \mu_1, \Sigma) &= \log \prod_{i=1}^N p(\mathbf{x}^{(i)}, y^{(i)}; \beta, \mu_0, \mu_1, \Sigma) \\ &= \log \prod_{i=1}^N p(\mathbf{x}^{(i)} | y^{(i)}; \beta, \mu_0, \mu_1, \Sigma) p(y^{(i)}; \beta)\end{aligned}$$

- The **maximum likelihood estimates** are:

$$\begin{aligned}\hat{\beta} &= \frac{1}{N} \sum_{i=1}^N 1\{y^{(i)} = 1\} \\ \hat{\mu}_k &= \frac{\sum_{i=1}^N 1\{y^{(i)} = k\} \mathbf{x}^{(i)}}{\sum_{i=1}^N 1\{y^{(i)} = k\}}, \quad k = \{0, 1\} \\ \hat{\Sigma} &= \frac{1}{N} \sum_{i=1}^N (\mathbf{x}^{(i)} - \hat{\mu}_{y^{(i)}})(\mathbf{x}^{(i)} - \hat{\mu}_{y^{(i)}})^T\end{aligned}$$

# Illustration of GDA



# Classification with Discriminant Functions

- Assume class-conditional densities  $p(\mathbf{x}|C_i)$  are Gaussian with common covariance  $\Sigma$ :

$$p(\mathbf{x}|C_i) = \frac{1}{(2\pi)^{D/2}|\Sigma|^{1/2}} \exp\left(-\frac{1}{2}(\mathbf{x} - \mu_i)^T \Sigma^{-1}(\mathbf{x} - \mu_i)\right)$$

- Discriminant functions:


$$\begin{aligned} g_i(\mathbf{x}) &= \log P(\mathbf{x}|C_i) + \log P(C_i) \\ &= -\frac{D}{2} \log 2\pi - \frac{1}{2} \log |\Sigma| - \frac{1}{2}(\mathbf{x} - \mu_i)^T \Sigma^{-1}(\mathbf{x} - \mu_i) \\ &\quad + \log P(C_i) \\ &= \mathbf{w}_i^T \mathbf{x} + w_{i0} \end{aligned}$$

where

$$\begin{aligned} \mathbf{w}_i &= \Sigma^{-1} \mu_i \\ w_{i0} &= -\frac{1}{2} \mu_i^T \Sigma^{-1} \mu_i + \log P(C_i) \end{aligned}$$



# Linear Discriminant Function



- If class-conditional densities are Gaussian with a common covariance, the discriminant functions are linear.
- Given samples  $\mathcal{X}$ , we can find ML estimates for  $\mu_i, \Sigma$  and plug them into the discriminant functions.

# GDA and Logistic Regression

- If we view  $P(y = 1|\mathbf{x}; \beta, \mu_0, \mu_1, \Sigma)$  as a function of  $\mathbf{x}$ , it can be expressed in the form:

$$P(y = 1|\mathbf{x}; \beta, \mu_0, \mu_1, \Sigma) = \frac{1}{1 + \exp(-\mathbf{w}^T \mathbf{x})},$$

where  $\mathbf{w}$  is some appropriate function of  $\beta, \mu_0, \mu_1, \Sigma$ . This is exactly the form that **logistic regression** (a discriminative algorithm) used to model  $P(y = 1|\mathbf{x})$ .

- GDA** makes stronger modeling assumptions (i.e.,  $p(\mathbf{x}|y)$  is multivariate Gaussian) about the data, and is more **data efficient**.
- Logistic regression** makes weaker assumptions, and is **more robust to deviations from modeling assumptions**.

# Email Spam Filter

- Classifying emails to **spam** ( $y = 1$ ) or **non-spam** ( $y = 0$ )
- One example of broader set of problems called **text classification**
- Feature vector:**

$$\mathbf{x} = \begin{bmatrix} 1 \\ 0 \\ 1 \\ \vdots \\ 0 \end{bmatrix} \begin{matrix} a \\ aardwolf \\ buy \\ \vdots \\ zygmurgy \end{matrix}$$

- $\mathbf{x} \in \{0, 1\}^D$ ,  $D$  is the size of the set of words (vocabulary).
- Word selection: include words that appear in the emails but not in a dictionary; exclude very high frequency words; etc.
- $x_j$ 's can take more general values in  $\{1, 2, \dots, k_j\}$ , then the model  $p(\mathbf{x}_j|y)$  is multinomial rather than binomial.

## Email Spam Filter (2)

- To model  $p(\mathbf{x}|y)$ , we assume that  $x_j$ 's are **conditionally independent** given  $y$ . This is **Naive Bayes assumption**.
- In Naive Bayes classifier:

$$p(x_1, \dots, x_D|y) = p(x_1|y)p(x_2|y) \dots p(x_D|y) = \prod_{j=1}^D p(x_j|y)$$

- Given a training set  $\{\mathbf{x}^{(i)}, y^{(i)}\}, i = 1, \dots, N$ , the **joint log likelihood** of the data:

$$\mathcal{L}(p(x_j|y), P(y)) = \log \prod_{i=1}^N p(\mathbf{x}^{(i)}, y^{(i)})$$

## Email Spam Filter (3)

- Maximizing  $\mathcal{L}$  w.r.t.  $p(x_j|y)$ ,  $P(y)$  gives the following estimates:

$$\begin{aligned}
 p(x_j = 1|y = 1) &= \frac{\sum_{i=1}^N 1\{x_j^{(i)} = 1 \wedge y^{(i)} = 1\}}{\sum_{i=1}^N 1\{y^{(i)} = 1\}} \\
 p(x_j = 1|y = 0) &= \frac{\sum_{i=1}^N 1\{x_j^{(i)} = 1 \wedge y^{(i)} = 0\}}{\sum_{i=1}^N 1\{y^{(i)} = 0\}} \\
 p(y = 1) &= \frac{\sum_{i=1}^N 1\{y^{(i)} = 1\}}{N}
 \end{aligned}$$

## Email Spam Filter (4)

- To make a prediction on a new example  $\mathbf{x}$ , we simply calculate:

$$\begin{aligned} p(y = 1|\mathbf{x}) &= \frac{p(\mathbf{x}|y = 1)p(y = 1)}{p(\mathbf{x})} \\ &= \frac{(\prod_{j=1}^D p(x_j|y = 1))p(y = 1)}{(\prod_{j=1}^D p(x_j|y = 1))p(y = 1) + (\prod_{j=1}^D p(x_j|y = 0))p(y = 0)} \end{aligned}$$

- When the original, continuous-valued attributes are not well-modeled by a multivariate normal distribution, discretizing the features and using Naive Bayes (instead of GDA) will often result in a better classifier.

# Laplace Smoothing

- If you haven't seen some word before in your finite training set, the Naive Bayes classifier does not know how to make a decision.
- To estimate the mean of **multinomial random variable**  $x$  taking values  $\{1, \dots, k\}$ , given  $N$  observations  $\{x^{(1)}, \dots, x^{(N)}\}$ , the maximum likelihood estimates are:

$$p(x = j) = \frac{\sum_{i=1}^N \mathbf{1}\{x^{(i)} = j\}}{N}, \quad j = 1, \dots, k$$

- **Laplace smoothing:**

$$p(x = j) = \frac{\sum_{i=1}^N \mathbf{1}\{x^{(i)} = j\} + 1}{N + k}, \quad j = 1, \dots, k$$

- $\sum_{j=1}^k p(x = j) = 1$  still holds
- $p(x = j) \neq 0$  for all  $j$

## Laplace Smoothing (2)

- Returning to **Naive Bayes classifier**, with **Laplace smoothing** we obtain the following estimates:

$$p(x_j = 1 | y = 1) = \frac{\sum_{i=1}^N 1\{x_j^{(i)} = 1 \wedge y^{(i)} = 1\} + 1}{\sum_{i=1}^N 1\{y^{(i)} = 1\} + 2}$$

$$p(x_j = 1 | y = 0) = \frac{\sum_{i=1}^N 1\{x_j^{(i)} = 1 \wedge y^{(i)} = 0\} + 1}{\sum_{i=1}^N 1\{y^{(i)} = 0\} + 2}$$

- In practice, it usually doesn't matter whether we apply Laplace smoothing to  $p(y = 1)$  or not, since we typically have a fair fraction between spam and non-spam emails.



# Event Models for Text Classification

- **Naive Bayes** uses **multivariate Bernoulli event model**: the probability of an email was given by  $P(y) \prod_{j=1}^D p(x_j|y)$ .
- A different way to represent emails:  $\mathbf{x} = (x_1, \dots, x_M)$ ,  $x_j$  denotes the  $j^{th}$  word in the email, taking values in  $\{1, \dots, |V|\}$ ;  $V$  is the vocabulary;  $M$  is the length of the email.
- **Multinomial event model**:  $p(x_j|y)$  is now multinomial. The overall probability of an email is  $P(y) \prod_{j=1}^M p(x_j|y)$ .

## Event Models for Text Classification (2)

- Given a training set  $\{\mathbf{x}^{(i)}, y^{(i)}\}, i = 1, \dots, N$ , where  $\mathbf{x}^{(i)} = (x_1^{(i)}, \dots, x_{M_i}^{(i)})^T$  ( $M_i$  is the number of words in  $i^{th}$  training example), the maximum likelihood estimates of the model are:

$$p(x_j = k | y = 1) = \frac{\sum_{i=1}^N \sum_{j=1}^{M_i} 1\{x_j^{(i)} = k \wedge y^{(i)} = 1\}}{\sum_{i=1}^N 1\{y^{(i)} = 1\} M_i}, \text{ for any } j$$

$$p(x_j = k | y = 0) = \frac{\sum_{i=1}^N \sum_{j=1}^{M_i} 1\{x_j^{(i)} = k \wedge y^{(i)} = 0\}}{\sum_{i=1}^N 1\{y^{(i)} = 0\} M_i}, \text{ for any } j$$

$$P(y = 1) = \frac{\sum_{i=1}^N 1\{y^{(i)} = 1\}}{N}$$

# Naive Bayes vs. Logistic Regression

- **Asymptotic comparison** (# training examples  $\rightarrow$  infinite)
- When model assumption is correct:
  - Naive Bayes (NB), Logistic regression (LR) produce identical classifier
- When model assumption is incorrect:
  - LR is less biased - does not assume conditional independence
  - therefore expected to outperform NB.

## Naive Bayes vs. Logistic Regression (2)

- **Non-asymptotic comparison** (see [Ng and Jordan, 2002])
- Convergence rate of parameter estimation - how many training examples needed to achieve good estimates?
  - NB order  $\log D$  (where  $D = \#$  of attributes in  $\mathbf{X}$ )
  - LR order  $D$
- NB converges more quickly to its asymptotic estimates.

# K-fold Cross Validation

- The data set  $\mathcal{X}$  is randomly partitioned into  $K$  equal-sized subsets  $\mathcal{X}_i, i = 1, \dots, K$ , called folds.
- **Stratification**: the class distributions in different subsets are kept roughly the same (stratified cross validation).
- $K$  training/validation set pairs  $\{\mathcal{T}_i, \mathcal{V}_i\}_{i=1}^K$ :

$$\mathcal{T}_1 = \mathcal{X}_2 \cup \mathcal{X}_3 \cup \dots \cup \mathcal{X}_K, \quad \mathcal{V}_1 = \mathcal{X}_1$$

$$\mathcal{T}_2 = \mathcal{X}_1 \cup \mathcal{X}_3 \cup \dots \cup \mathcal{X}_K, \quad \mathcal{V}_2 = \mathcal{X}_2$$

$$\vdots$$

$$\mathcal{T}_K = \mathcal{X}_1 \cup \mathcal{X}_2 \cup \dots \cup \mathcal{X}_{K-1}, \quad \mathcal{V}_K = \mathcal{X}_K$$

# Leave-One-Out

- If  $N$  is small,  $K$  should be large to allow large enough training sets.
- One extreme case of  $K$ -fold cross validation is **leave-one-out** where only one instance is left out as the validation set and the remaining  $N - 1$  for training, giving  $N$  separate training/validation set pairs.

# Bootstrapping

- When the sample size  $N$  is very small, a better alternative to cross validation is **bootstrapping**.
- Bootstrapping generates new samples, each of size  $N$ , by drawing instances randomly from the original sample **with replacement**.
- The bootstrap samples usually overlap more than the cross validation samples and hence their estimates are more dependent.
- Probability that an instance is not chosen after  $N$  random draws:

$$\left(1 - \frac{1}{N}\right)^N \approx e^{-1} = 0.368$$

So each bootstrap sample contains only approximately **63.2%** of the instances.

- Multiple bootstrap samples are used to maximize the chance that the system is trained on all the instances.

# Error Measure

- If 0 – 1 loss is used, error calculations will be based on the **confusion matrix**:

True Class	Predicted Class	
	Yes	No
Yes	TP: true positive	FN: false negative
No	FP: false postivie	TN: true negative

- Error rate:

$$\text{Error rate} = \frac{|FN| + |FP|}{|TP| + |FP| + |TN| + |FN|} = \frac{|FN| + |FP|}{N}$$

where  $N$  is the total number of instances in the validation set.


- For more general loss functions, the risks should be measured instead.
- For  $K > 2$  classes, the class confusion matrix is a  $K \times K$  matrix such that its  $(i, j)$ -th entry contains the number of instances that belong to  $C_i$  but are assigned to  $C_j$ .



# ROC Curve

- Receiver operating characteristics (ROC) curve:

$$\text{Hit rate } \frac{|TP|}{|TP| + |FN|} \text{ vs. False alarm rate } \frac{|FP|}{|FP| + |TN|}$$



- The curve is obtained by varying a certain parameter (e.g., threshold for making decision) of the classification algorithm.
- The **area under curve (AUC)** is often used as a performance measure.

# Point Estimation vs. Interval Estimation

- A **point estimator** (e.g., MLE) specifies a value for a parameter  $\theta$ .
- An **interval estimator** specifies an **interval** within which  $\theta$  lies with a certain **degree of confidence**.
- The probability distribution of the point estimator is used to obtain an interval estimator.

# Example: Estimation of Mean of Normal Distribution

- Given an i.i.d. sample  $\mathcal{X} = \{x^{(i)}\}_{i=1}^N$ , where

$$x^{(i)} \sim \mathcal{N}(\mu, \sigma^2)$$

- Point estimation of the mean  $\mu$ :

$$m = \frac{\sum_i x^{(i)}}{N} \sim \mathcal{N}(\mu, \sigma^2/N)$$

since  $x^{(i)}$  are i.i.d.

- Define a statistic  $\mathcal{Z}$  with a **unit normal distribution**:

$$\mathcal{Z} = \frac{\sqrt{N}(m - \mu)}{\sigma} \sim \mathcal{N}(0, 1)$$

# Two-Sided Confidence Interval


- For a unit normal distribution, about 95% of  $Z$  lies in  $(-1.96, 1.96)$ :

$$P(-1.96 < \frac{\sqrt{N}(m - \mu)}{\sigma} < 1.96) = 0.95$$

or

$$P(m - 1.96 \frac{\sigma}{\sqrt{N}} < \mu < m + 1.96 \frac{\sigma}{\sqrt{N}}) = 0.95$$

meaning that with 95% confidence,  $\mu$  lies within  $1.96\sigma/\sqrt{N}$  units of the sample mean  $m$ .



## Two-Sided Confidence Interval (2)

- Let us denote  $z_\alpha$  such that

$$P(\mathcal{Z} > z_\alpha) = \alpha, \quad 0 < \alpha < 1$$

- Because  $\mathcal{Z}$  is symmetric around the mean, we have

$$P(\mathcal{Z} < -z_{\alpha/2}) = P(\mathcal{Z} > z_{\alpha/2}) = \alpha/2$$

- Hence, for any specified level of confidence  $1 - \alpha$ , we have


$$P(-z_{\alpha/2} < \frac{\sqrt{N}(m - \mu)}{\sigma} < z_{\alpha/2}) = 1 - \alpha$$

or

$$P(m - z_{\alpha/2} \frac{\sigma}{\sqrt{N}} < \mu < m + z_{\alpha/2} \frac{\sigma}{\sqrt{N}}) = 1 - \alpha$$

# One-Sided Upper Confidence Interval

- $P(Z < 1.64) = 0.95$



- 95% one-sided upper confidence interval for  $\mu$ :

$$P\left(\frac{\sqrt{N}(m - \mu)}{\sigma} < 1.64\right) = 0.95$$

or

$$P\left(m - 1.64 \frac{\sigma}{\sqrt{N}} < \mu\right) = 0.95$$

# One-Sided Upper Confidence Interval (2)

- The **one-sided upper confidence interval** defines a **lower bound** for  $\mu$ .
- In general, a  $100(1 - \alpha)$  percent one-sided upper confidence interval for  $\mu$  can be computed from

$$P\left(m - z_{\alpha} \frac{\sigma}{\sqrt{N}} < \mu\right) = 1 - \alpha$$

- The one-sided lower confidence interval that defines an upper bound can be calculated similarly.

# $t$ Distribution

- When the variance  $\sigma^2$  is not known, it can be replaced by the sample variance

$$S^2 = \frac{\sum_i (x^{(i)} - m)^2}{N - 1}$$

- The statistic  $\sqrt{N}(m - \mu)/S$  follows a  $t$  distribution with  $N - 1$  degrees of freedom:

$$\sqrt{N}(m - \mu)/S \sim t_{N-1}$$

- For any  $\alpha \in (0, 1/2)$ ,

$$P(-t_{\alpha/2, N-1} < \frac{\sqrt{N}(m - \mu)}{S} < t_{\alpha/2, N-1}) = 1 - \alpha$$

or

$$P(m - t_{\alpha/2, N-1} \frac{S}{\sqrt{N}} < \mu < m + t_{\alpha/2, N-1} \frac{S}{\sqrt{N}}) = 1 - \alpha$$



## $t$ Distribution (2)

- One-sided confidence intervals can also be defined.
- The  $t$  distribution has **larger spread** (longer tails) than the unit normal distribution, and generally the **interval** given by the  $t$  distribution is **larger** due to additional uncertainty about the unknown variance.
- As  $N$  increases, the  $t$  distribution will get closer to the normal distribution.

# Hypothesis Testing

- Instead of explicitly estimating some parameters, hypothesis testing uses the sample to test some hypothesis concerning the parameters, e.g., whether the mean is  $< 0.02$ .
- **Hypothesis testing:**
  - We define a statistic that obeys a certain distribution if the hypothesis is correct.
  - If the random sample is consistent with the hypothesis under consideration (i.e., if the statistic calculated from the sample has a high enough probability of being drawn from the probability), the hypothesis is accepted.
  - Otherwise the hypothesis is rejected.
- There are parametric and nonparametric tests. Only **parametric tests** are considered here.

# Null Hypothesis

- Given a sample from a normal distribution  $\mathcal{N}(\mu, \sigma^2)$  where  $\sigma$  is known but  $\mu$  is unknown.
- We want to test a specific hypothesis about  $\mu$ , called the **null hypothesis**, e.g.,

$$H_0 : \mu = \mu_0$$

against the **alternative hypothesis**

$$H_1 : \mu \neq \mu_0$$

for some specified constant  $\mu_0$

# Two-Sided Test

- We accept  $H_0$  if the sample mean  $m$ , a point estimate of  $\mu$ , is not too far from  $\mu_0$ .
- We accept  $H_0$  with level of significance  $\alpha$  if  $\mu_0$  lies in the  $100(1 - \alpha)$  percent confidence interval, i.e.,

$$\frac{\sqrt{N}(m - \mu_0)}{\sigma} \in (-z_{\alpha/2}, z_{\alpha/2})$$

# One-Sided Test

- Null and alternative hypotheses:

$$H_0 : \mu \leq \mu_0$$

$$H_1 : \mu > \mu_0$$

- We accept  $H_0$  with level of significance  $\alpha$  if

$$\frac{\sqrt{N}(m - \mu_0)}{\sigma} \in (-\infty, z_\alpha)$$

# $t$ Tests

- If the variance  $\sigma^2$  is not known, the sample variance  $S^2$  will be used instead and so

$$\frac{\sqrt{N}(m - \mu_0)}{S} \sim t_{N-1}$$

- Two-sided  $t$  test:

$$H_0 : \mu = \mu_0$$

$$H_1 : \mu \neq \mu_0$$

We accept at significance level  $\alpha$  if

$$\frac{\sqrt{N}(m - \mu_0)}{S} \in (-t_{\alpha/2, N-1}, t_{\alpha/2, N-1})$$

- One-side  $t$  test can be defined similarly.

# Binomial Test

- Single training/validation set pair: a classifier is trained on a training set  $\mathcal{T}$  and tested on a validation set  $\mathcal{V}$ .
- Let  $p$  be the (unknown) probability that the classifier makes a misclassification error.
- For the  $i$ th instance in  $\mathcal{V}$ , we define  $x^{(i)}$  as a **Bernoulli variable** to denote the correctness of the classifier's decision:

$$x^{(i)} = \begin{cases} 1 & \text{with probability } p \\ 0 & \text{with probability } 1 - p \end{cases}$$

- **Point estimate** of  $p$ :

$$\hat{p} = \frac{\sum_i x^{(i)}}{|\mathcal{V}|} = \frac{\sum_i x^{(i)}}{N}$$

## Binomial Test (2)

- Hypothesis test:

$$H_0 : p \leq p_0 \text{ vs. } H_1 : p > p_0$$

Given that the classifier makes  $e$  errors on  $\mathcal{V}$  with  $|\mathcal{V}| = N$ , can we say that the classifier has error probability  $p_0$  or less?

- Let  $X = \sum_{i=1}^N x^{(i)}$  denote the number of errors on  $\mathcal{V}$ .
- Because  $x^{(i)}$  are independent Bernoulli variables,  $X$  is **binomial**:

$$P(X = j) = \binom{N}{j} p^j (1 - p)^{N-j}$$

- Under the null hypothesis  $p \leq p_0$ , so the probability that there are  $e$  errors or less is

$$P(X \leq e) = \sum_{j=1}^e \binom{N}{j} p_0^j (1 - p_0)^{N-j}$$

- Binomial test**: accept  $H_0$  if  $P(X \leq e) < 1 - \alpha$ ; reject otherwise.



# Paired $t$ Test

- Multiple training/validation set pairs: if we run the algorithm  $K$  times on  $K$  training/validation set pairs, we get  $K$  error probabilities,  $p_k, k = 1, \dots, K$ , on the  $K$  validation sets.
- We can use the **paired  $t$  test** to determine whether to accept the null hypothesis  $H_0$  that the classifier has error probability  $p_0$  or less at significance level  $\alpha$ .
- Bernoulli variables:

$$x_k^{(i)} = \begin{cases} 1 & \text{if classifier trained on } \mathcal{T}_k \text{ makes an error on instance } i \text{ of } \mathcal{V}_k \\ 0 & \text{otherwise} \end{cases}$$

- **Test statistic:**

$$\frac{\sqrt{K}(m - p_0)}{S} \sim t_{K-1}$$

where

$$p_i = \frac{\sum_{i=1}^N x^{(i)}}{N}, m = \frac{\sum_{k=1}^K p_k}{K}, S^2 = \frac{\sum_{k=1}^K (p_k - m)^2}{K - 1}$$

# K-Fold Cross-Validated Paired $t$ Test

- Given two classification algorithms and a data set, we want to compare and test whether the two algorithms construct classifiers that have the same expected error rate on a new instance.
- K-fold cross validation** is used to obtain  $K$  training/validation set pairs,  $\{(\mathcal{T}_k, \mathcal{V}_k)\}_{k=1}^K$ .
- For each training/validation set pair  $(\mathcal{T}_k, \mathcal{V}_k)$ , the two classification algorithms are trained on  $\mathcal{T}_k$  and tested on  $\mathcal{V}_k$ , with error probabilities  $p_k^1$  and  $p_k^2$ .
- We define

$$p_k = p_k^1 - p_k^2, k = 1, \dots, K$$

The distribution of  $p_k$  is used for hypothesis testing.

- if  $p_k^1$  and  $p_k^2$  are both normally distributed, then  $p_k$  should also be **normal** (with mean  $\mu$ ).

# $K$ -Fold Cross-Validated Paired $t$ Test (2)

- Null and alternative hypotheses:

$$H_0 : \mu = 0$$

$$H_1 : \mu \neq 0$$

- Test statistic:

$$\frac{\sqrt{K}(m - 0)}{S} = \frac{\sqrt{K}m}{S} \sim t_{K-1}$$

where

$$m = \frac{\sum_{k=1}^K p_k}{K}, S^2 = \frac{\sum_{k=1}^K (p_k - m)^2}{K - 1}$$

- $K$ -fold CV paired  $t$  test:** accept  $H_0$  at significance level  $\alpha$  if  $\sqrt{K}m/S \in (-t_{\alpha/2, K-1}, t_{\alpha/2, K-1})$ ; reject otherwise.